

Introduction to **MACHINE LEARNING**

BY DATAROCKIE  X DATA SCIENCE CHILL CHILL 

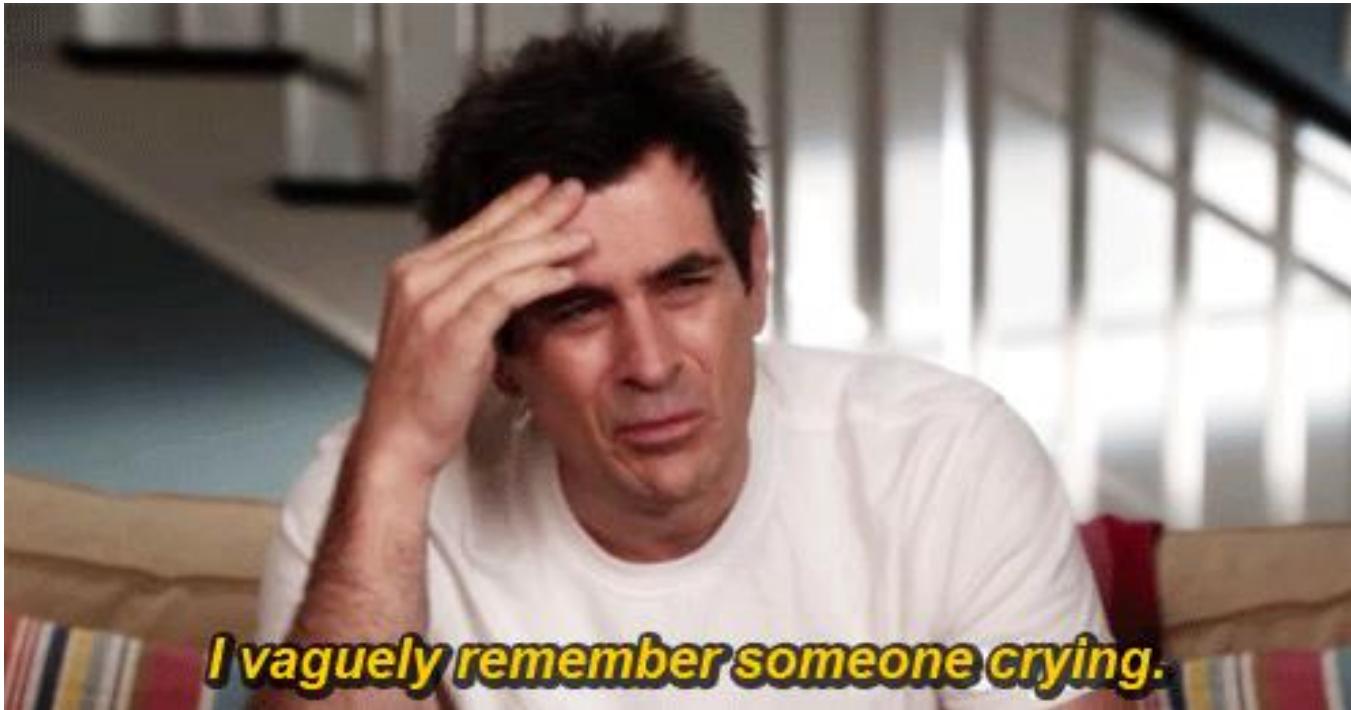
What will you get from this class?

- Theory (morning)
- Workshop (afternoon)



We assume
you've finished
R programming
course already :)

สมัครเรียนพรีໄเด้นท์ DataRockie School
<https://datarockie.teachable.com/>



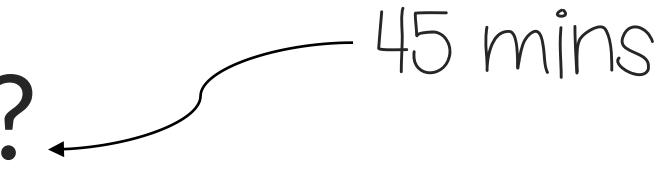
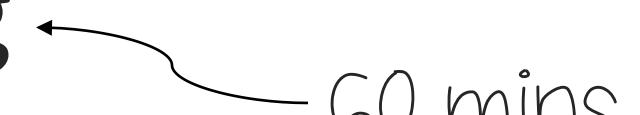
I vaguely remember someone crying.

Don't worry
today we discuss
minimum math &
statistics

**We'll do our best in teaching you
today, but let's face the truth**

- Only practice will make <you> better

Let's get started :)

- What is Data Science?  15 mins
- What is Machine Learning?  45 mins
- 4 Steps in Machine Learning  60 mins
- Algorithms You Need to Know  90 mins

**What is
DATA SCIENCE?**

Data + Science

What is Data?
Data is new electricity.

Science as we know it.
Scientists use data to test hypothesis, theory formation.



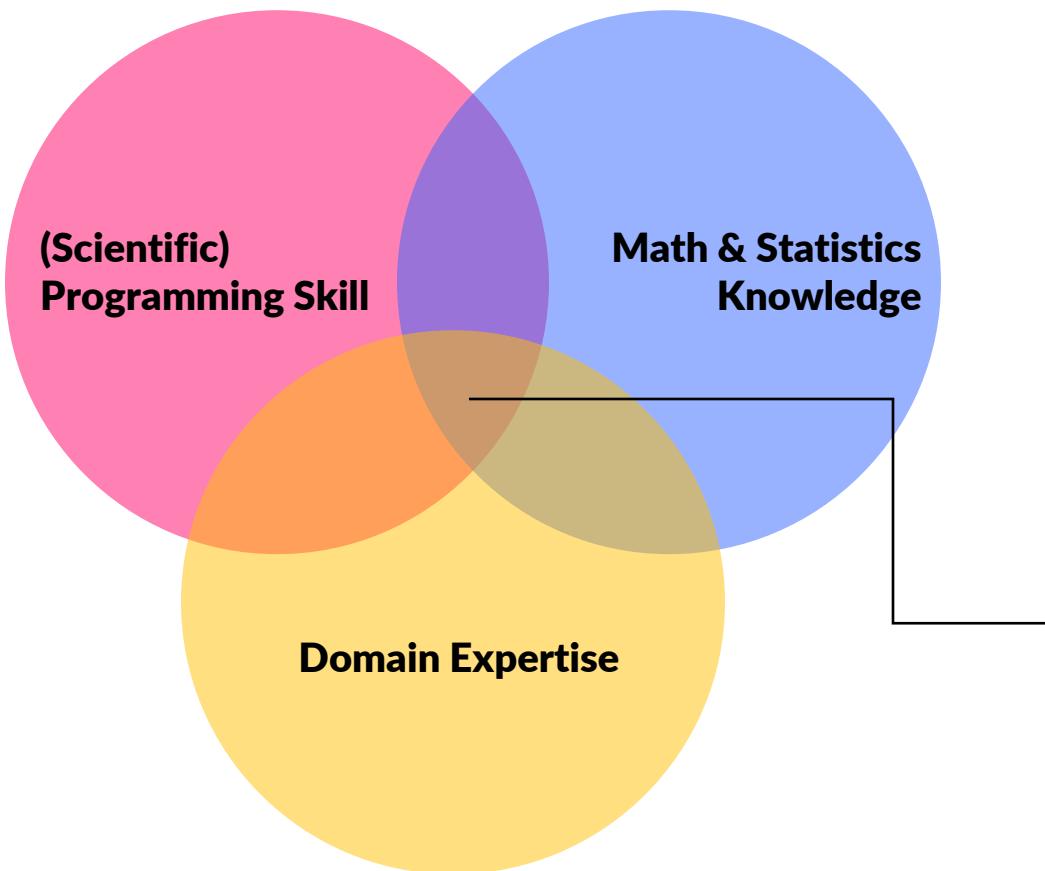
Andrew Ng

“

Data is the past, and the past is Data

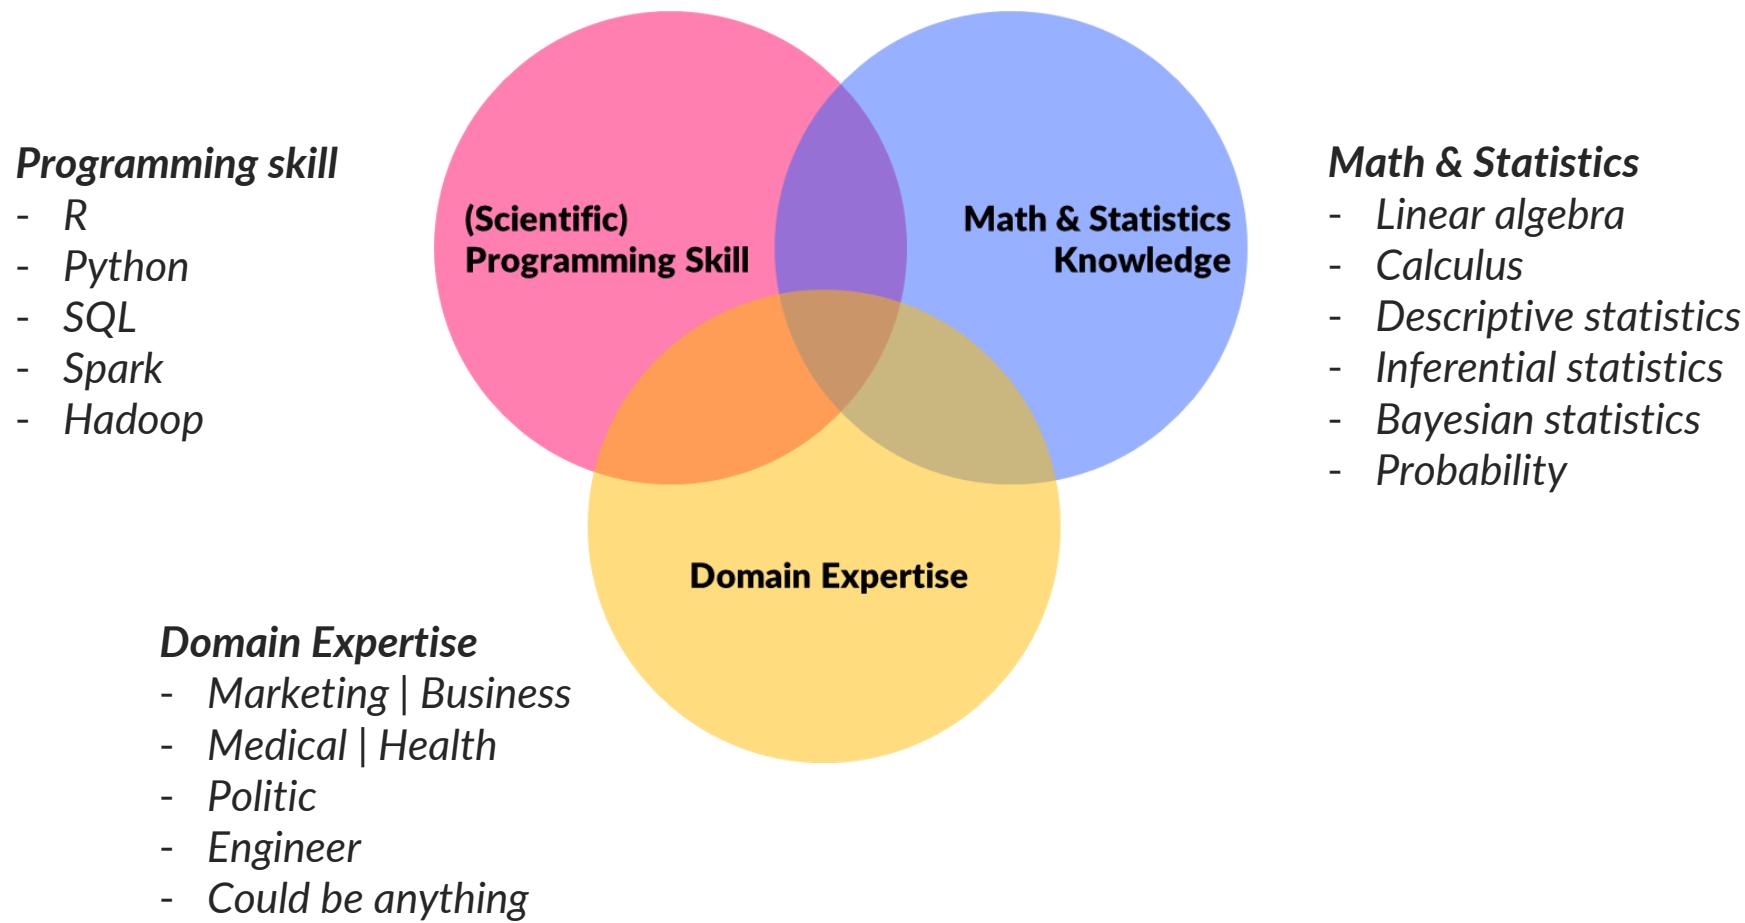
Kirill Eremenko

DS Venn Diagram



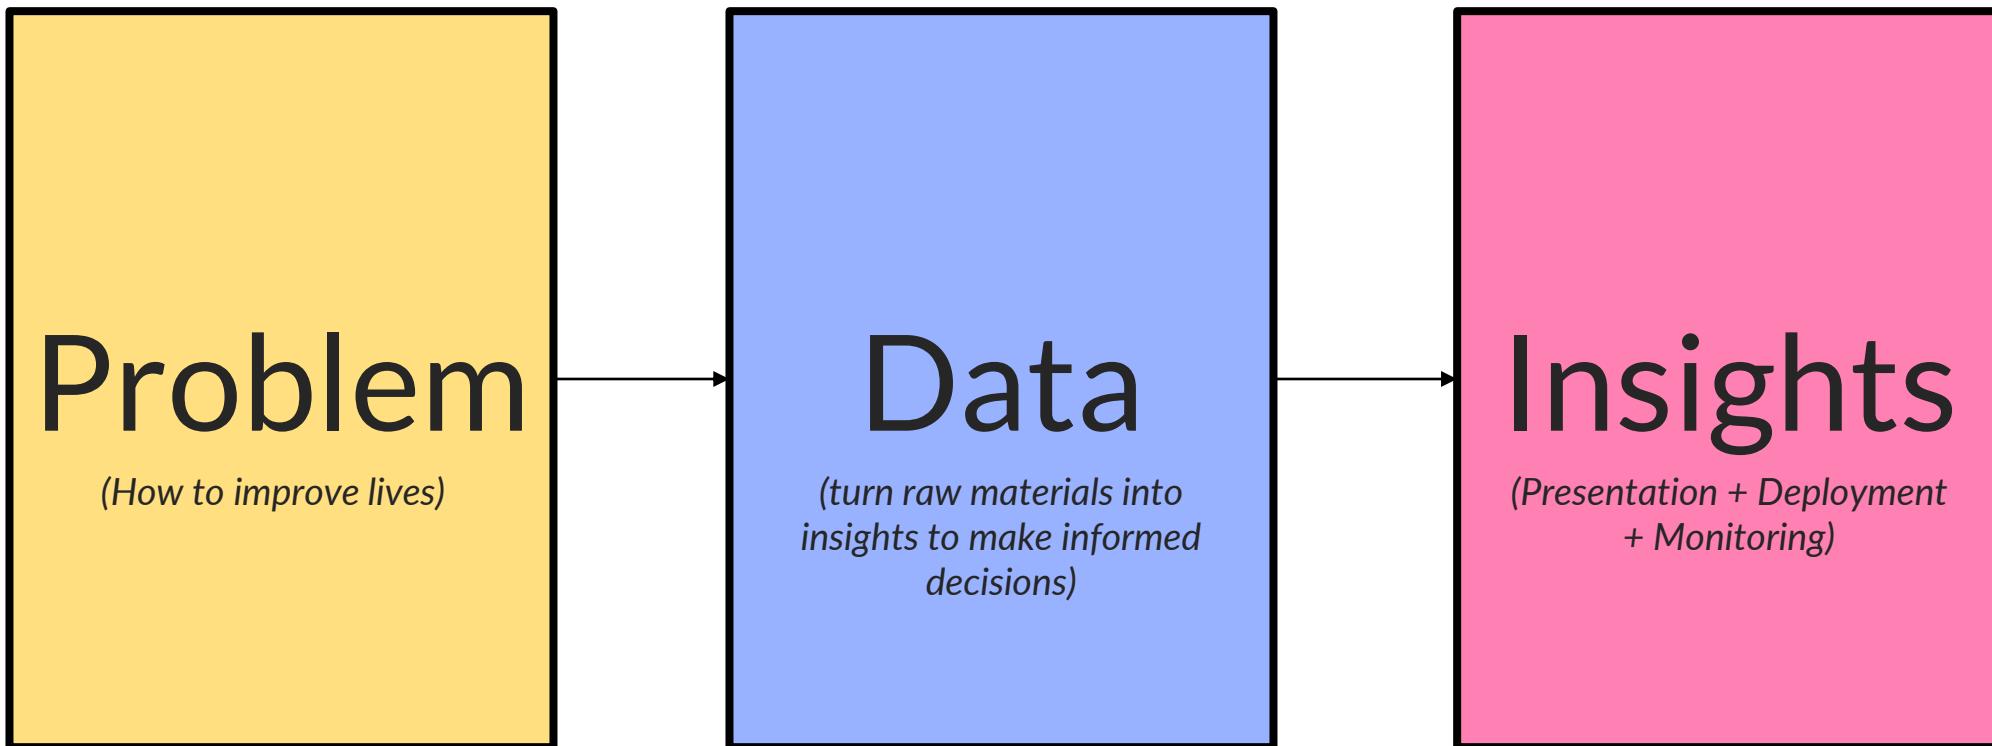
Data Science
is the intersection of three skills

Essential Knowledge



Always

Good data science starts with good question



**What is
MACHINE LEARNING?**

Can you guess?

x (input)	1	2	3	4	5	6	7	8	9	10
y (output)	1	4	9	16	25	36	49	64	81	<input type="text"/>

You guess the function is x^2

A young woman with long brown hair, wearing a yellow beanie and glasses, is smiling while pouring beer from a tap into a glass. She is wearing a grey ribbed sweater. The background shows a blurred bar or pub interior.

Ask her for a date?

Right! Can you guess this one?

	X1	X2	X3	X4	Y
Occasion 1	Restaurant	Many	Friday	Good	Yes
Occasion 2	Pub Bar	Many	Saturday	OK	Yes
Occasion 3	Pub Bar	Few	Monday	OK	No
Occasion 4	Restaurant	Few	Friday	Bad	Yes
Occasion 5	Restaurant	Many	Friday	Good	No
Occasion 6	Pub Bar	Few	Saturday	OK	Yes
Occasion 7	Pub Bar	Many	Monday	Good	No
Occasion 8	Restaurant	Few	Saturday	OK	No
Occasion 9	Pub Bar	Many	Friday	Bad	No
Occasion 10	Restaurant	Few	Friday	Good	

Yes or No?

Machine Learning
is *function approximation*

Machine Learning is *inductive reasoning*

Learning from
experience



ML algorithms trying to MAP inputs (x's) to output (y)

	X1	X2	X3	X4	Y
Occasion 1	Restaurant	Many	Friday	Good	Yes
Occasion 2	Pub Bar	Many	Saturday	OK	Yes
Occasion 3	Pub Bar	Few	Monday	OK	
Occasion 4	Restaurant	Few	Friday	Bad	
Occasion 5	Restaurant	Many	Friday	Good	
Occasion 6	Pub Bar	Few	Saturday	OK	
Occasion 7	Pub Bar	Many	Monday	Good	
Occasion 8	Restaurant	Few	Saturday	OK	
Occasion 9	Pub Bar	Many	Friday	Bad	
Occasion 10	Restaurant	Few	Friday	Good	



It's good to have lots data but ...

REPRESENTATION

is the quality we desire

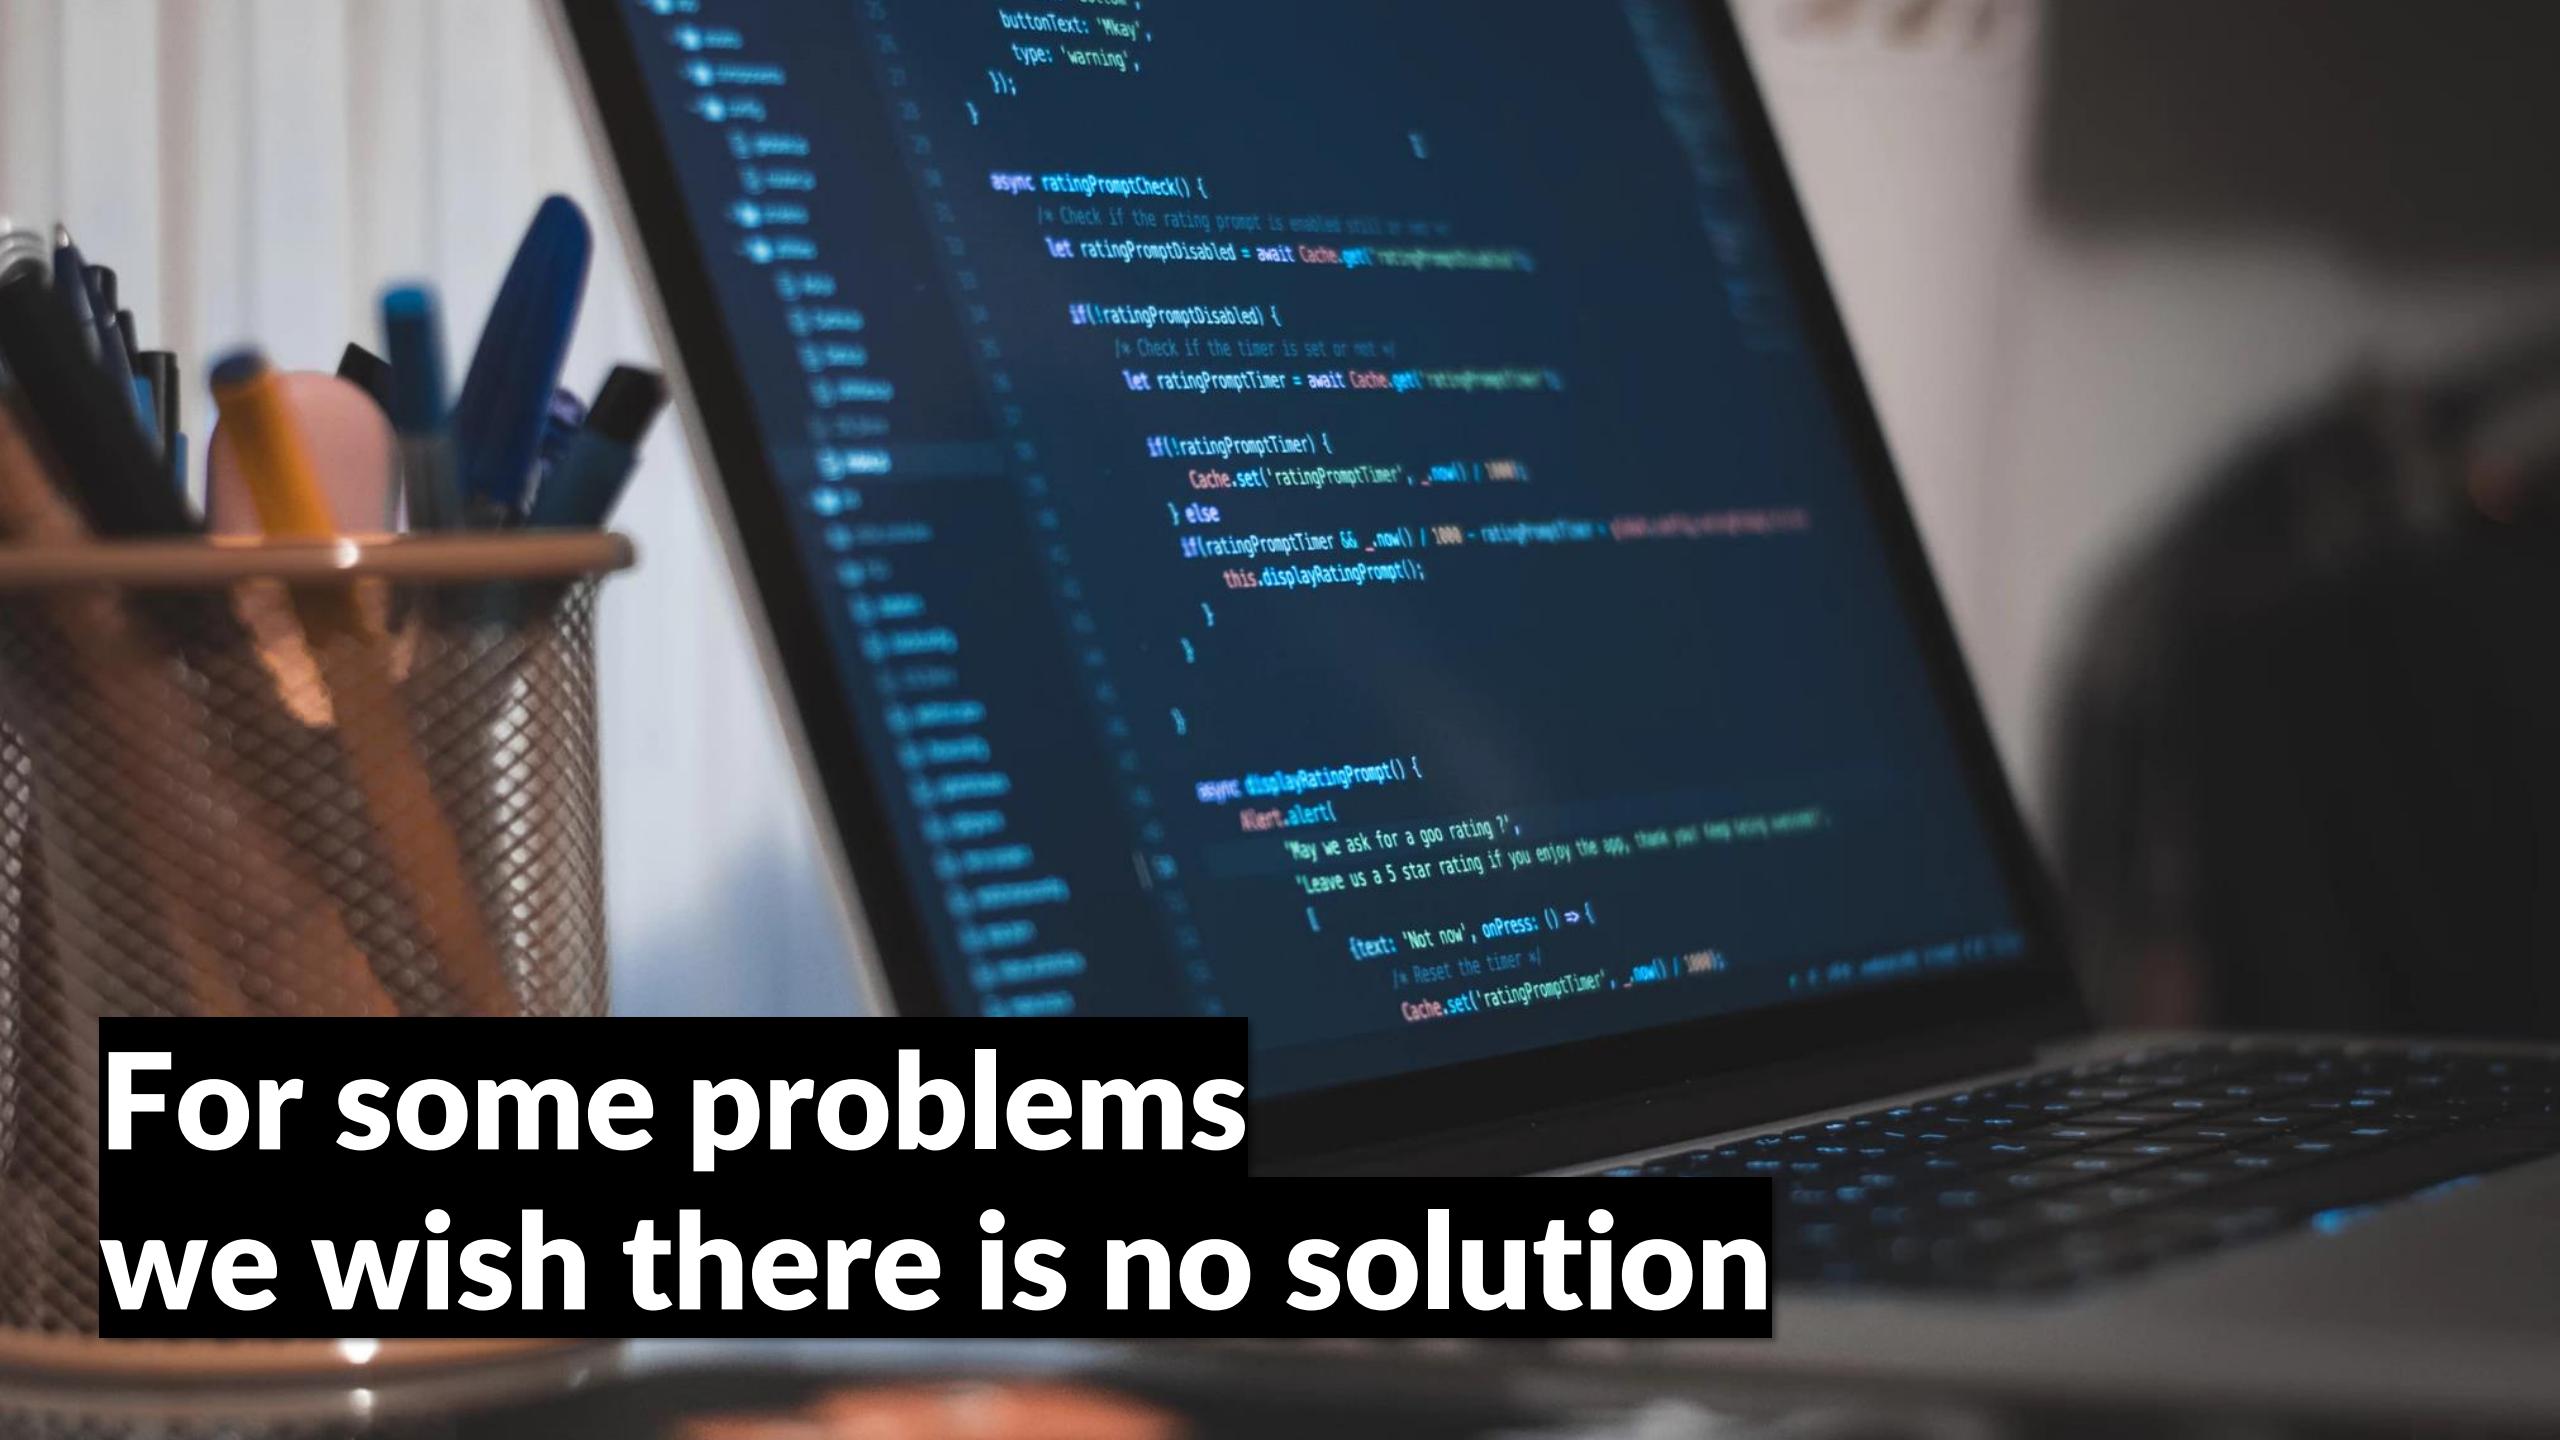
ML Example

- Predict breast cancer
- Weather forecast
- Watson analytics
- Google flu trend
- Obama winning the US election

 *Predicting crime before it happens*



Can you get rid
of poverty?

A close-up photograph of a laptop screen displaying a block of code in a dark-themed code editor. The code is written in a programming language, likely JavaScript or TypeScript, and includes functions for rating prompts and alert dialogs. In the foreground, on the left, there is a blurred image of a pen holder containing several pens and pencils.

**For some problems
we wish there is no solution**

```
        buttonText: 'Mkay',
        type: 'warning',
      });

    }

    async ratingPromptCheck() {
      /* Check if the rating prompt is enabled still or not */
      let ratingPromptDisabled = await Cache.get('ratingPromptDisabled');

      if(!ratingPromptDisabled) {
        /* Check if the timer is set or not */
        let ratingPromptTimer = await Cache.get('ratingPromptTimer');

        if(!ratingPromptTimer) {
          Cache.set('ratingPromptTimer', '_now') / 1000;
        } else {
          if(ratingPromptTimer < _now / 1000 - ratingPromptTimeOut) {
            this.displayRatingPrompt();
          }
        }
      }

      await displayRatingPrompt();
      Alert.alert(
        'May we ask for a goo rating ?',
        'Leave us a 5 star rating if you enjoy the app, thank you very much !',
        [
          {text: 'Not now', onPress: () => {
            /* Reset the timer */
            Cache.set('ratingPromptTimer', '_now') / 1000;
          }}
        ]
      );
    }
  }
}
```

A proper definition of ML

“

Machine Learning: Field of study that gives computers *the ability to learn* without being explicitly programmed.



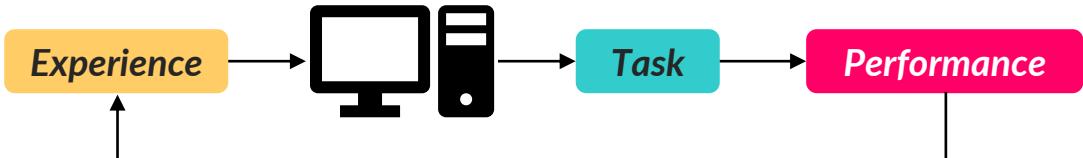
Arthur Samuel (1959)

A more modern def in 1998

Well posed Learning Problem: A computer program is said to learn from experience **E** with respect to some task **T** and some performance measure **P**, if its performance on **T**, as measured by **P**, improves with experience **E**.



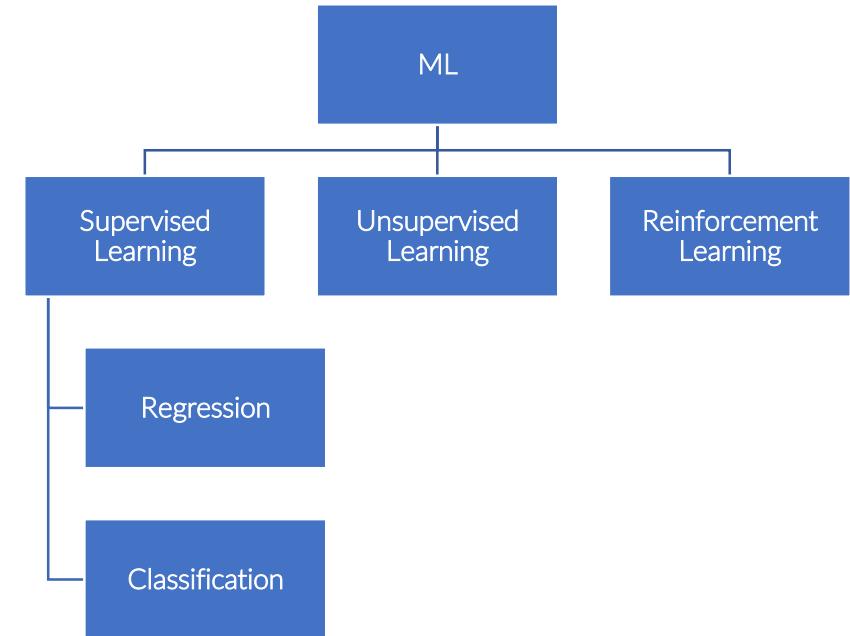
Tom Mitchell (1998)



Can we improve computer's ability to perform task **T** over time (without being explicitly programmed)?

Types of ML algorithms

1. Supervised Learning
 - i. Regression
 - ii. Classification
2. Unsupervised Learning
3. Reinforcement Learning



What's Artificial Intelligence



AI

ML

DEEP LEARNING

Function approximation

$$Y = f(X)$$

Model = Algorithm (Data)

Predicting new unseen data

In the past

We build models to represent the world (problem we're trying to solve)



And the goal of Machine Learning is **GENERALIZATION**

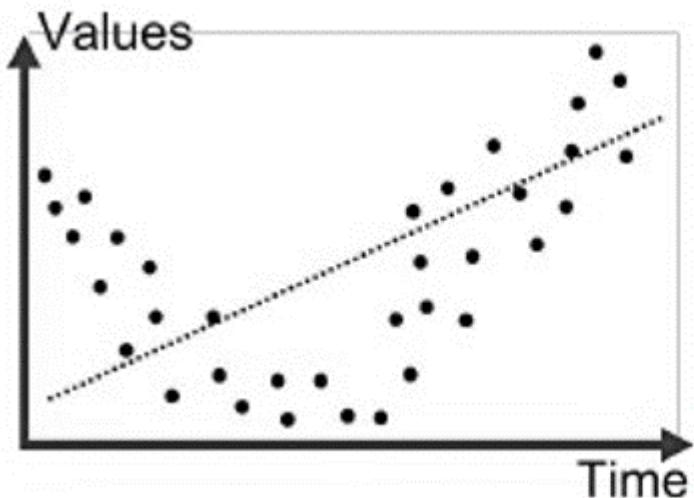
Beyond (past) data you used to build the model

**Generalization means we make sure
our models **don't overfit** the data**

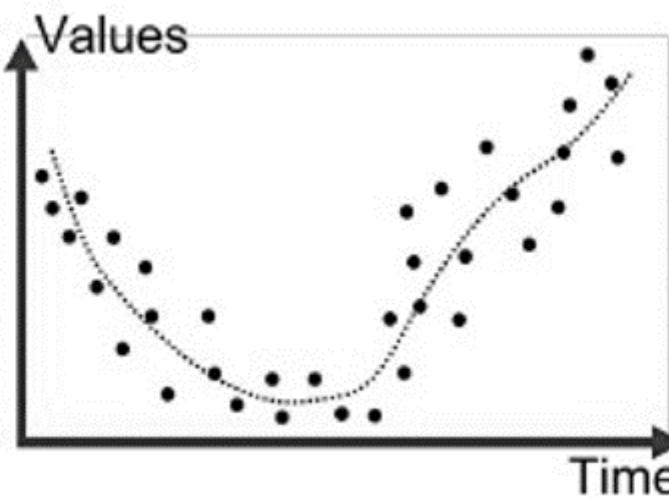
Also don't underfit

We mean
data in the
past

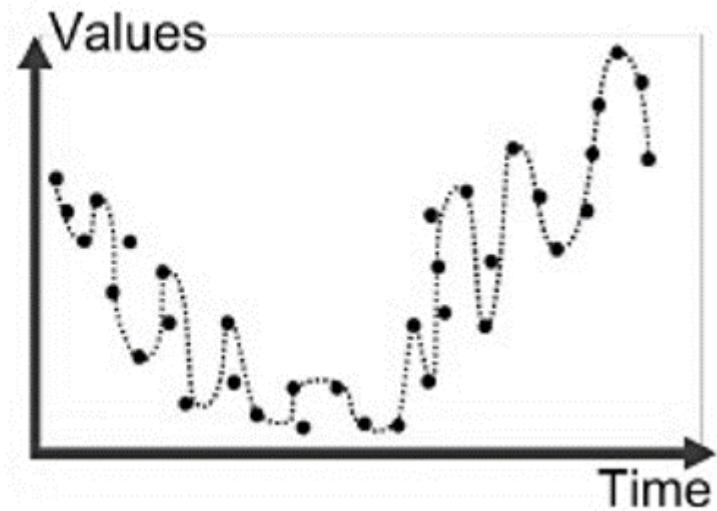
If your model ***can't generalize to new unseen data***, then you're in trouble.



Underfitting



Good fit



Overfitting

https://cdn-images-1.medium.com/max/1600/1*6vPGzBNppqMHllg1o_se8Q.png

This model looks okay as
it fit underlying patterns
in the data well. Though,
not 100% correct ?

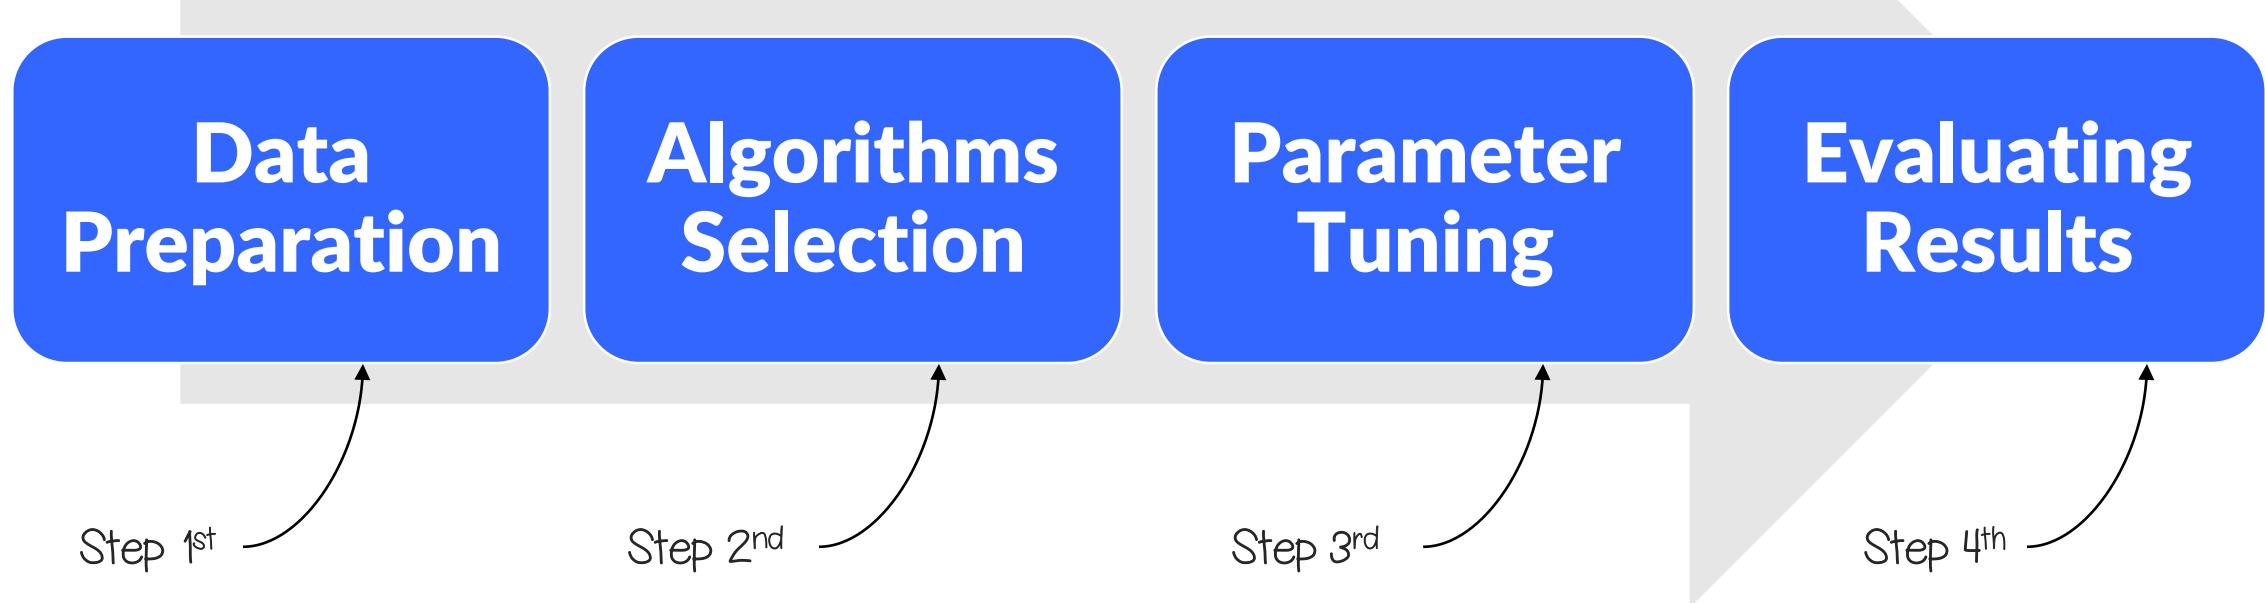
SUMMARY

- Data Science in Thailand is still very young
- ML is data scientist's arsenal tool
- ML is **Function Approximation**, learning from experience
- **Representation & Generalization** (avoid overfitting | underfitting)
- 3 types of ML algorithms including supervised, unsupervised and reinforcement learning
- Regression predicts numbers | Classification predicts categories



KICK START MACHINE LEARNING

Four key steps in ML



DATA PREPARATION

Data Format

The diagram illustrates the structure of a dataset. A horizontal double-headed arrow at the top is labeled "Variables | Features". Below it is a table with 10 rows, each representing a data point. The first column is labeled "Occasion 1" through "Occasion 10". The columns are labeled X1, X2, X3, X4, and Y. The last column, Y, is labeled "Target Variable". A vertical double-headed arrow on the right side of the table is labeled "Data Points | Observations".

	X1	X2	X3	X4	Y
Occasion 1	Restaurant	Many	Friday	Good	Yes
Occasion 2	Pub Bar	Many	Saturday	OK	Yes
Occasion 3	Pub Bar	Few	Monday	OK	No
Occasion 4	Restaurant	Few	Friday	Bad	Yes
Occasion 5	Restaurant	Many	Friday	Good	No
Occasion 6	Pub Bar	Few	Saturday	OK	Yes
Occasion 7	Pub Bar	Many	Monday	Good	No
Occasion 8	Restaurant	Few	Saturday	OK	No
Occasion 9	Pub Bar	Many	Friday	Bad	No
Occasion 10	Restaurant	Few	Friday	Good	No

Variable Types

1. Binary (1/0)
2. Categorical
3. Integer
4. Continuous

Feature Selection

- Choose relevant set of features that represent the problem
- Curse of dimensionality: 2^n

Feature Engineering

- Recode
- Combine multiple variables (PCA)
- Extract new information from existing feature

Gentle intro to Feature Engineering

ID	Name	DOB	Math	Stats	Art	Music	Design
1	Mr. David Beckham	13 / 05 / 1975	85	90	50	42	49
2	Ms. Angelina Jolie	20 / 09 / 1988	50	60	90	95	88
3	Mr. Khalid Khan	05 / 01 / 1950	32	65	70	72	65
4	Mr. Hibino Takamoto	31 / 12 / 1995	75	76	80	50	30

Question:

What new variable can we create from this dataset?

Think at least **5 new variables**

We can use
`dplyr::mutate()` to
create new variables

Missing Data

1. Approximated
2. Predicted (Supervised Learning)
3. Removed

Mean | Median
imputation

If we have large data points and NA is not
much, we may consider removing them

Approximated or Removed?

```
Console ~/Desktop/R PROGRAMMING/ ↵
> head(complete.cases(biopsy), 200)
[1] TRUE  TRUE
[14] TRUE  FALSE  TRUE  TRUE
[27] TRUE  TRUE
[40] TRUE  FALSE  TRUE  TRUE
[53] TRUE  TRUE
[66] TRUE  TRUE
[79] TRUE  TRUE
[92] TRUE  TRUE
[105] TRUE  TRUE
[118] TRUE  TRUE
[131] TRUE  FALSE  TRUE  TRUE
[144] TRUE  TRUE  FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
[157] TRUE  TRUE  FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  FALSE  TRUE  TRUE  TRUE  TRUE
[170] TRUE  TRUE
[183] TRUE  TRUE
[196] TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

R functions:

`complete.cases(df)` สำหรับตรวจสอบว่าข้อมูลมี NA หรือเปล่า?

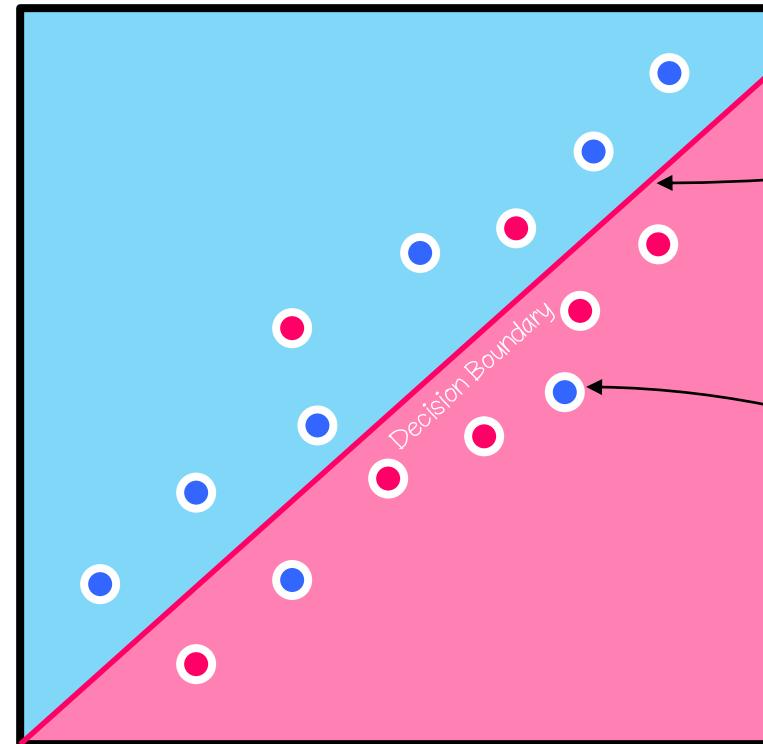
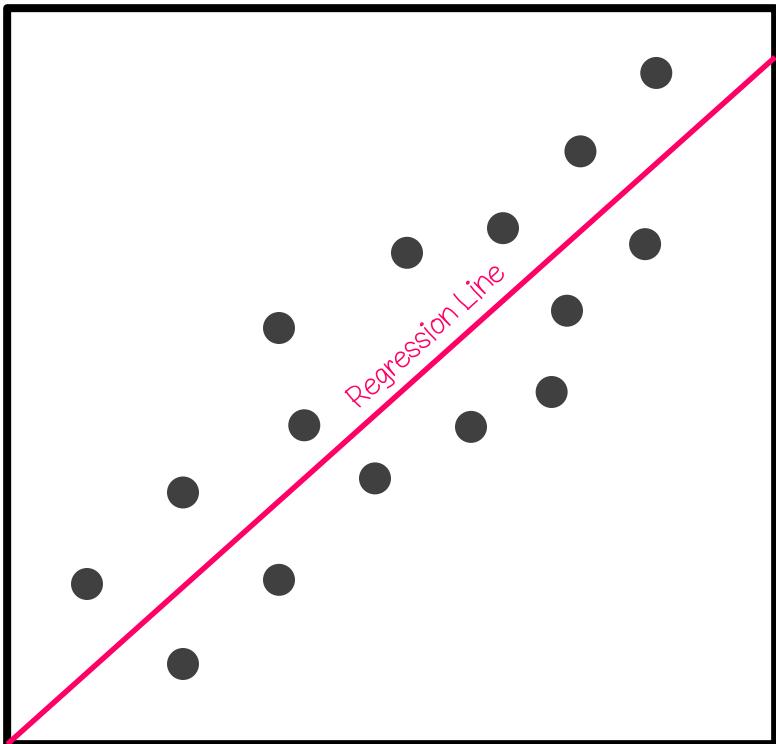
`tidyr::drop_na(df)` สำหรับ drop case ที่มีข้อมูลไม่ครบถ้วน

ALGORITHM SELECTION

There are 3 types of algorithms

	Algorithms	Main Tasks
Unsupervised learning	K-Means Principal Component Analysis (PCA) Association Rules Social Network Analysis (SNA)	Tell me what patterns exist in my data - descriptive
Supervised learning	Linear Regression Logistic Regression Decision Tree Random Forests KNN Support Vector Machine Neural Networks	Use the pattern in my data to make predictions - predictive
Reinforcement learning	Multi-Armed Bandits (UBC, Thompson Sampling)	Use the pattern in my data to make predictions and - improve - these predictions as more results come in

Regression vs. Classification



Decision Boundary:
The line that best split
red vs. blue data points

100% in reality may not
be feasible, some blue
points are mistakenly
classified as red

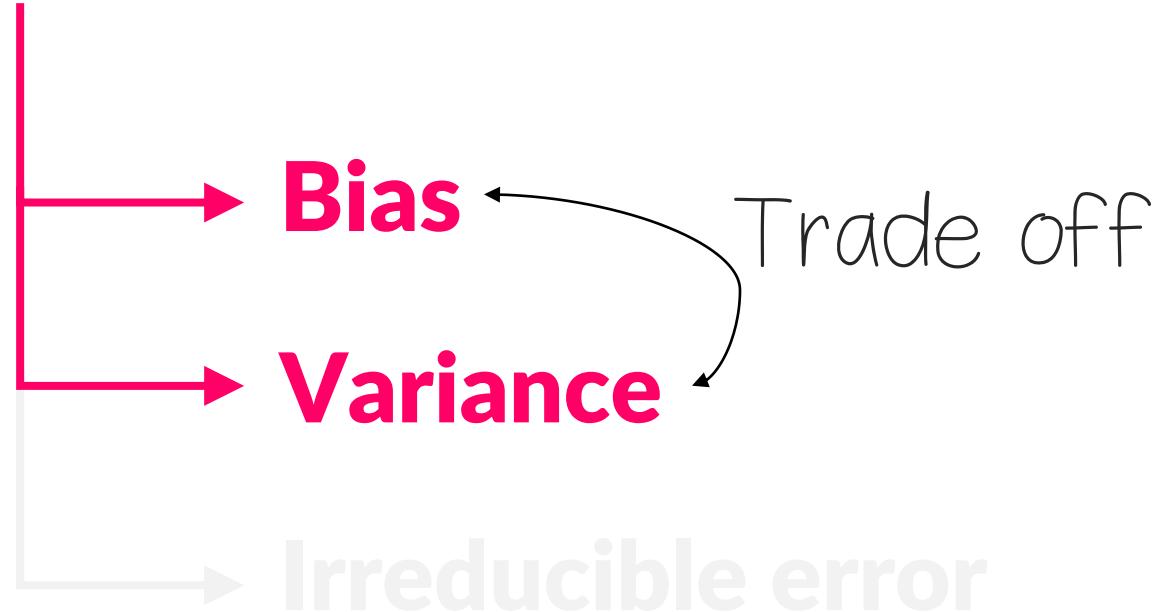
NO FREE LUNCH

- Speed
- Accuracy
- Bias & Variance Tradeoff



Our model consists of 3 types of error

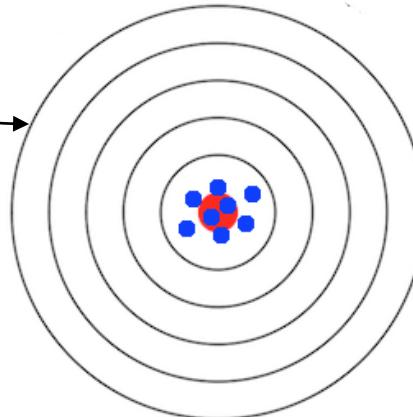
$$Y = f(X) + \text{Error}$$



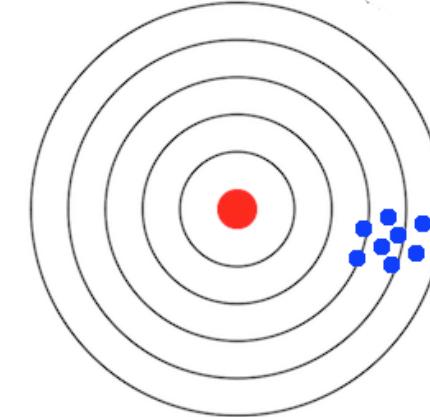
This is ideal,
but difficult
in practice

Low bias

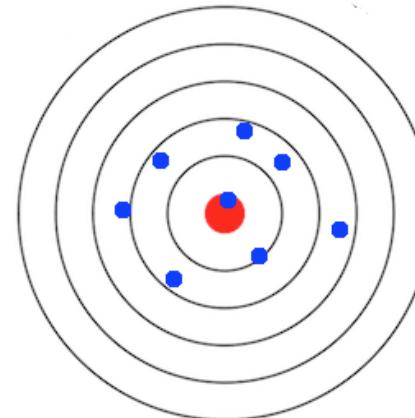
Low
variance



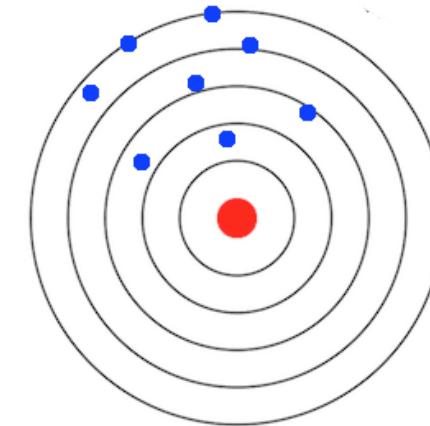
High bias



High
variance



This must
be avoided



Example Models

High Bias

Having lots of assumptions to do function approximation

Linear Regression
Logistic Regression
Linear Discriminant Analysis

Low Bias

Having little or few assumptions about the function approximation

Decision Tree
KNN
SVM

Low Variance

New data does not cause model to change much

High Variance

New data causes model to change much i.e. prediction results vary significantly compared to low variance algorithms

PARAMETER TUNING



Music Composer is
a very good example

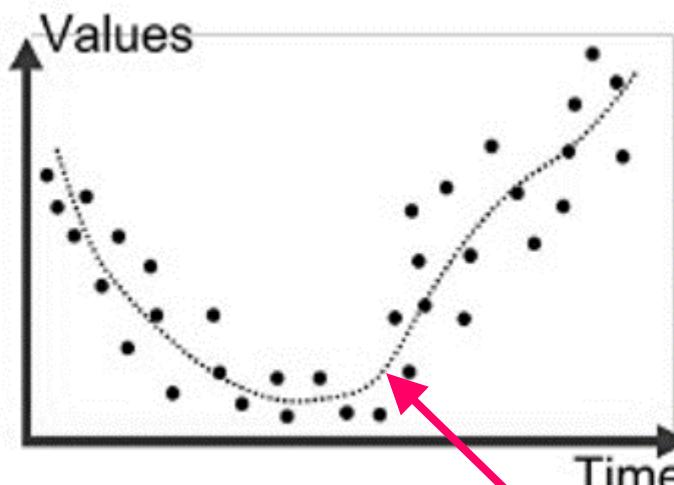
There are so many learners.

**Even the same learner can have
different performances.**

We tune parameter to get the best results (**GOOD FIT**) from our learner.



Underfitting



Good fit



Overfitting

The parameters that can be tuned
called "**HYPERPARAMETERS**"

Some models have hyperparameters, some don't.

Stochastic Gradient Boosting

```
method = 'gbm'
```

Type: Regression, Classification

Tuning parameters:

- `n.trees` (# Boosting Iterations)
- `interaction.depth` (Max Tree Depth)
- `shrinkage` (Shrinkage)
- `n.minobsinnode` (Min. Terminal Node Size)



Hyperparameters of GBM algorithms

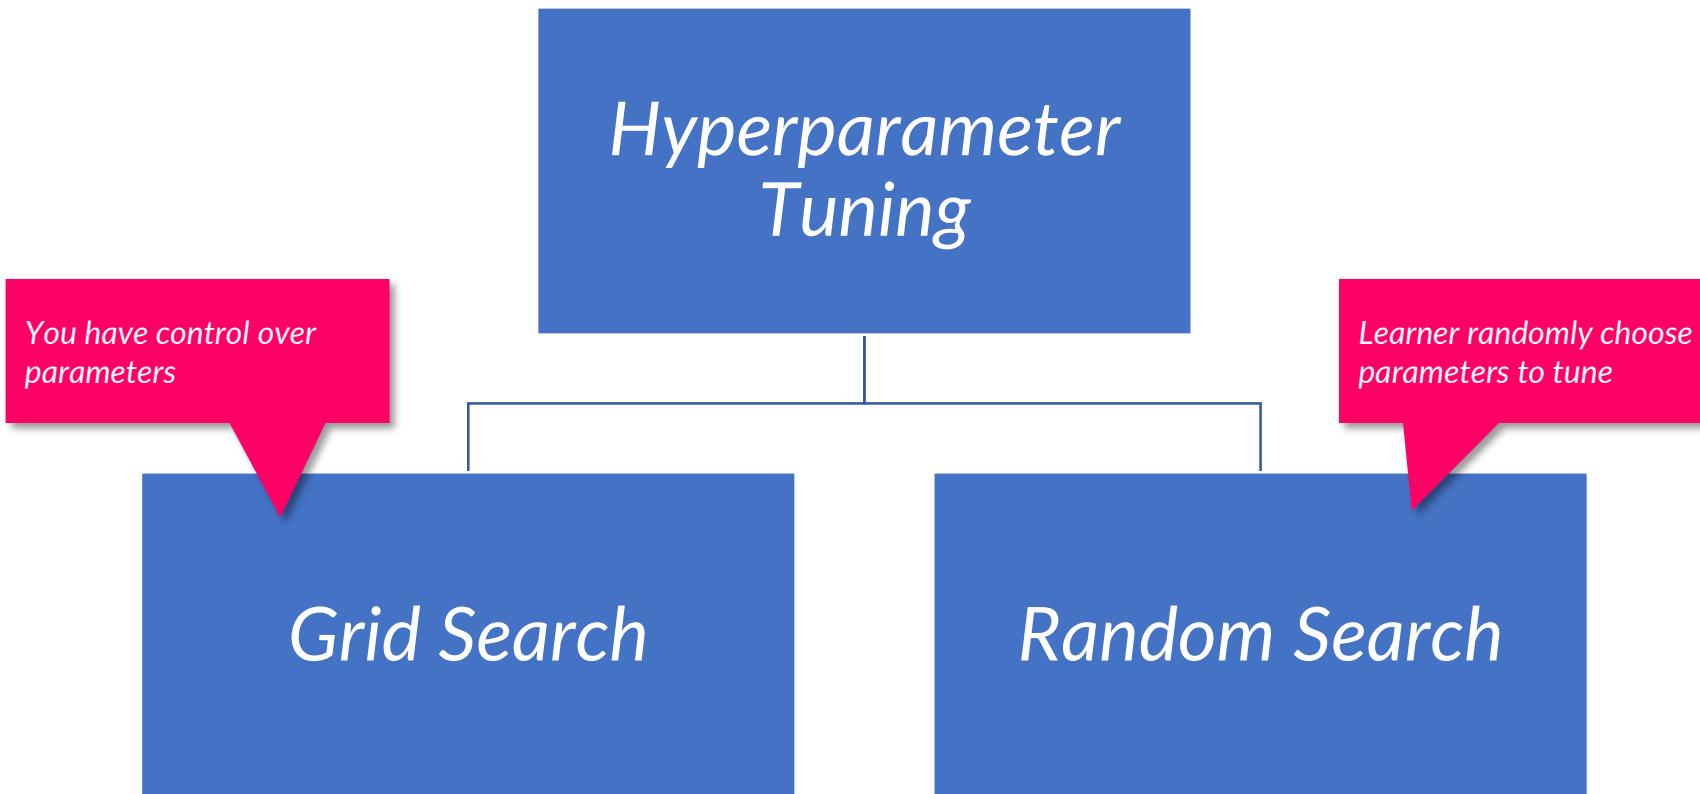
Required packages: `gbm` , `plyr`

A model-specific variable importance metric is available.

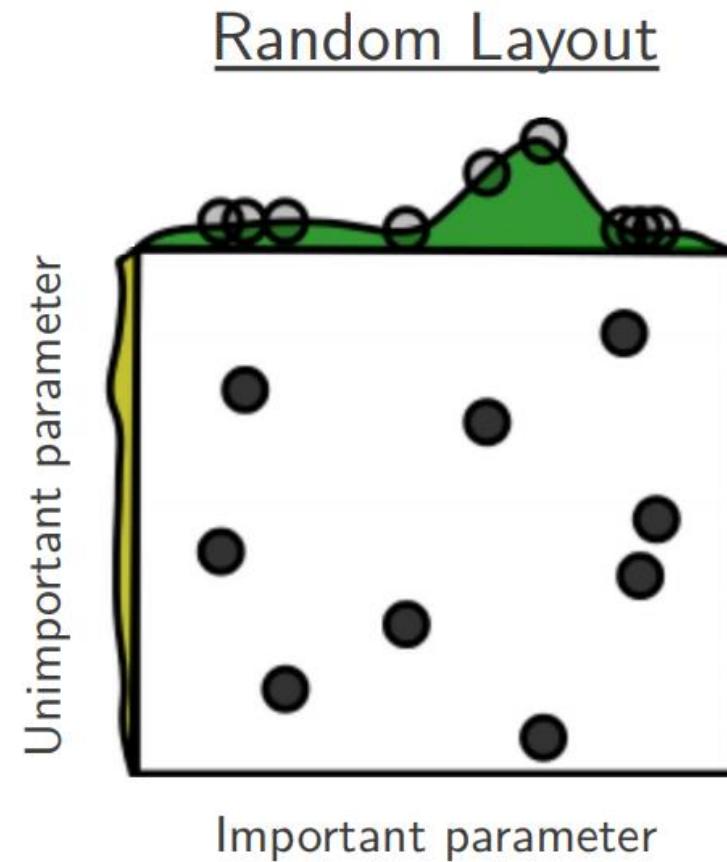
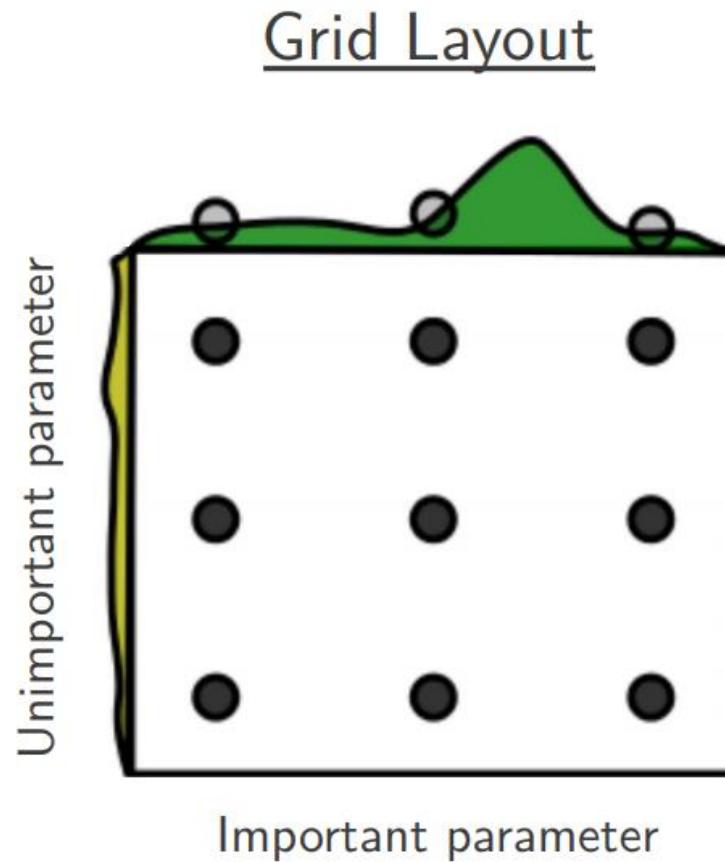
<http://topepo.github.io/caret/train-models-by-tag.html>

SEARCH for the best tune

ML is optimization problem



It really depends on the learners. Sometimes it's more efficient to use random search, sometimes it's not.



EVALUATING RESULTS

Performance Metrics

Classification problems

- *Confusion Matrix*
- *% Accuracy*
- *Area Under Curve (AUC)*
- *Logarithmic Loss (Log Loss)*
- *Cohen's KAPPA #for imbalanced problems*

Regression problems

- *Mean Absolute Error (MAE)*
- *Root Mean Squared Error (RMSE)*

Cancer Detection



Confusion Matrix

		Reference	
		Cancer	No Cancer
Predicted	Cancer	60	3
	No Cancer	2	35

Accuracy is the sum of diagonal divided by total sample size

$$60+35 / 100 = 95\%$$

Confusion Matrix

		Reference	
		Cancer	No Cancer
Predicted	Cancer	60 (A)	3 (B)
	No Cancer	2 (C)	35 (D)

ถ้าผู้ป่วยเป็นมะเร็ง เราหายฉุกเฉิน (%)

$$\text{Sensitivity} = A / A+C$$

ถ้าผู้ป่วยไม่ได้เป็นมะเร็ง เราหายฉุกเฉิน (%)

$$\text{Specificity} = D / B+D$$

ผล prediction "cancer" ของเราถูกกี่ %

$$\text{Precision} = A / A+B$$

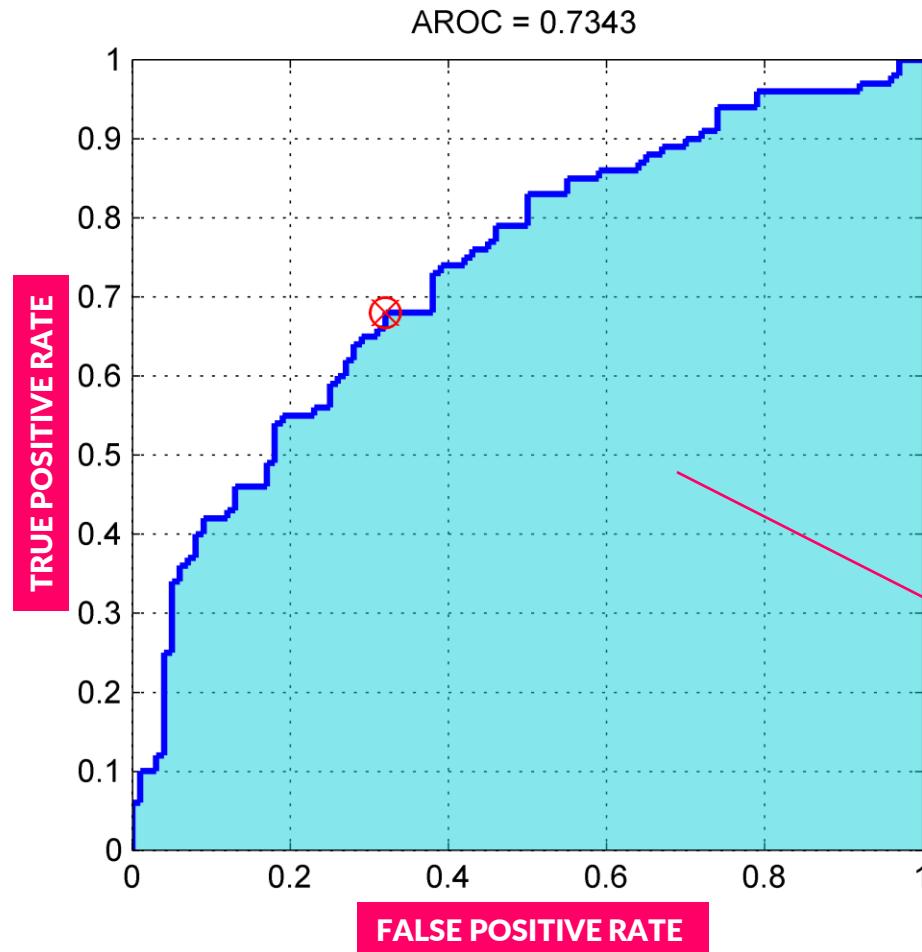
ในผู้ป่วยโรคมะเร็งทั้งหมด เรายำนาຍฉุกเฉิน (%)

$$\text{Recall} = A / A+C$$

Measure of overall test's accuracy

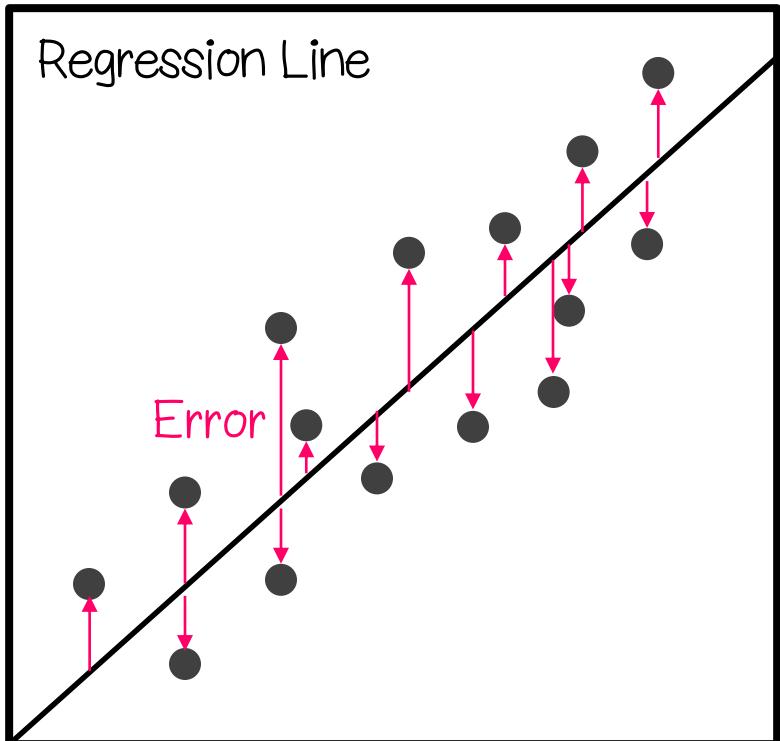
$$F1 = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

Area Under Curve (AUC)



For binary classification, sometimes we want to maximize AUC

Root Mean Squared Error (RMSE)



Mathematical Expression always look scary

$$RMSE = \sqrt{\frac{\sum_{i=1}^n (Actual - Predicted y_i)^2}{N}}$$

$$RMSE = \sqrt{\frac{sum(error^2)}{N}}$$

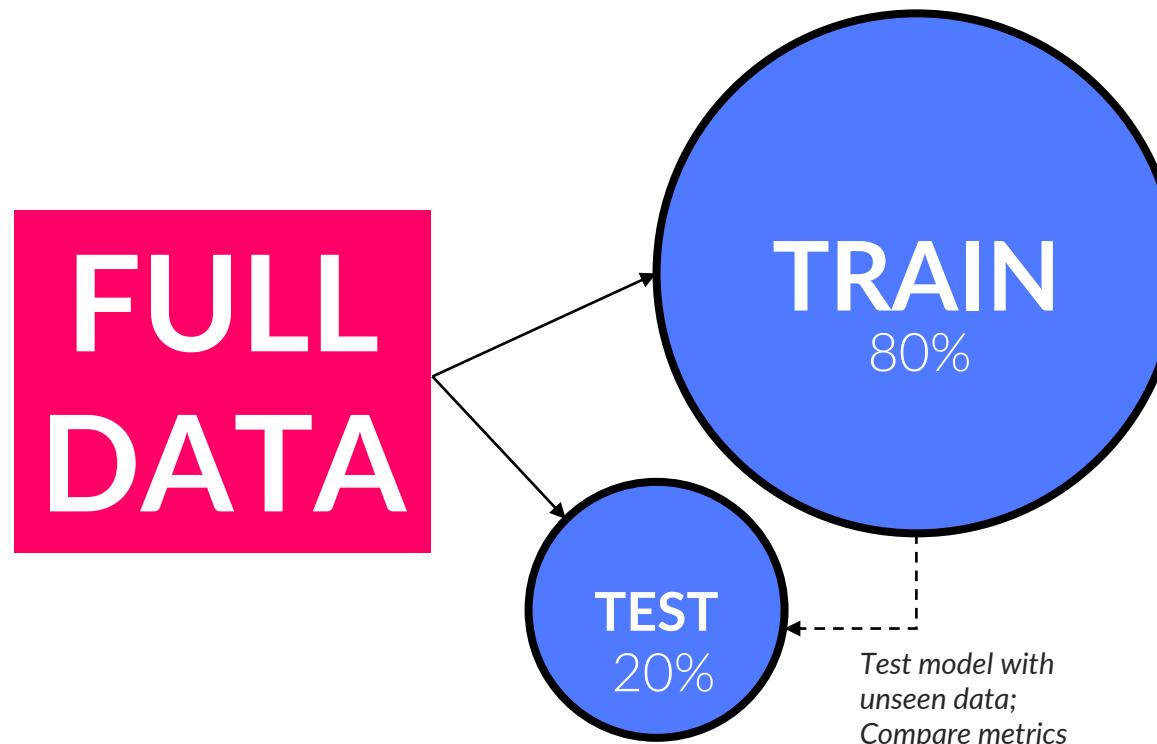
Remember this one

$$RMSE = \sqrt{mean(error^2)}$$

**Recall the goal of ML is
“GENERALIZATION”**

**We must test our model on new
unseen dataset**

Is our model good to go?



Training set (build model)

Testing set (evaluate model)

Decision Criterion:

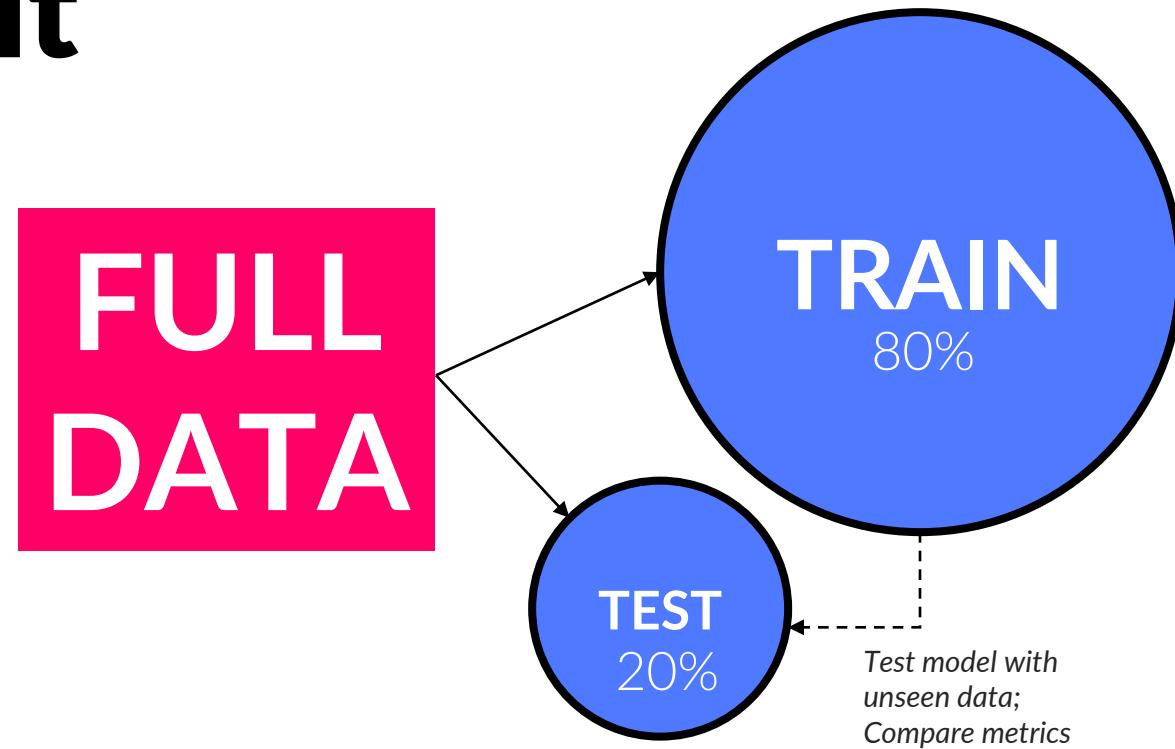
*The model **passes the test** if
tested performance is similar
trained performance [metric]*

Resampling techniques

1. Train-test split
2. *k*-Fold Cross-Validation
3. Repeated *k*-Fold Cross-Validation
4. Bootstrap
5. Leave One Out Cross-Validation

*This method is preferred
but it's very expensive for
large dataset*

The most efficient method is train-test split



K-Fold Cross Validation

FULL DATASET

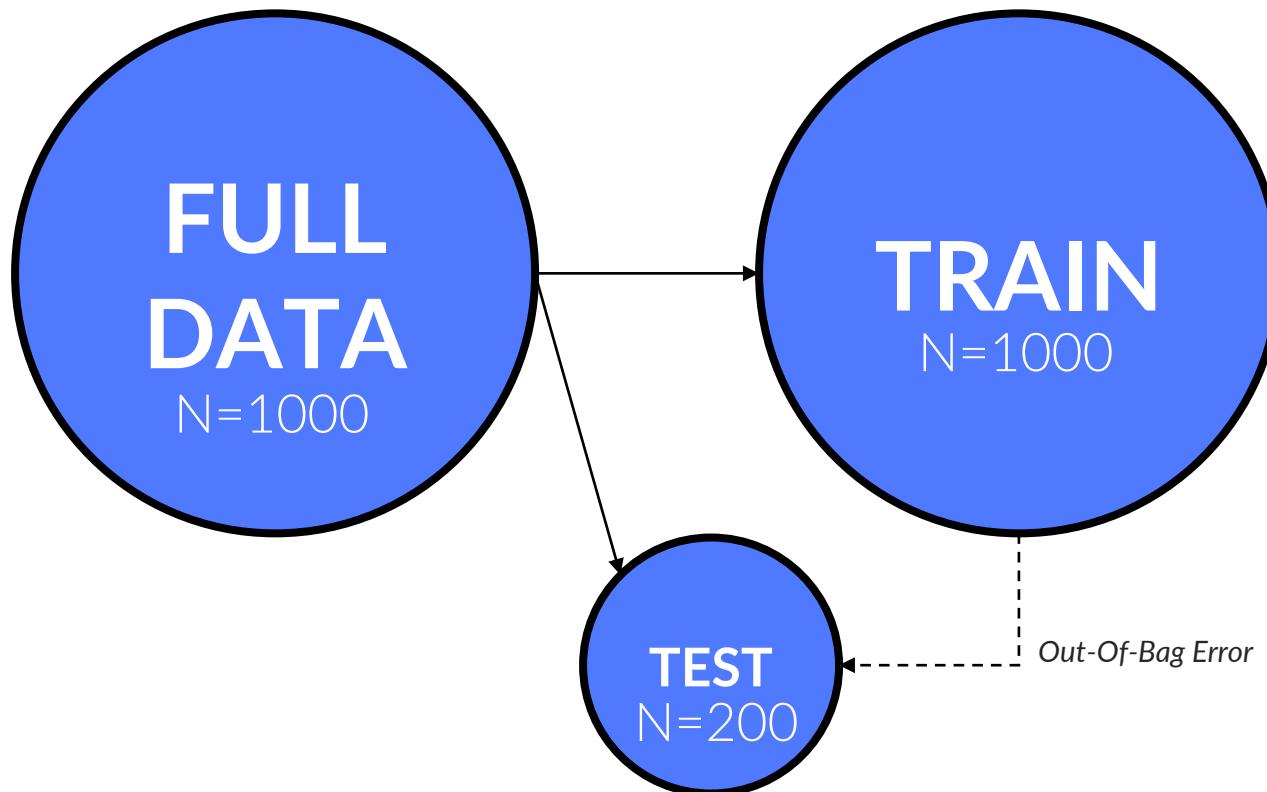


[1]	Train	Train	Train	Train	Test
[2]	Train	Train	Train	Test	Train
[3]	Train	Train	Test	Train	Train
[4]	Train	Test	Train	Train	Train
[5]	Test	Train	Train	Train	Train

People normally
do 5, 10 K-Fold

Bootstrap Sampling

(aka. Random sampling with replacement)

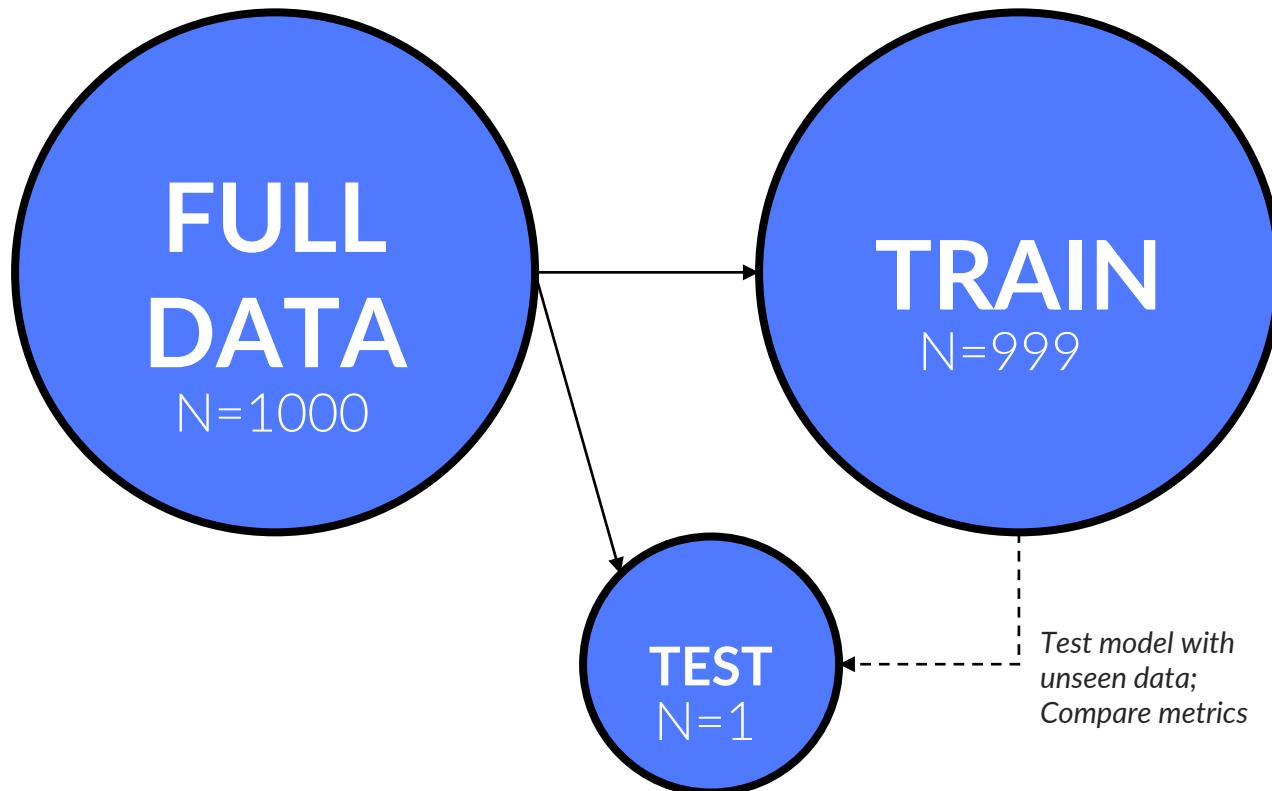


Explanation:

1000 may come from 800 unique data points + 200 repeated data points

Leave One Out CV

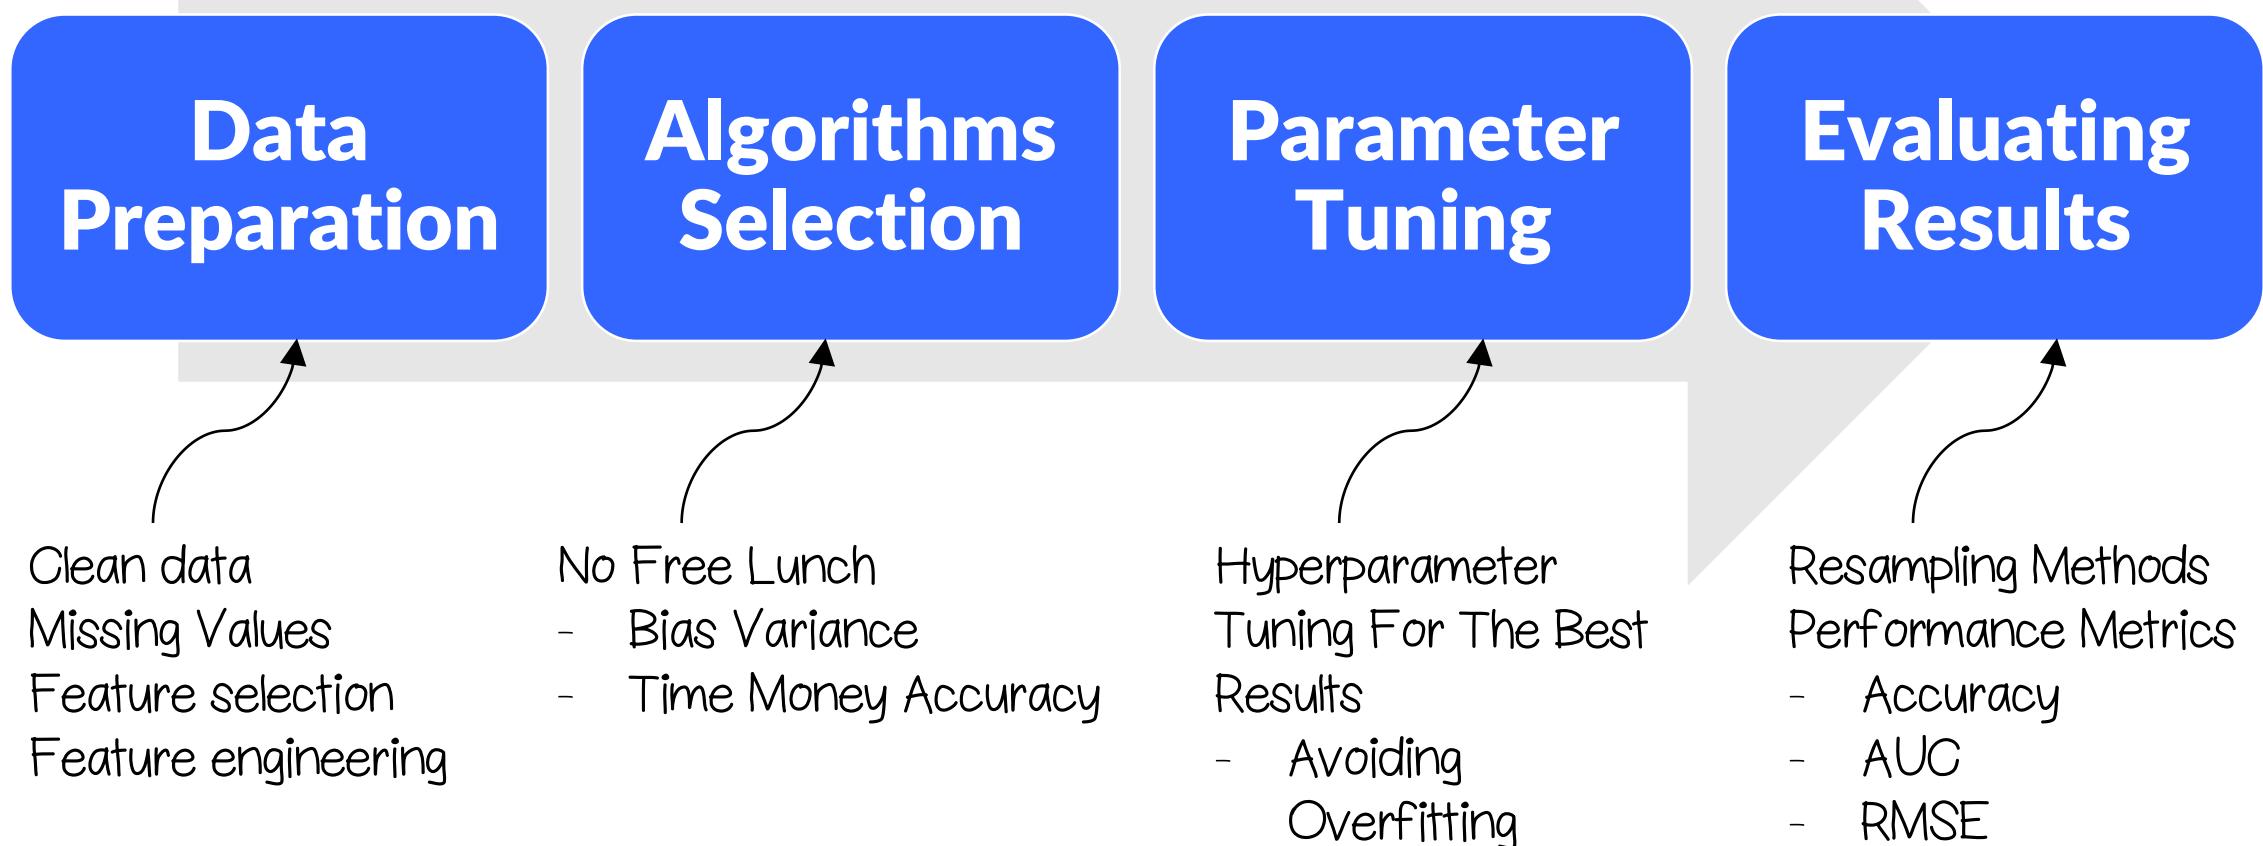
This is the most expensive method – time consuming



Explanation:

Leave 1 data point to test the model. For $n=1000$, this means we need to run model 1000 times

#SIMPLE :)



SUMMARY

- Data preparation takes 70-80% of our time.
 - “Garbage in, Garbage out”
 - Feature selection | Feature engineering | Missing Value
- No free lunch theorem
 - Bias vs. Variance Tradeoff
 - Speed vs. Accuracy Tradeoff
- [Hyper]Parameter tuning to get the best out of our learning algorithms
- Resampling techniques used in ML research incl. train-test, k-fold CV, repeated k-fold CV, bootstrap, and LOOCV

OKAY !! Let's recap a bit

- ML is function approximation
- Learn from past data to predict future (unseen) data point
- The goal is representation & generalization (avoid overfitting)
- Regression or Classification problems depend target variable (y)
- No Free Lunch: bias vs. variance tradeoff
- Algorithms can be tuned via hyperparameter
- K-Fold cross validation is popular resampling technique
- Our model passes the test if training vs. testing performances are similar (not overfit)

ALGORITHMS YOU SHOULD KNOW

We cover 8 algorithms

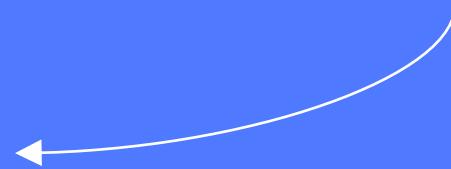
Linear Regression	Regression
Logistic Regression	Classification
Decision Tree	Regression & Classification
Random Forest	(mainly) Classification
KNN	(mainly) Classification
Support Vector Machine	(mainly) Classification
Neural Networks	Regression & Classification
K-Means	Unsupervised Learning

Supervised Learning

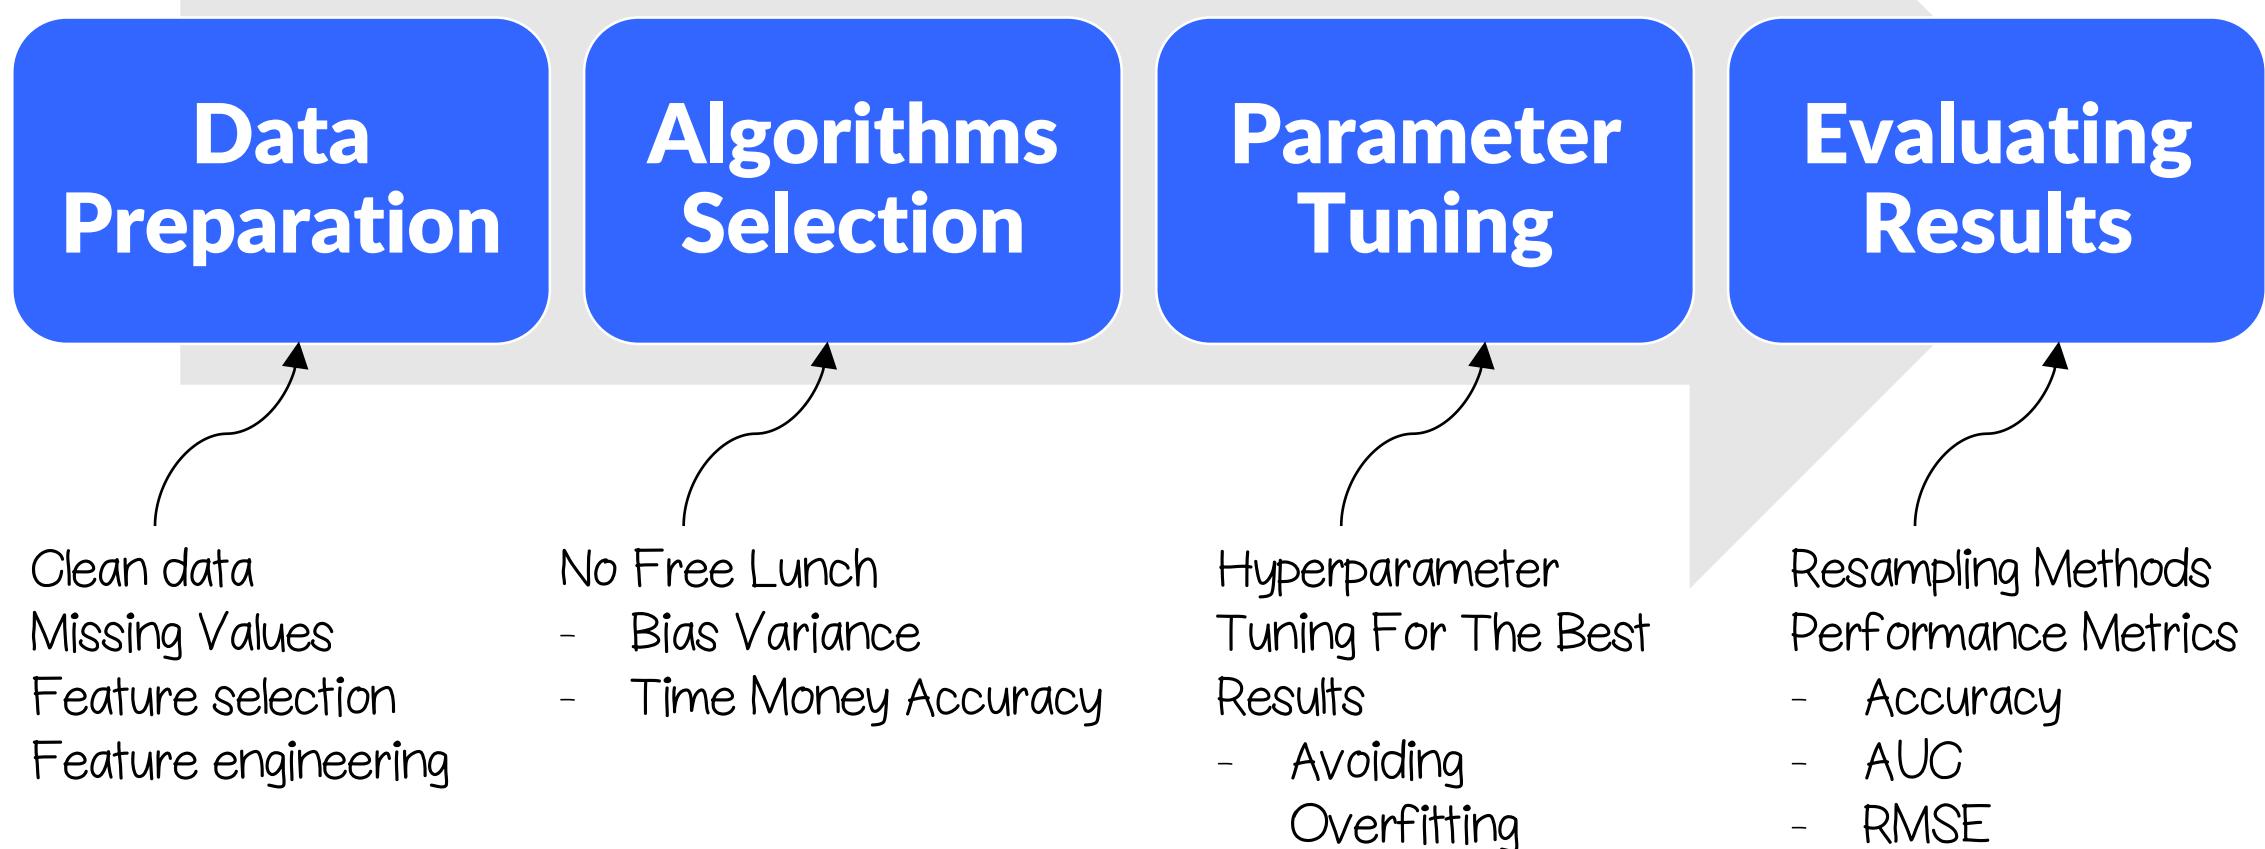
NN is special algorithm, it works like reinforcement learning

LINEAR REGRESSION

Hello World in
Machine Learning

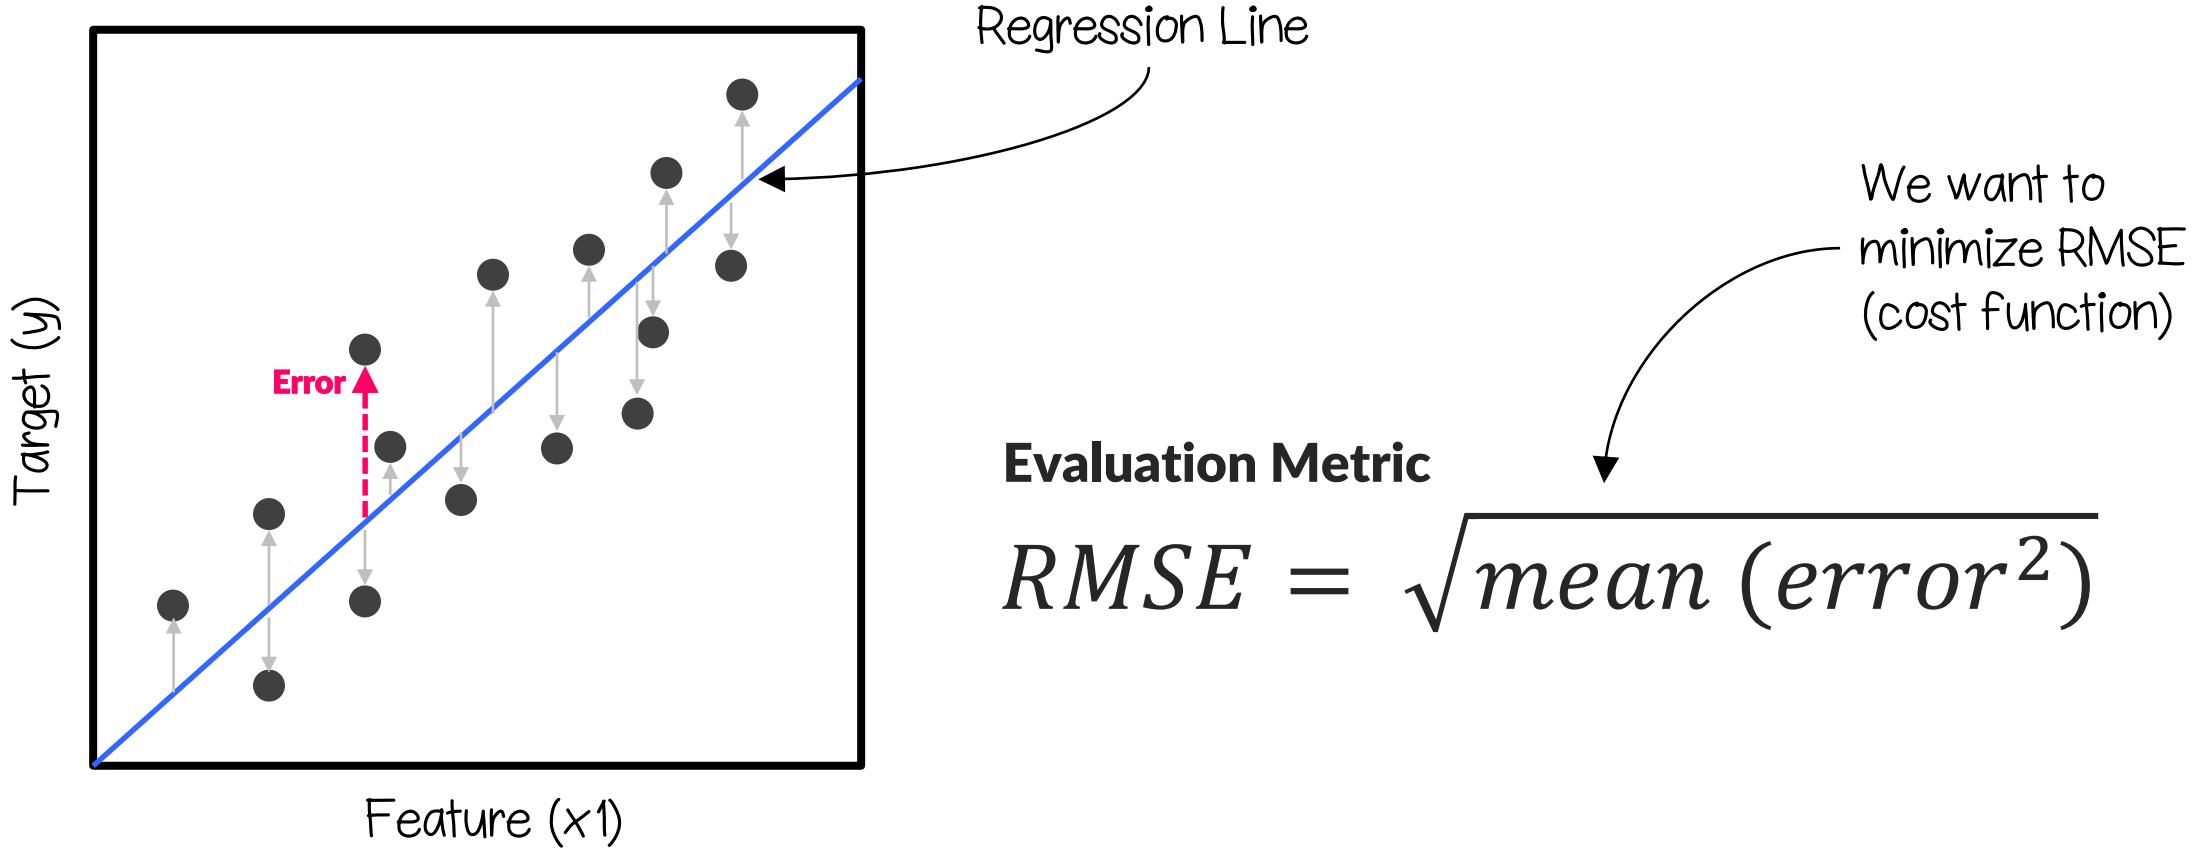


Recall these steps

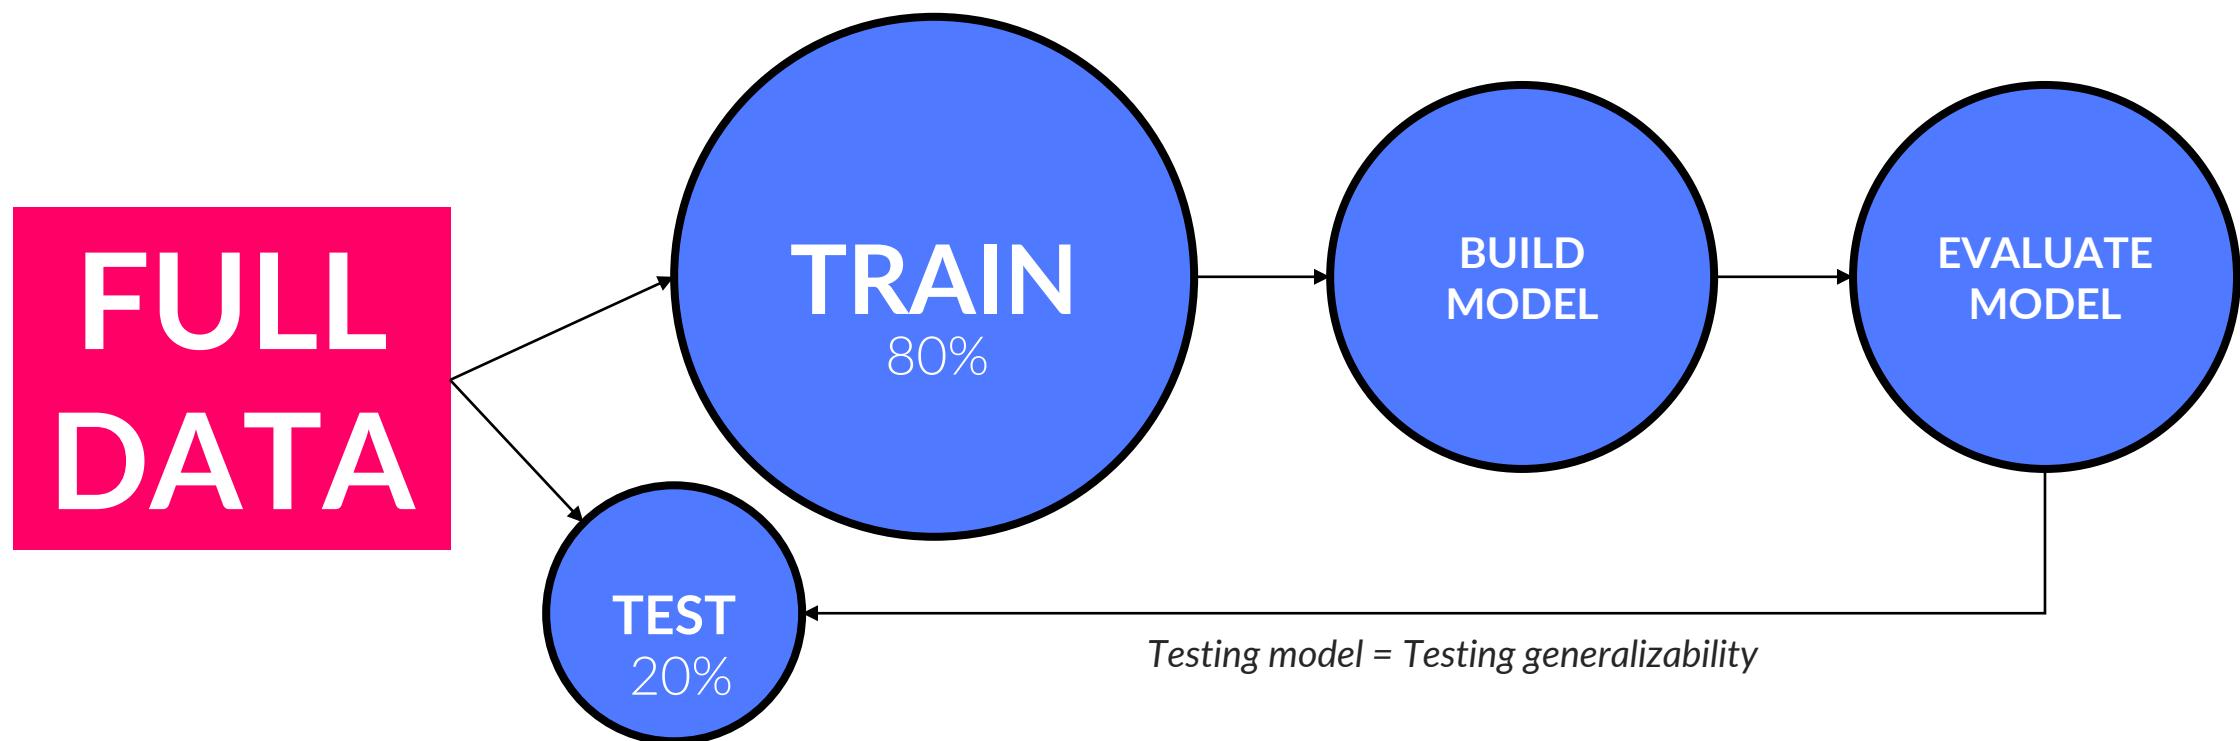


Algorithm explained:

How to draw a single straight line that fit best with given data points



Let's do a **train-test split** for linear regression problem



Advantages:

- *Easy to train*
- *Always find the optimum solutions (least squared)*
- *Foundation of advanced algorithms like neural nets*

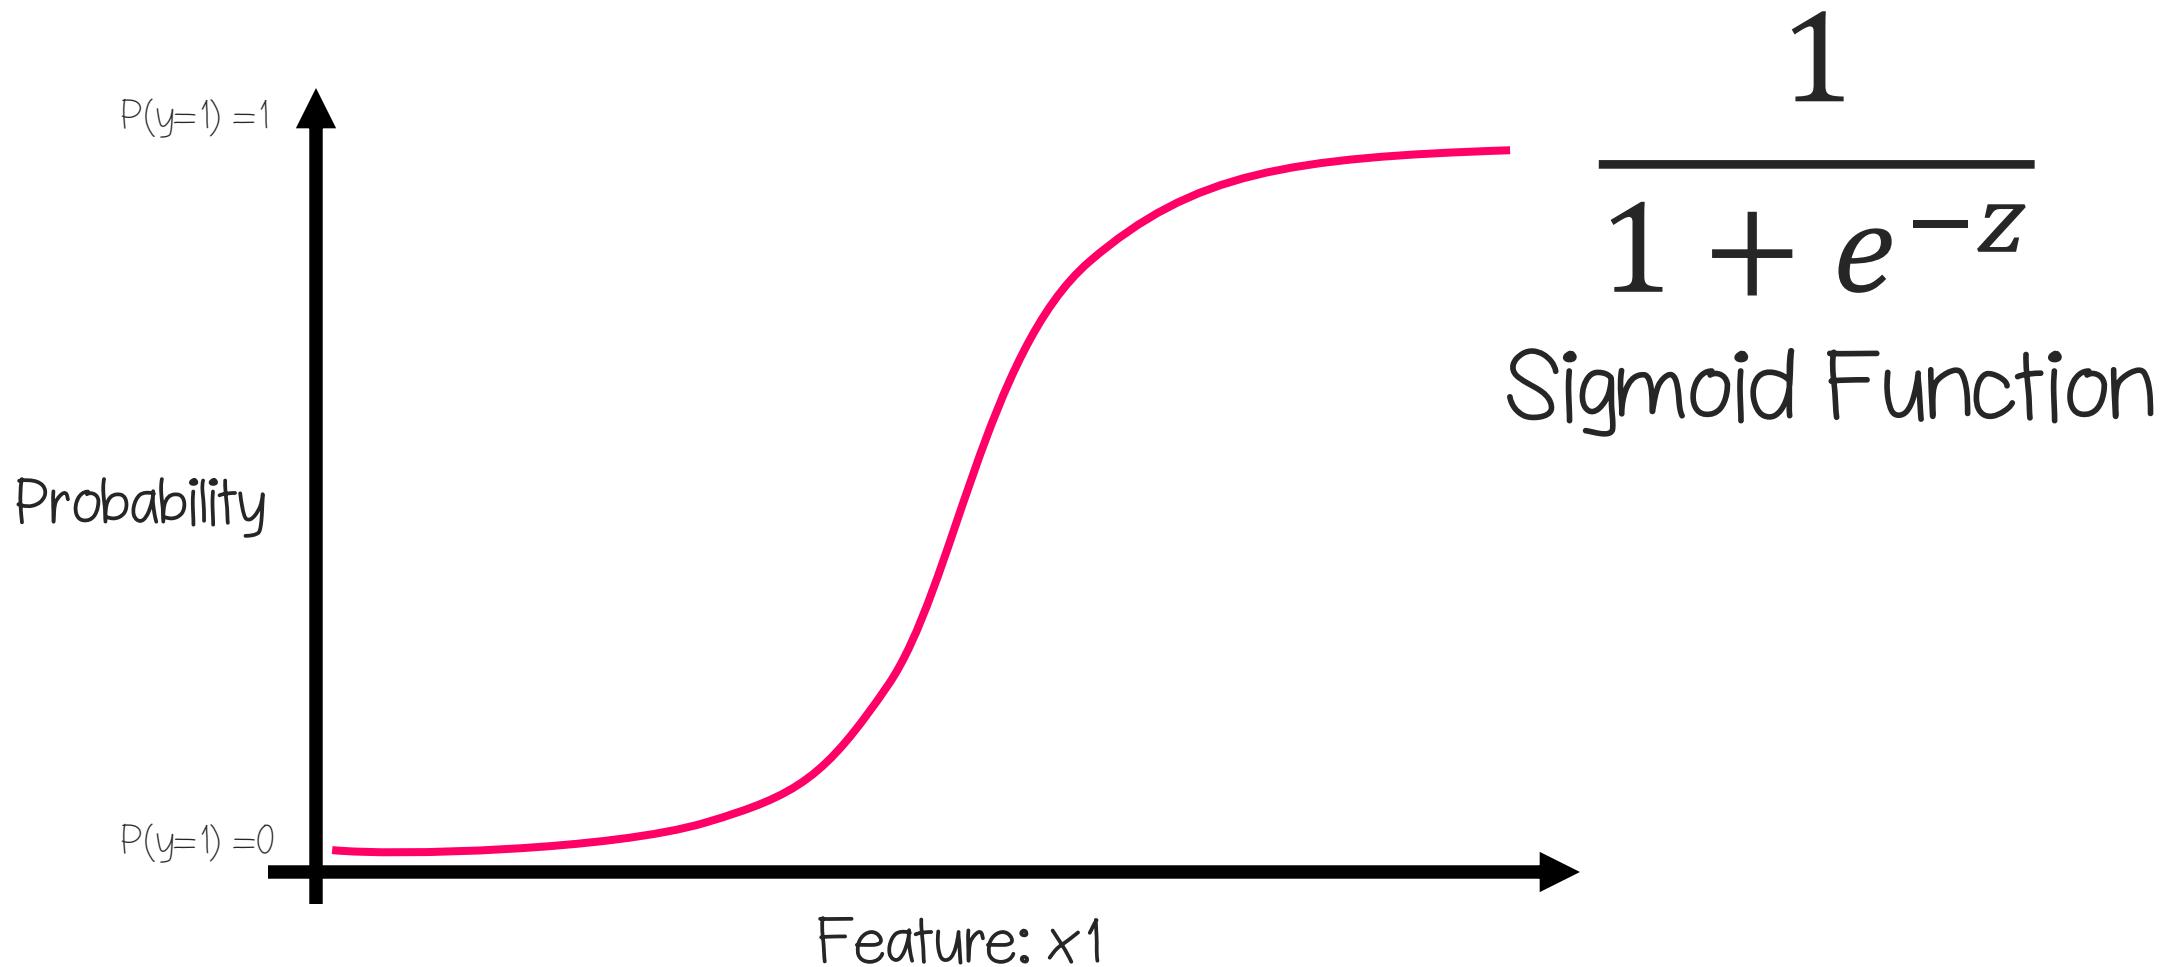
Disadvantages:

- *Only work for regression problems (continuous y)*
- *Full of assumptions (high bias algorithms)*

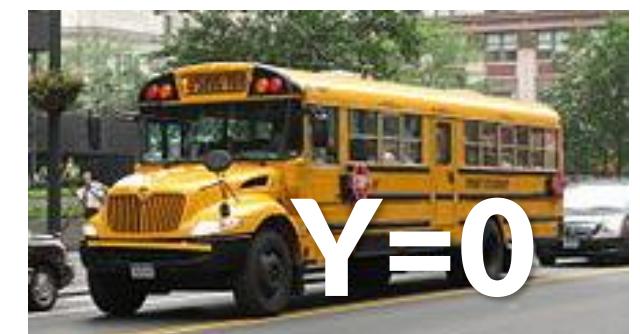
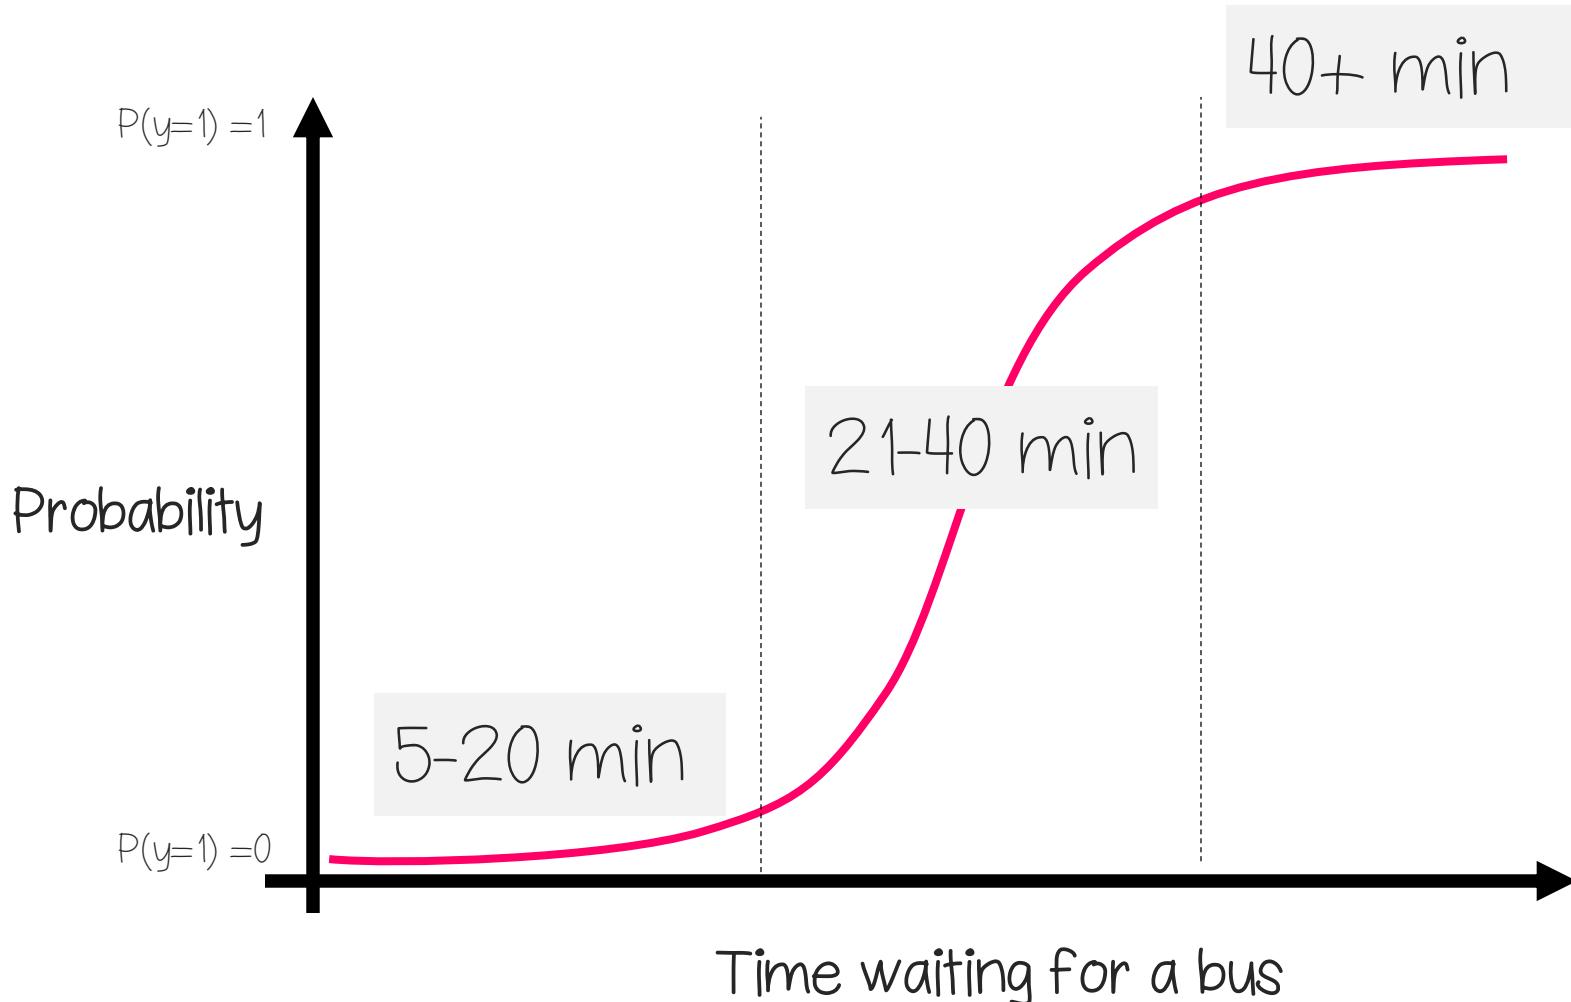
LOGISTIC REGRESSION

Algorithm explained:

Classify data points by calculating the **probability** that Y=1



Should I get taxi or a bus?



The model is very similar to linear regression

(assuming linear relationship between x and y)

$z = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + b_4x_4 \rightarrow$ then $1 / 1 + e^{-z}$ to get probability scores

ID	X1 (minute)	X2 (crowd)	X3 (time)	X4 (weather)	P(Y=1) (get Taxi)	Prediction
1	50	Many	Friday	Good	.95	Yes
2	20	Many	Saturday	OK	.32	No
3	20	Few	Monday	OK	.31	No
4	25	Few	Friday	Bad	.85	Yes
5	30	Many	Friday	Good	.68	Yes
6	35	Few	Saturday	OK	.25	No
7	12	Many	Monday	Good	.52	Yes
8	5	Few	Saturday	OK	.12	No
9	10	Many	Friday	Bad	.20	No
10	65	Few	Friday	Good	.98	Yes



We set threshold = 0.5
If $p(y=1) \geq 0.5$, then we predicted YES get taxi!



Advantages:

- One of the most powerful algorithms to date
- Easy to train
- Can be developed to create complex decision boundaries (e.g. x^2 y^2)

Disadvantages:

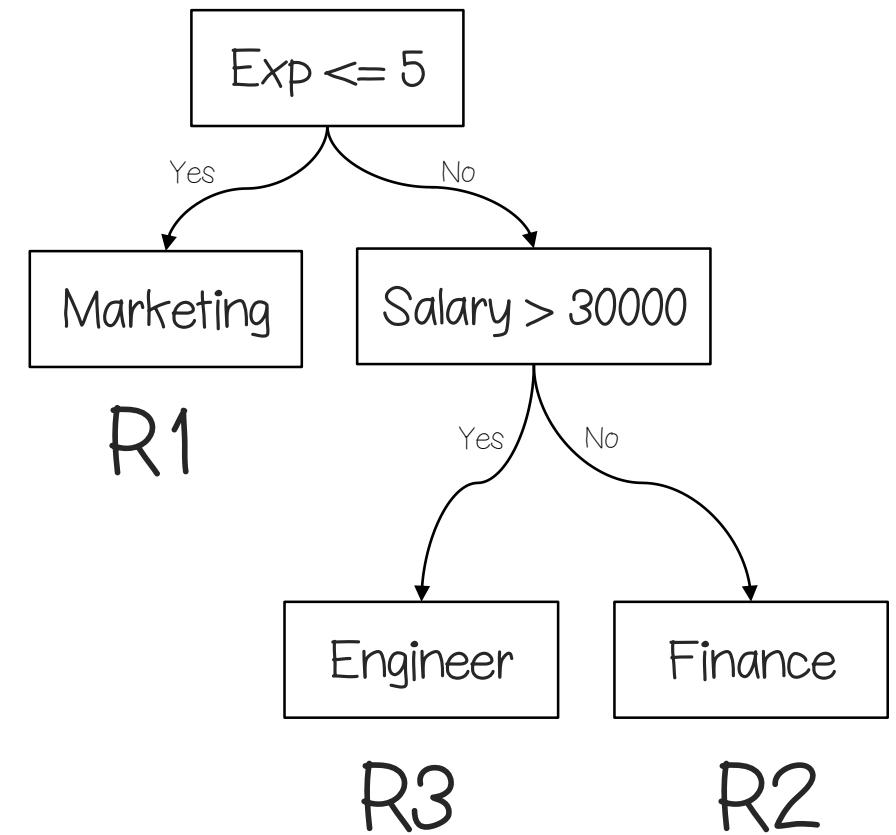
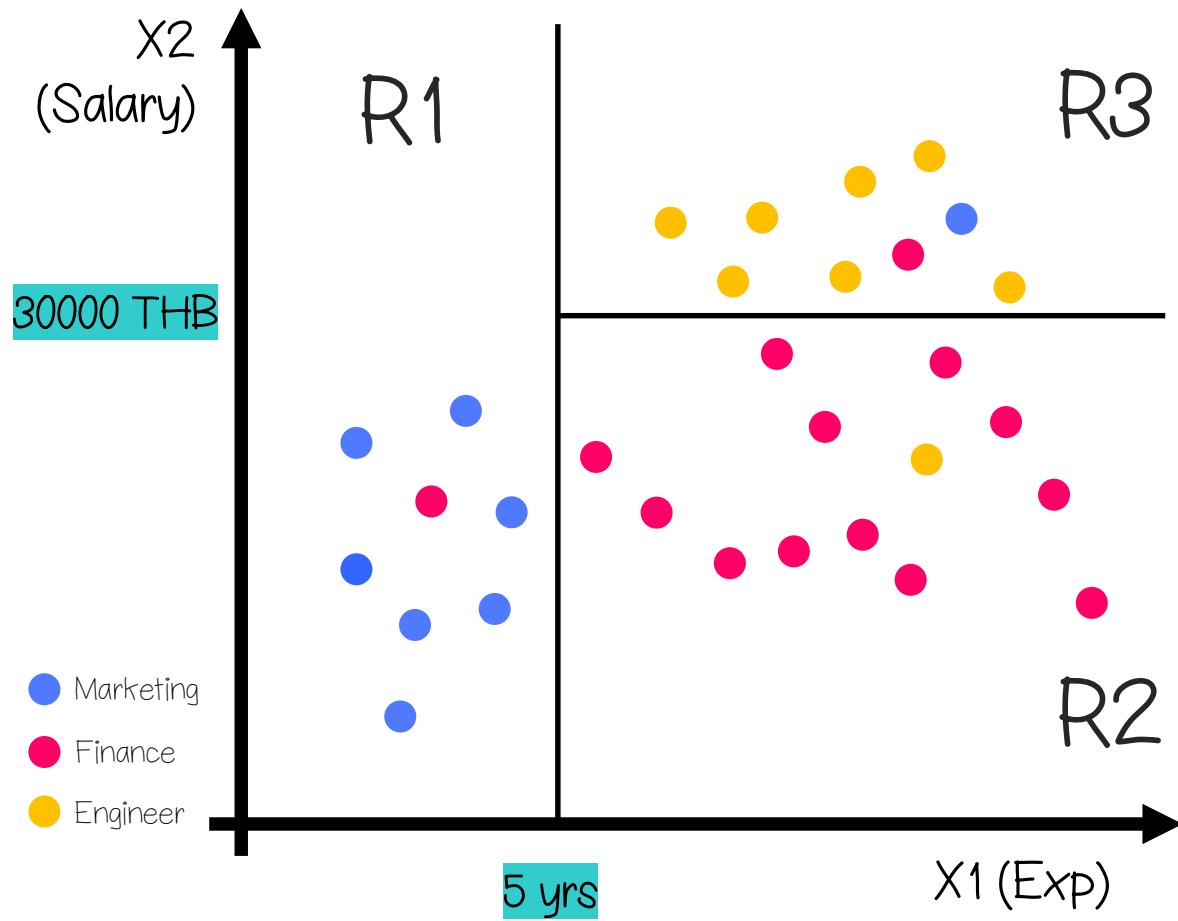
- Also a lot of assumption e.g. no outlier, no noise
- Logistic regression is also a linear algorithm
- Prone to multicollinearity (check correlation first)

(sigmoid function transform an output into non-linear)

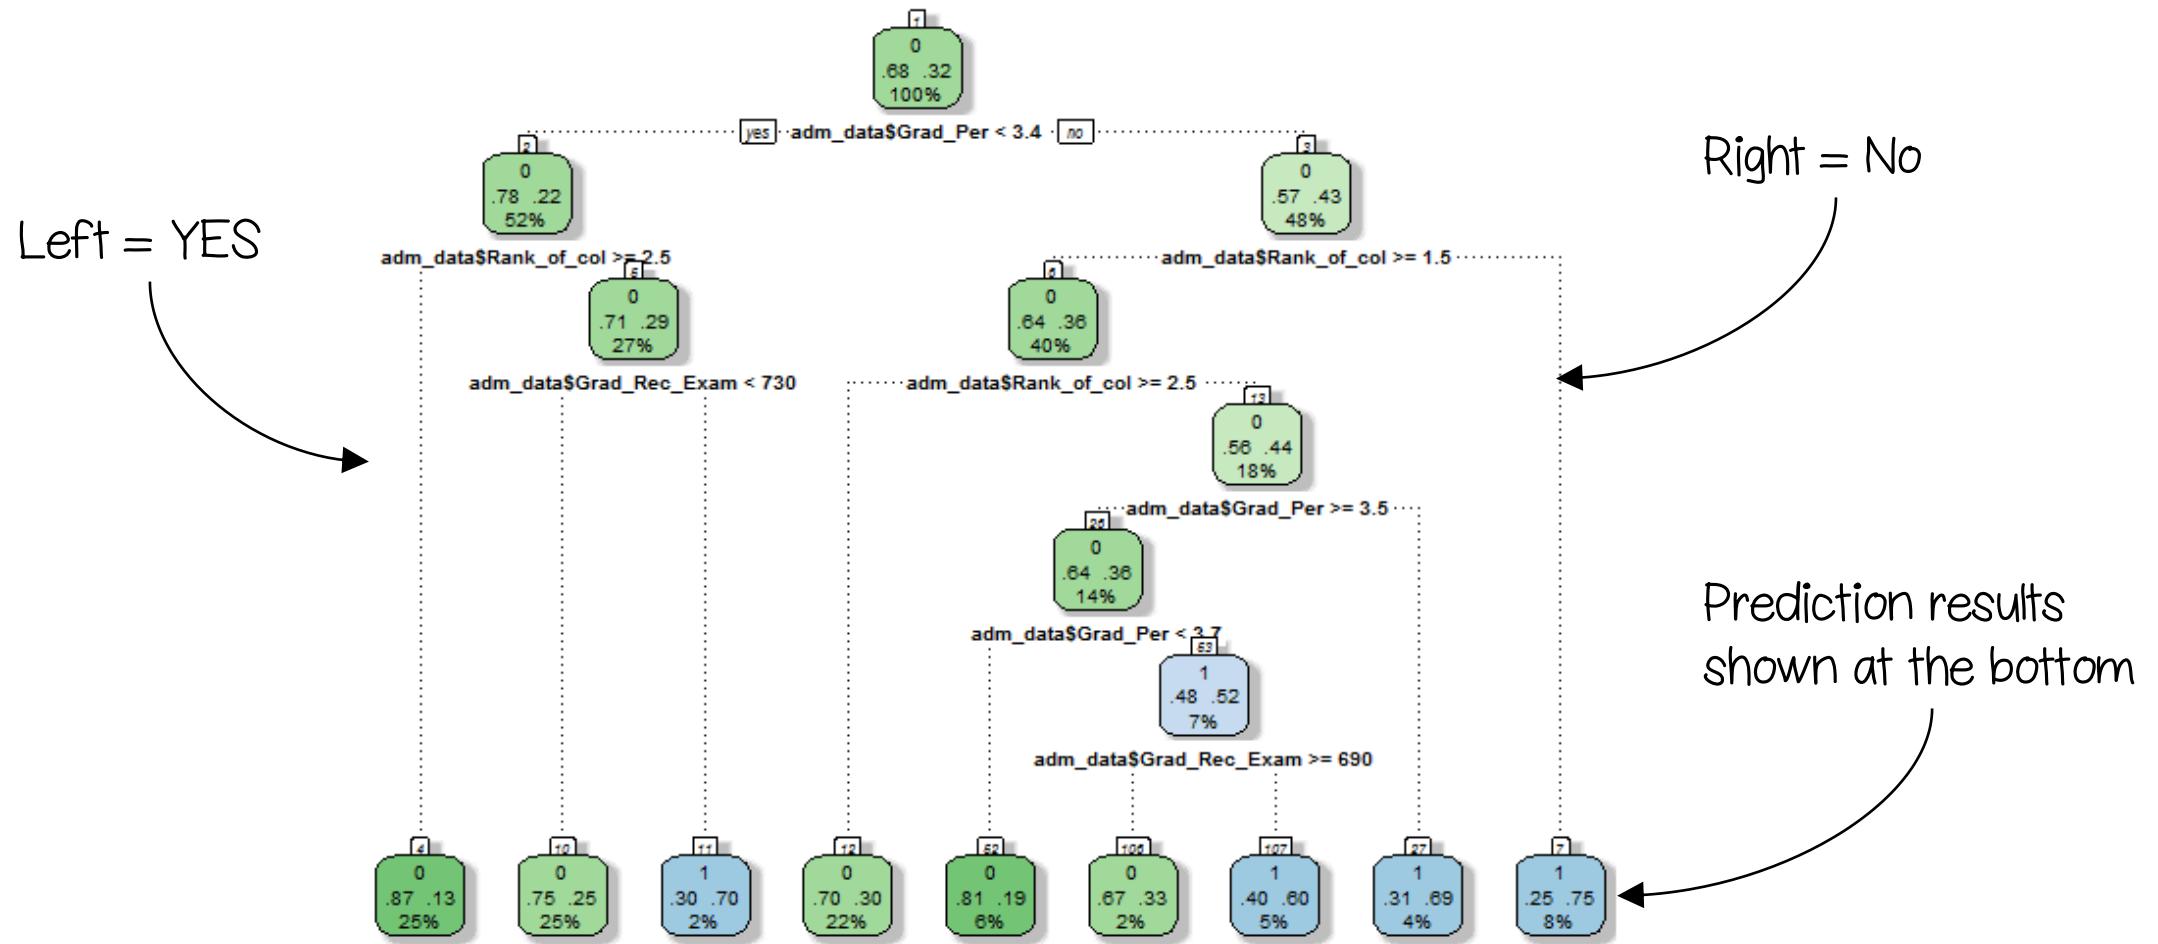
DECISION TREE

Algorithm explained:

Divide space into regions that best predict data points (classes)



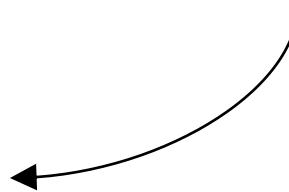
Decision Tree is easily presented using diagram



Advantages:

- *Easy to train*
- *Easy to interpret results*
- *Widely used in many disciplines*

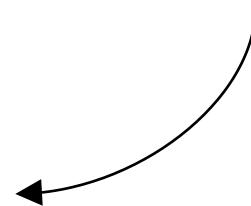
The main reason why people use decision tree



Disadvantages:

- *Prone to overfitting: Greedy algorithms*

Locally optimal decision are made at each node, but can't guarantee globally optimal decision trees



Required Packages:

`library(caret)` – train & test models

`library(caTools)` – split data into training & testing sets

`library(rpart)` – train decision tree (CART): recursive partitioning and regression tree

CART

```
method = 'rpart'
```

Type: Regression, Classification

Tuning parameters:

- `cp` (Complexity Parameter)

Required packages: `rpart`

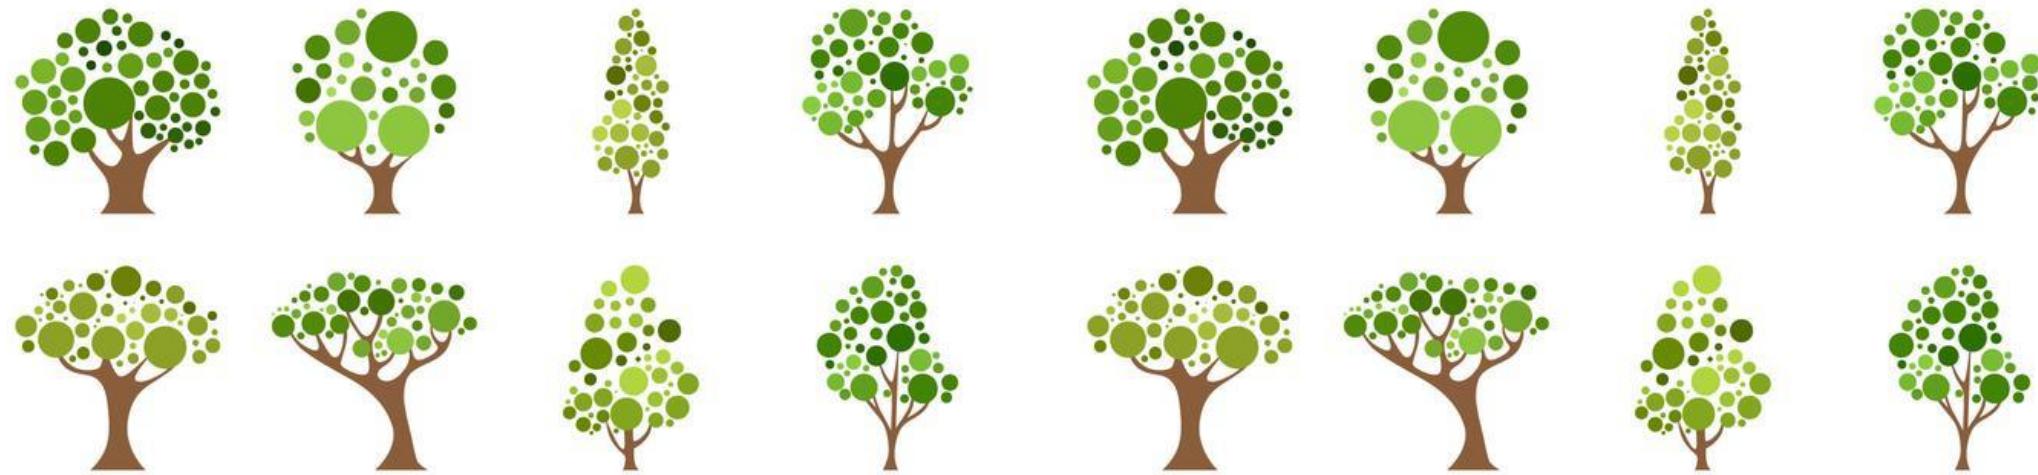
A model-specific variable importance metric is available.

Cp: complexity parameter
to penalize larger trees

RANDOM FOREST

Algorithm explained:

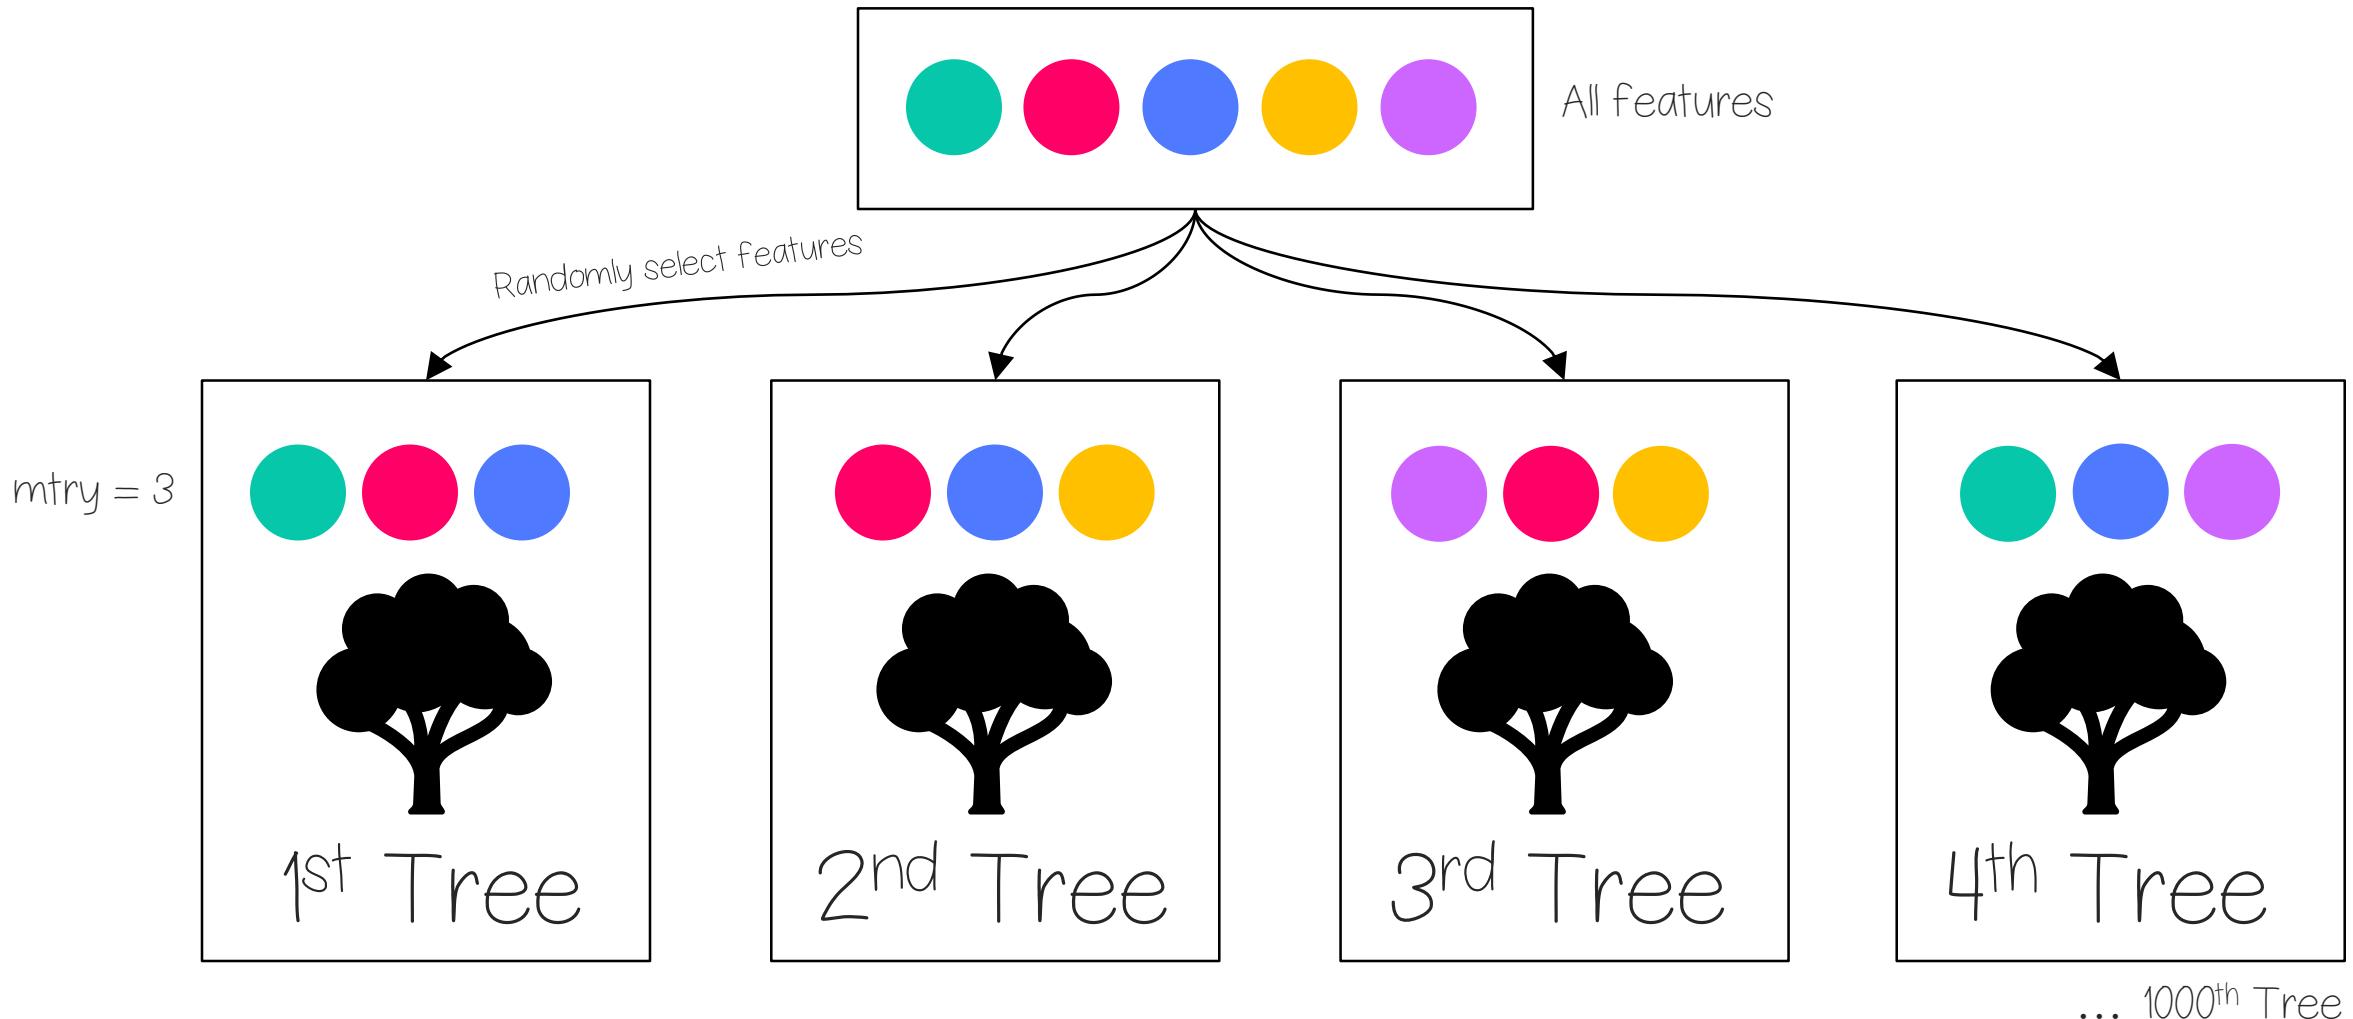
Grow thousand trees and average their prediction results



Grow *uncorrelated* 1000 trees

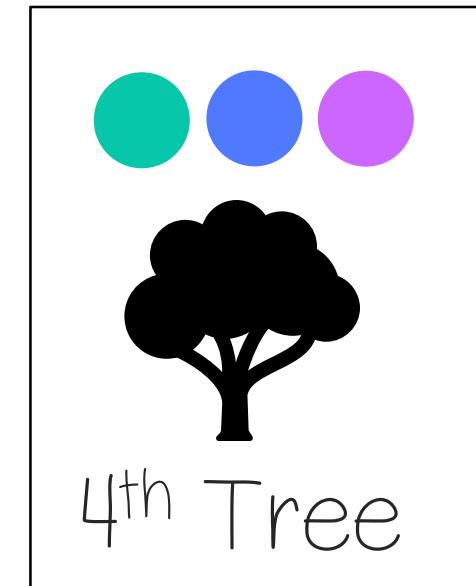
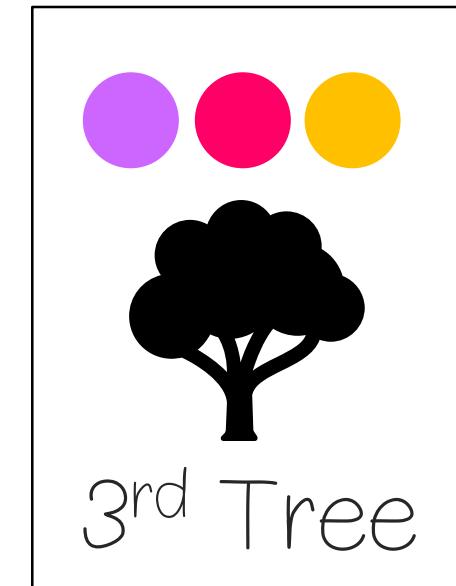
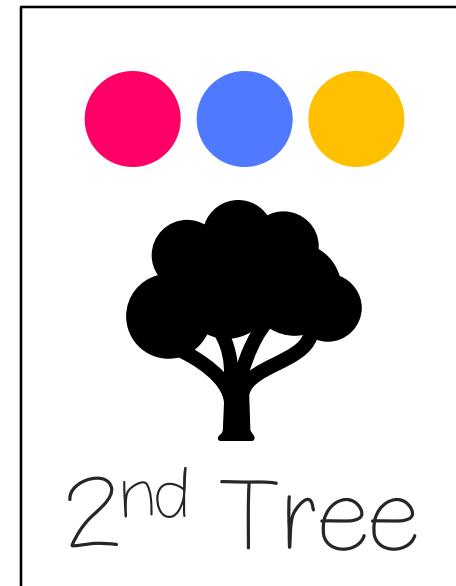
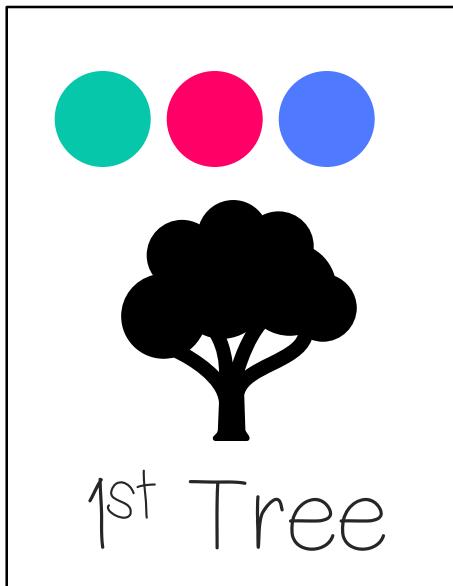
How do we grow 1000 trees?

Randomness is your friend, hyperparameter is **mtry**



Wisdom of the crowds

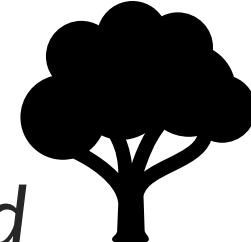
Find the average prediction results of all grown trees



... 1000th Tree

Advantages:

- *Amongst the best learners in the world*
- *Easy to train (Improvement from Decision Tree)*



Disadvantages:

- *Can be slow to train for large dataset*
- *Black Box Algorithm (not easy to explain)*

Required Packages:

`library(caret)` – train & test models

`library(caTools)` – split data into training & testing sets

`library(randomForest)` – train random forest in caret

Random Forest

```
method = 'rf'
```

Type: Classification, Regression

Tuning parameters:

- `mtry (#Randomly Selected Predictors)`

Required packages: `randomForest`

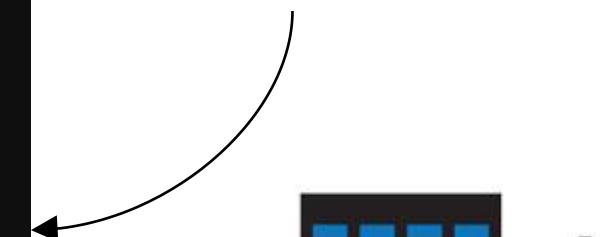
A model-specific variable importance metric is available.

Tune `mtry` parameter to
find the best prediction
from all trees

Let's talk about the hardware side



Graphic Cards
have GPU 1000
cores



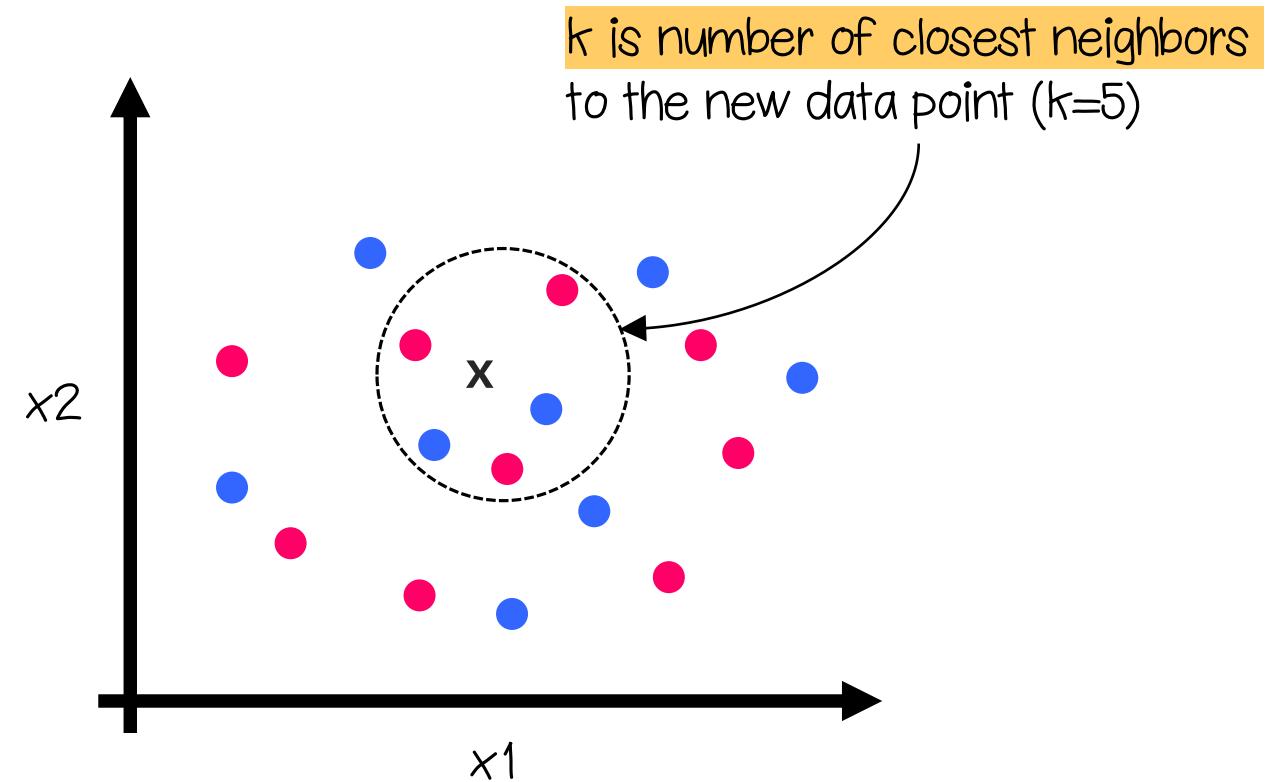
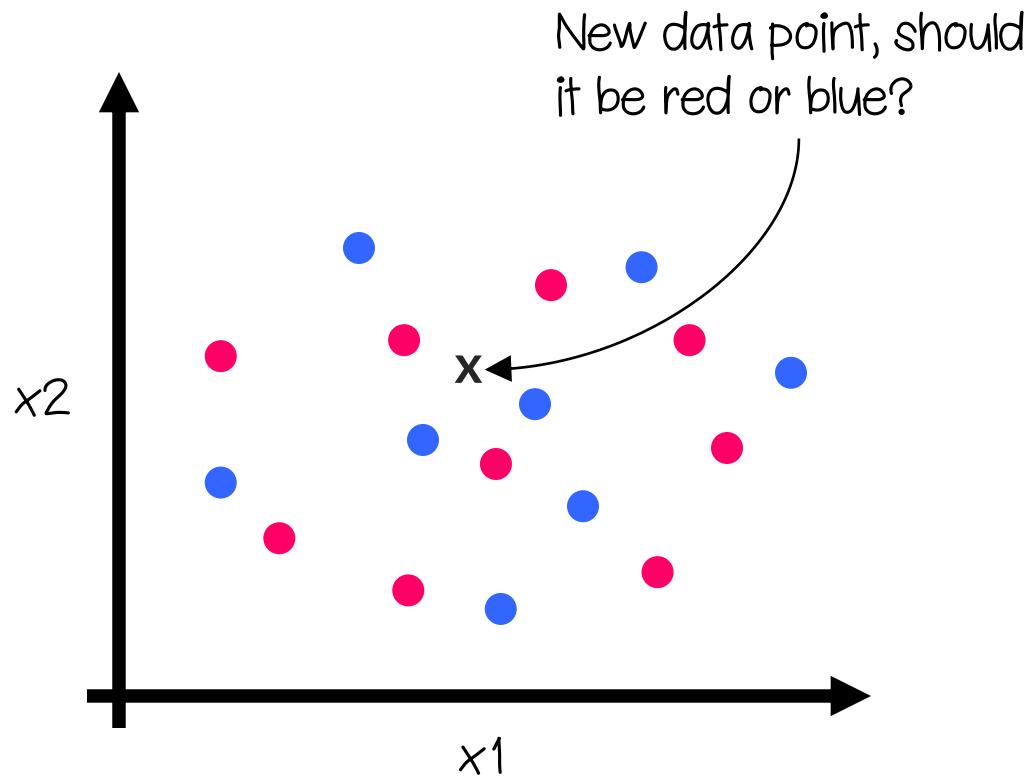
Nvidia has CUDA technology

Read more: <https://developer.nvidia.com/about-cuda>

KNN (K-NEAREST NEIGHBOR)

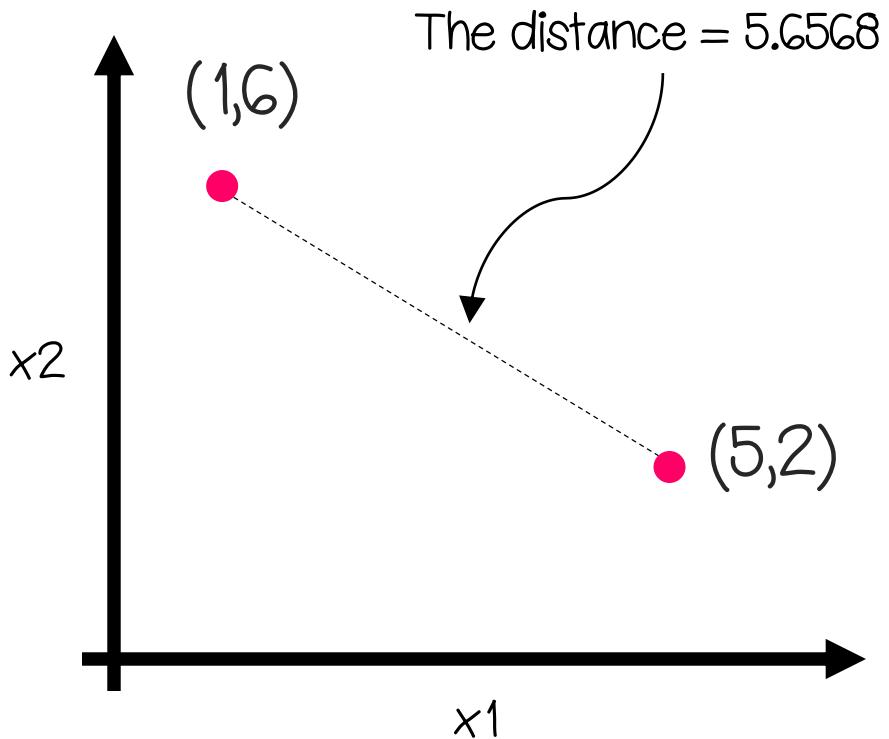
Algorithm explained:

Predict new data according to the distance from its neighbors



Euclidean Distance:

The most popular distance calculation method



$$d = \sqrt{\sum (x_i - y_i)^2}$$
$$d = \sqrt{(1 - 5)^2 + (6 - 2)^2}$$
$$d = \sqrt{16 + 16}$$

$d = 5.656854$

Very simple
to compute

Another Example

Consider the `nutrient` dataset provided with the `flexclust` package. The dataset contains measurements on the nutrients of 27 types of meat, fish, and fowl. The first few observations are given by

```
> data(nutrient, package="flexclust")
> head(nutrient, 4)
```

	energy	protein	fat	calcium	iron
BEEF BRAISED	340	20	28	9	2.6
HAMBURGER	245	21	17	9	2.7
BEEF ROAST	420	15	39	7	2.0
BEEF STEAK	375	19	32	9	2.6

We can calculate many dimensions (more than two)

and the Euclidean distance between the first two (beef braised and hamburger) is

$$d = \sqrt{(340 - 245)^2 + (20 - 21)^2 + (28 - 17)^2 + (9 - 9)^2 + (2.6 - 2.7)^2} = 95.64$$

Advantages:

- *Very simple*
- *Easy to train (remember all data)*

Disadvantages:

- *Worst for large dataset*
- *Not good with high dimensional data (many features)*
- *Don't work well with categorical data*

Example Code:

```
5 ## load library  
6 library(caret)  
7  
8 ## train model  
9 ## define trControl  
10 my_trcontrol = trainControl(method = "cv",  
11                               number = 3,  
12                               verboseIter = TRUE,  
13                               search = "random") }  
14  
15 knn_model <- train(Species ~ . , data = m,  
16                      method = "knn",  
17                      metric = "Accuracy",  
18                      tuneLength = 10,  
19                      trControl = my_trcontrol) }
```

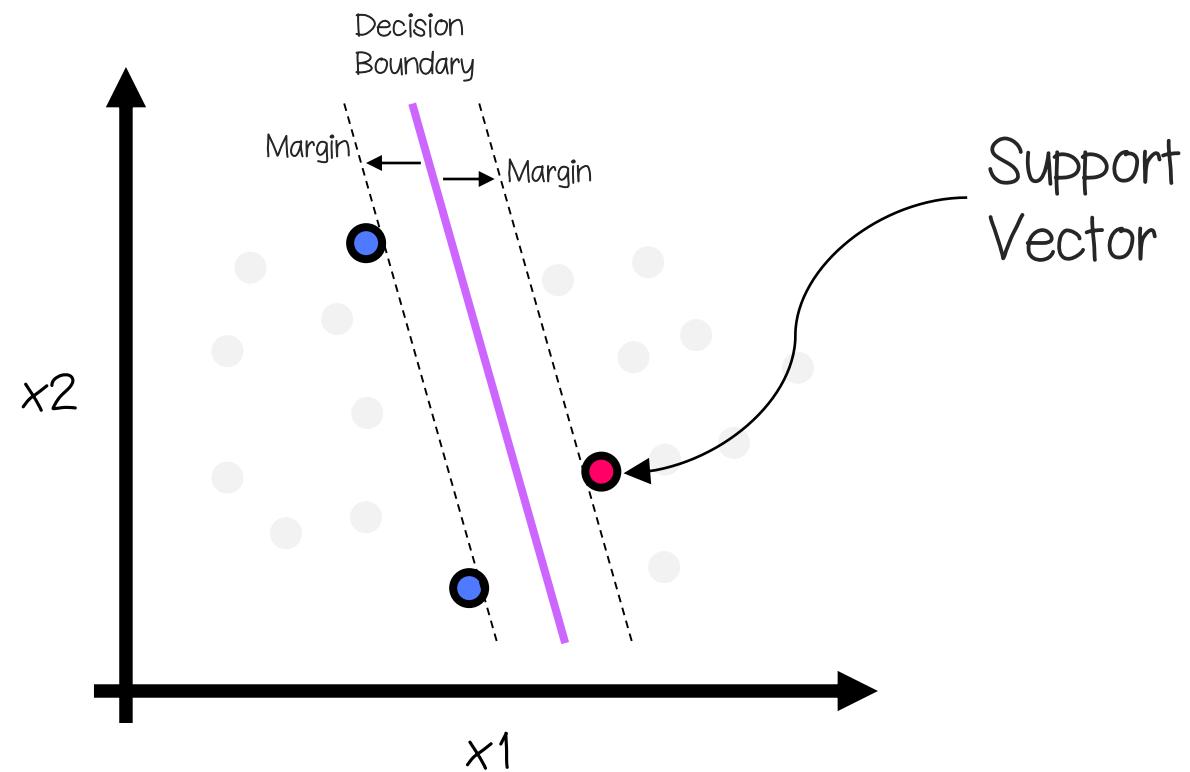
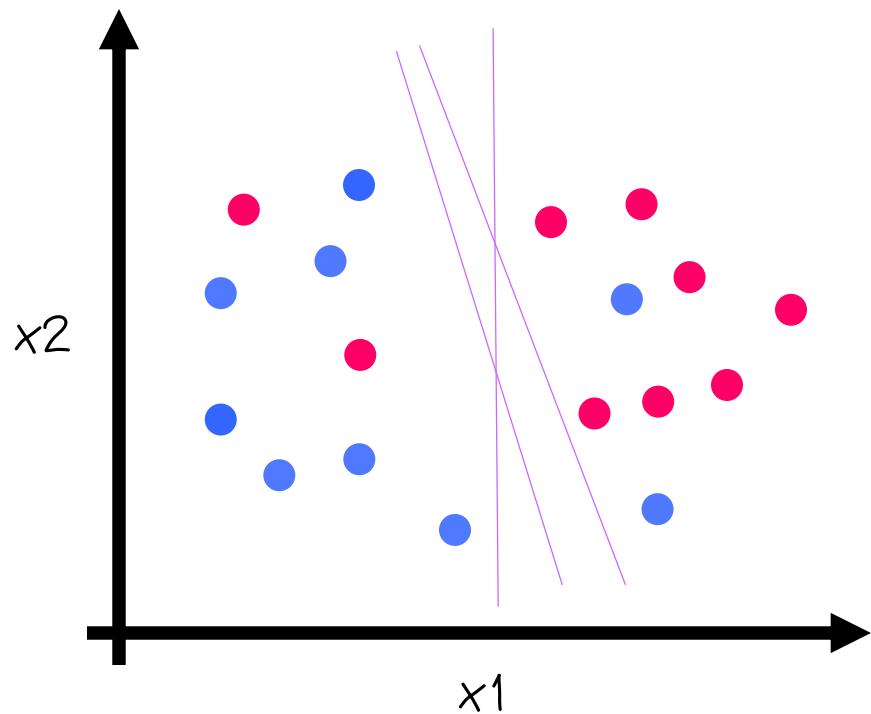
Define our technique
(to train the model)

Train the model using
random search +
tuneLength = 10

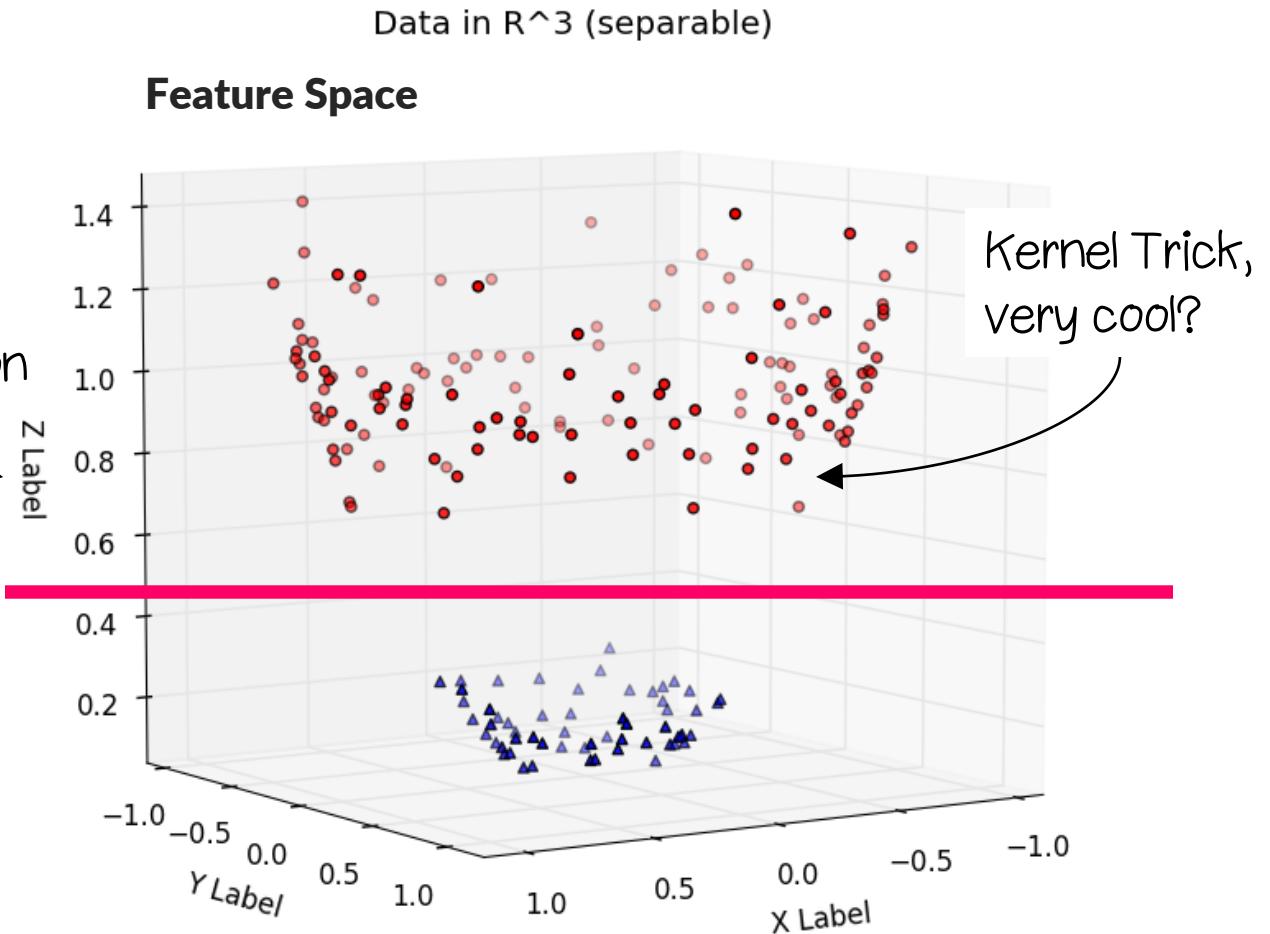
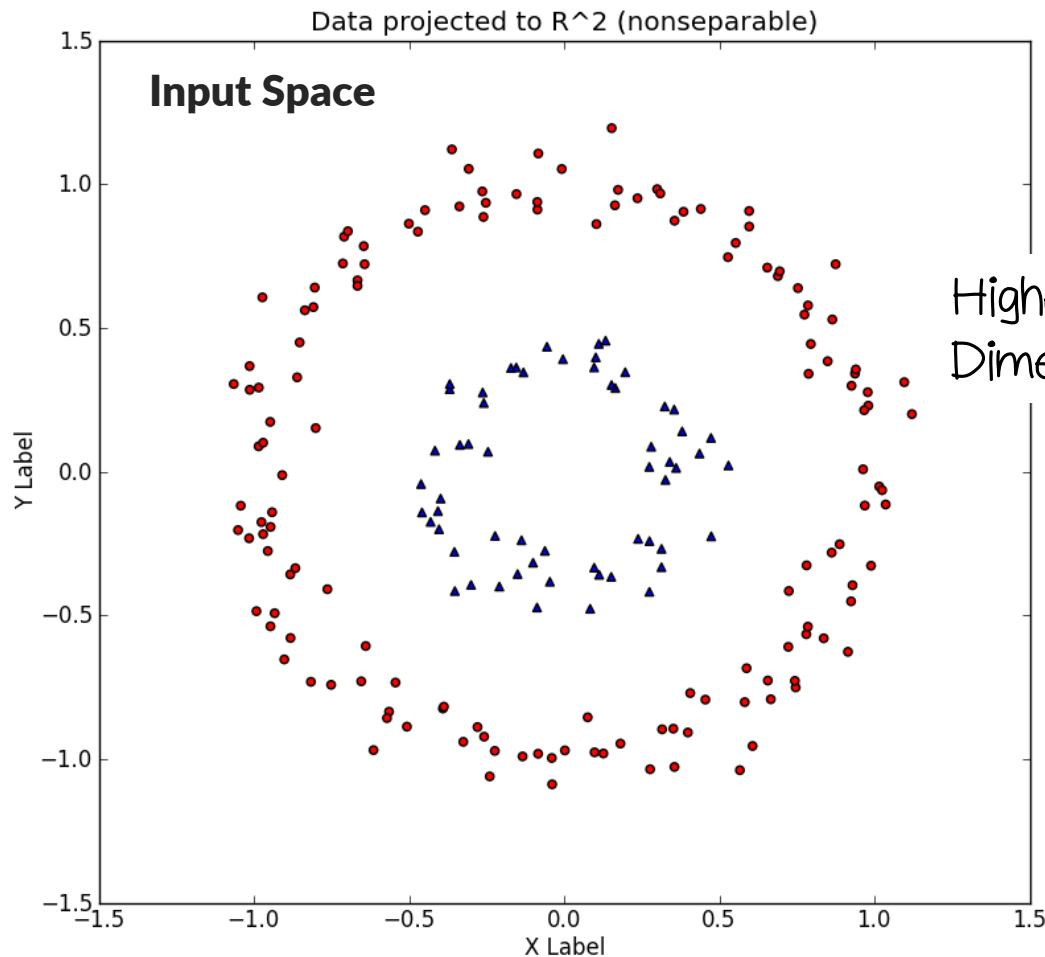
SUPPORT VECTOR MACHINE

Algorithm explained:

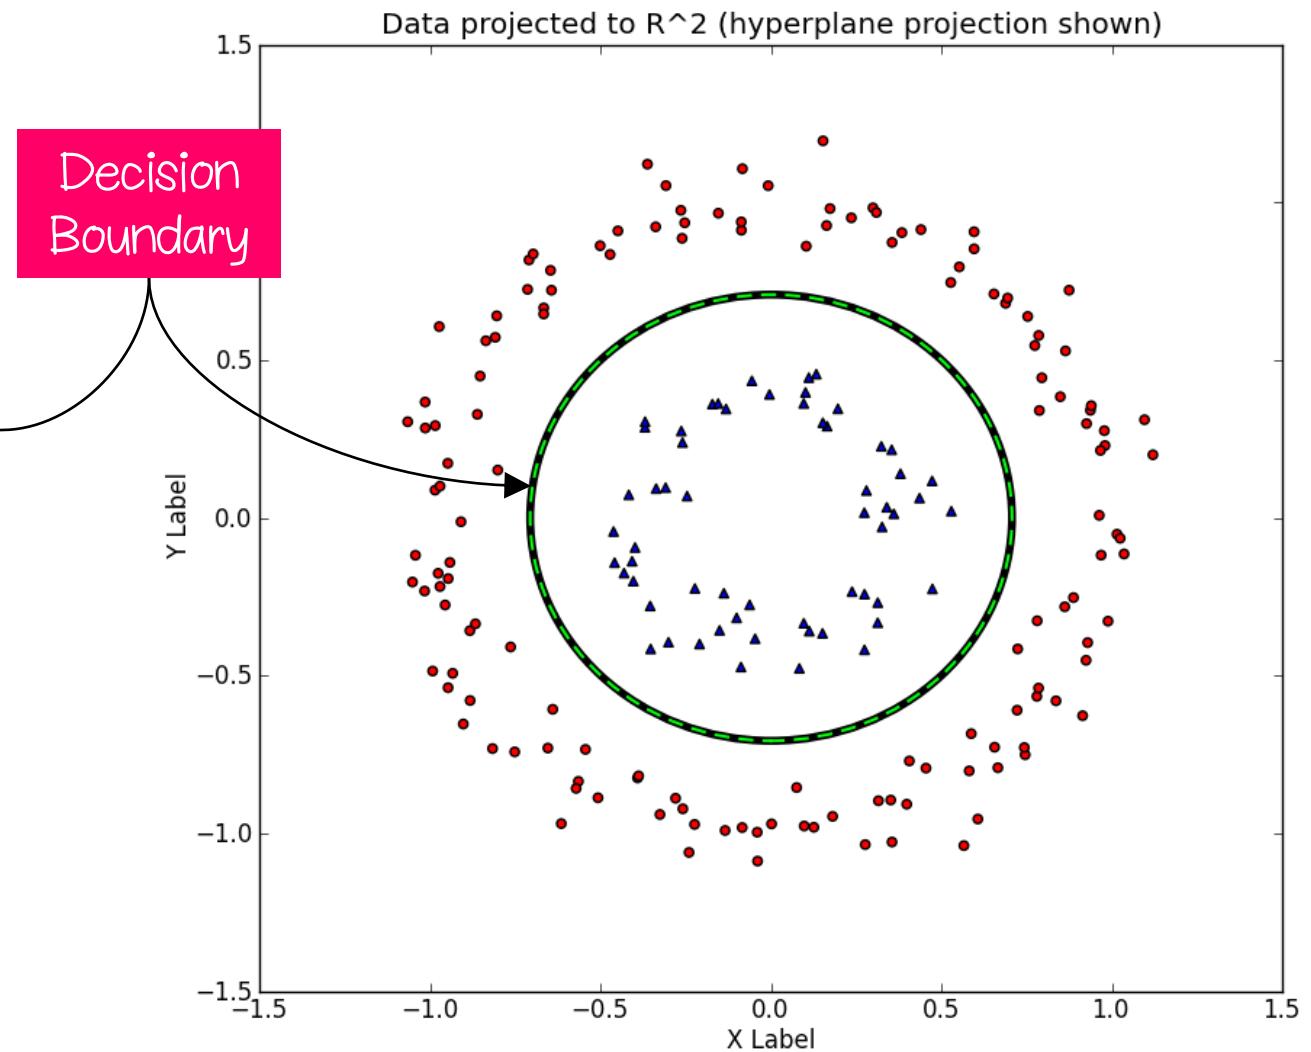
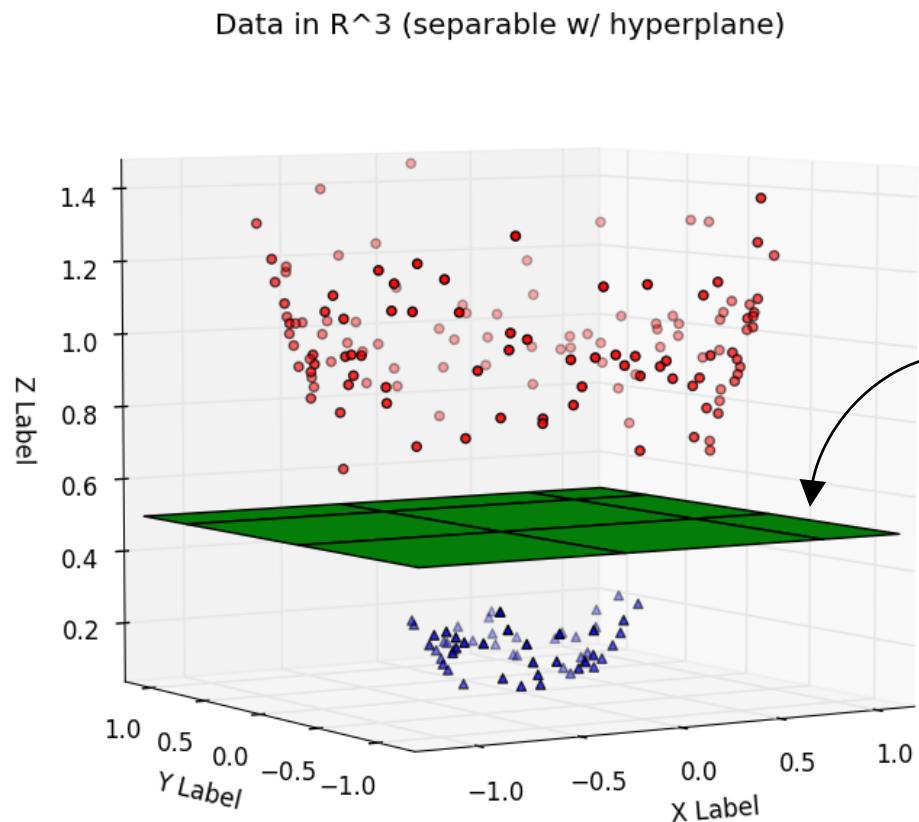
We want to maximize the **margin** from decision boundary



But can you divide this with a single straight line?

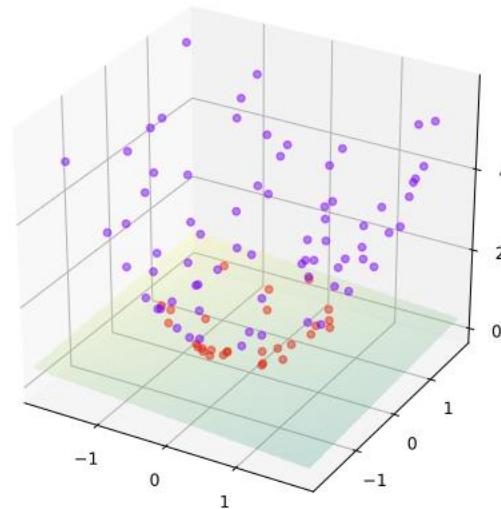
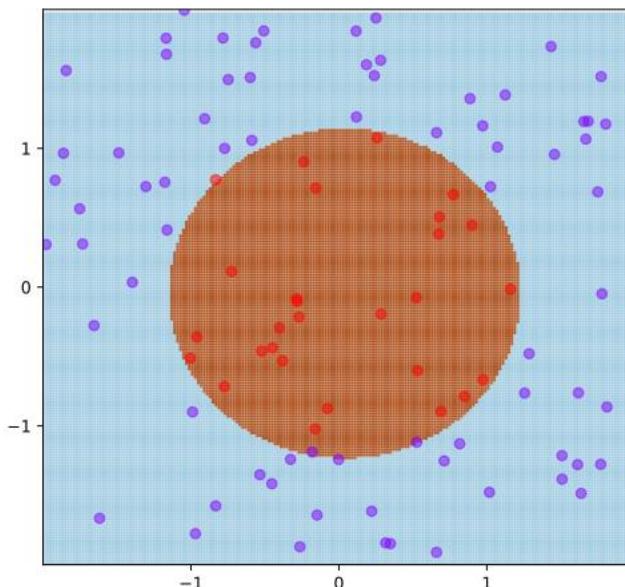


Impressive Result!!



Let's look closely at kernel trick

- we have (x,y) coordinates
- we project a new higher dimension: $(x, y, x^2 + y^2)$



3rd dimension

https://en.wikipedia.org/wiki/Kernel_method

Advantages:

- *Work with Regression | Classification problems*
- *Superior Performance*
- *Having lots of available kernel tricks*

<https://topepo.github.io/caret/train-models-by-tag.html#support-vector-machines>

Disadvantages:

- *Expensive*
- *It's lazy algorithm (remember data like KNN)*

Required Packages:

`library(caret)` – train & test models

`library(caTools)` – split data into training & testing sets

`library(e1071)` – load some algorithms e.g. SVM

Support Vector Machines with Linear Kernel

```
method = 'svmLinear2'
```

Type: Regression, Classification

Tuning parameters:

- `cost` (Cost)

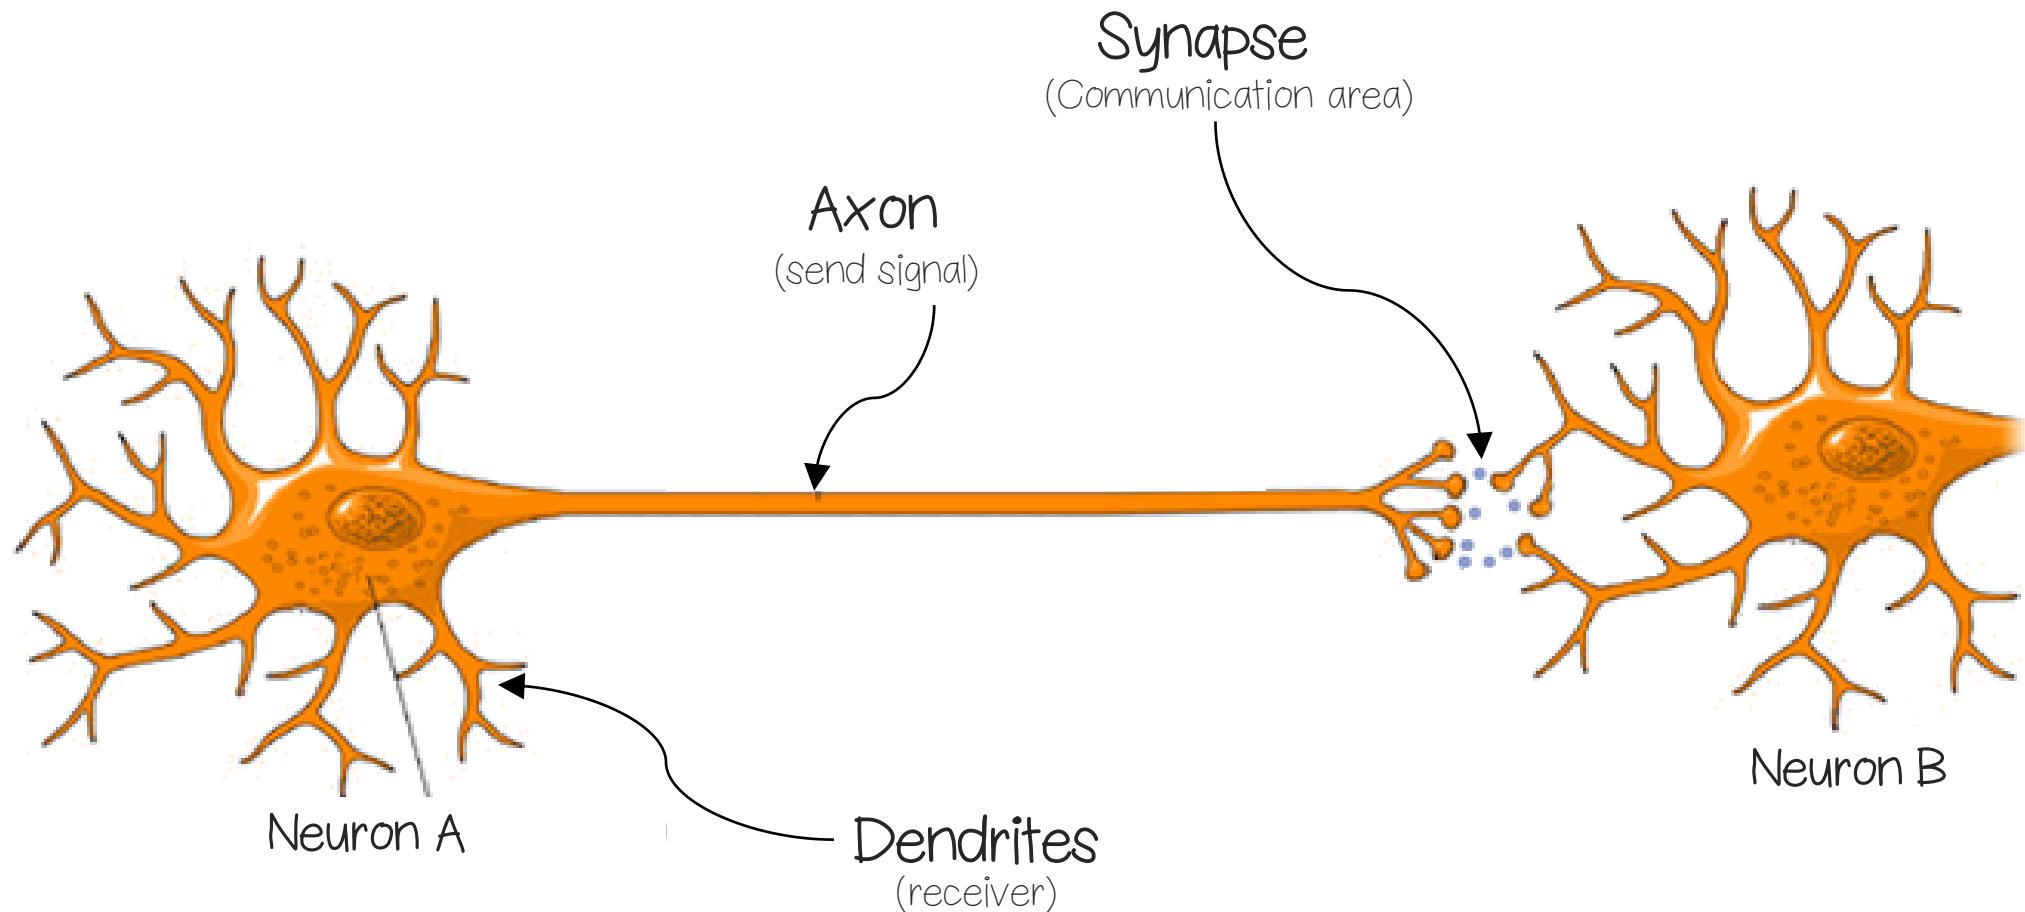
Required packages: `e1071`

NEURAL NETWORK



Neuron →

How your brain cells communicate



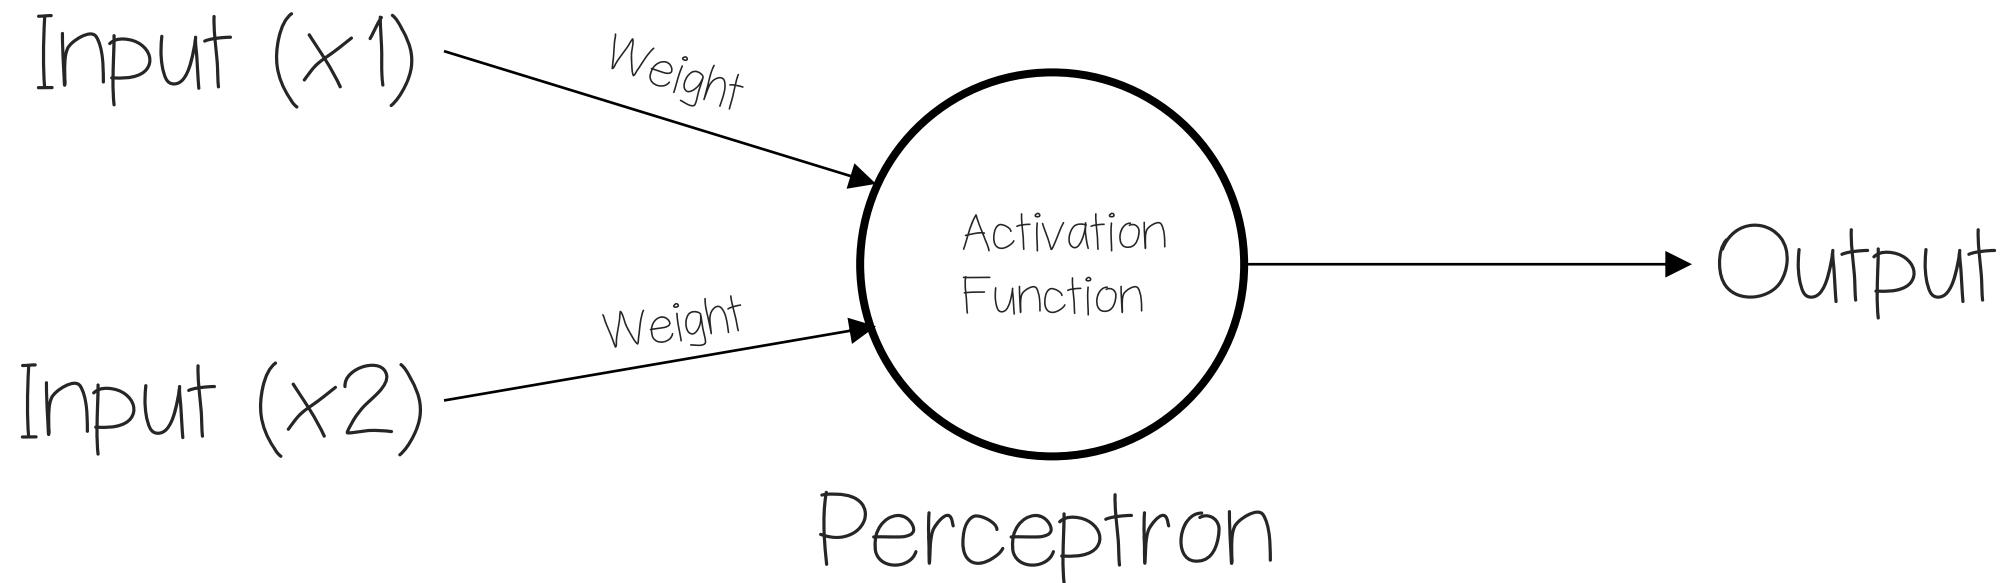


If communicated successfully,
the neuron fire!



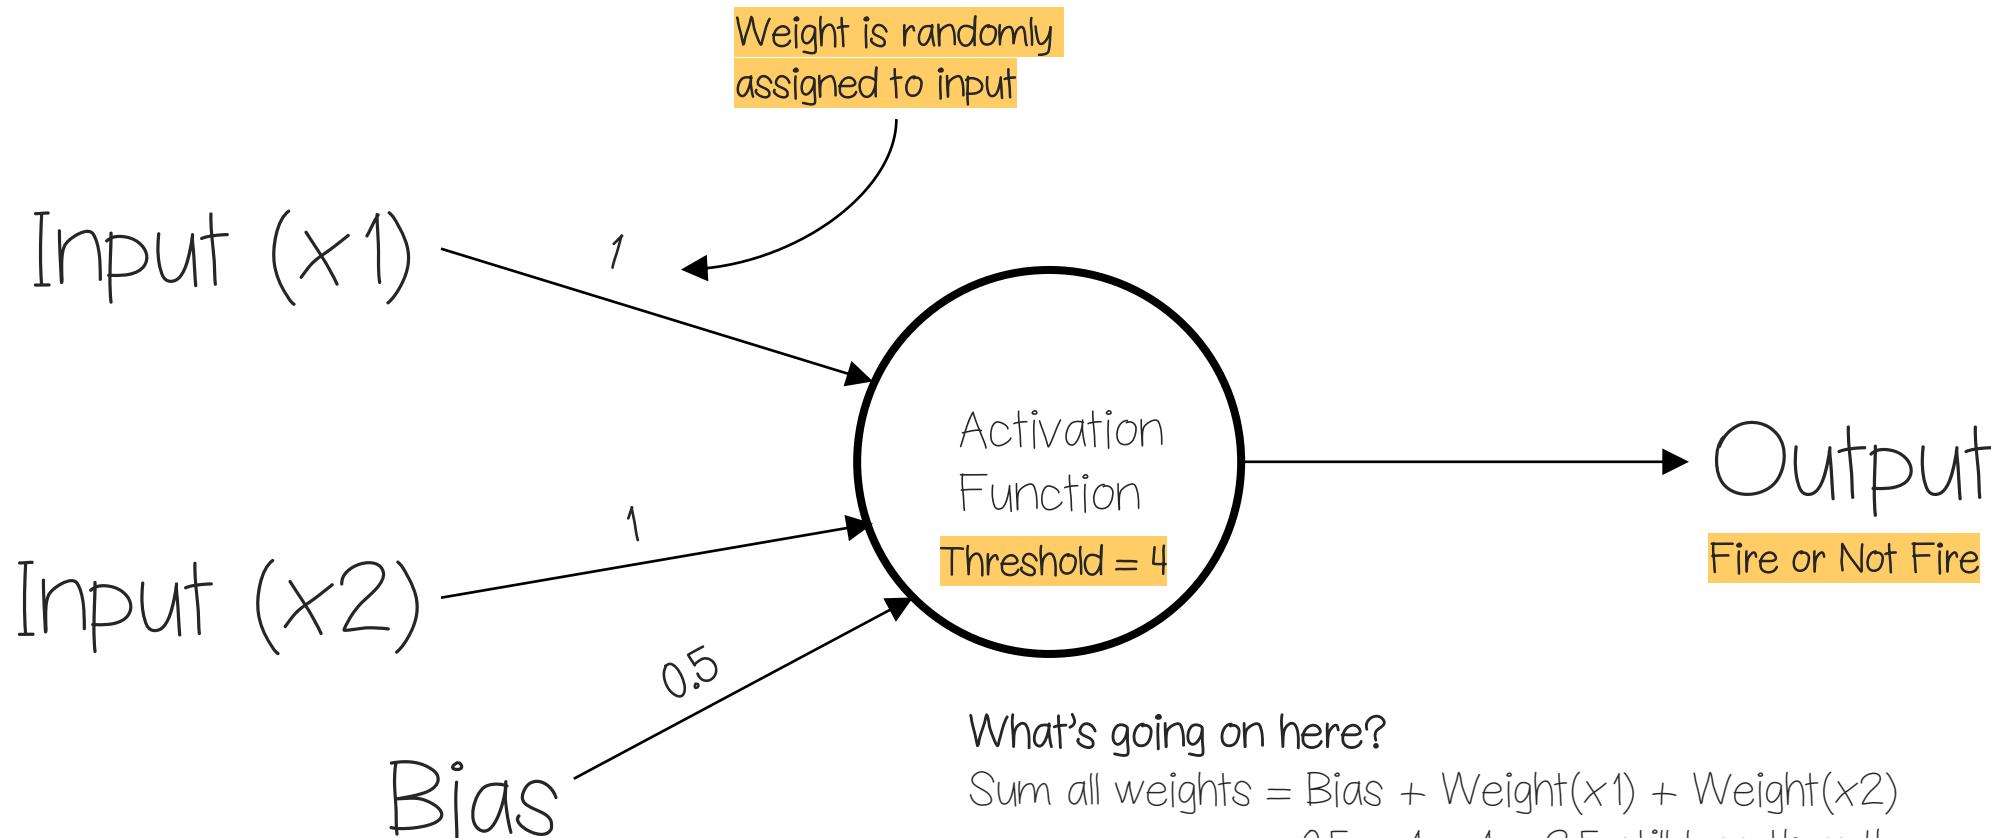
Algorithm Explained:

Neurons fire only when it's the signal strong enough to activate them



Algorithm Explained:

Neurons fire only when it's the signal strong enough to activate them

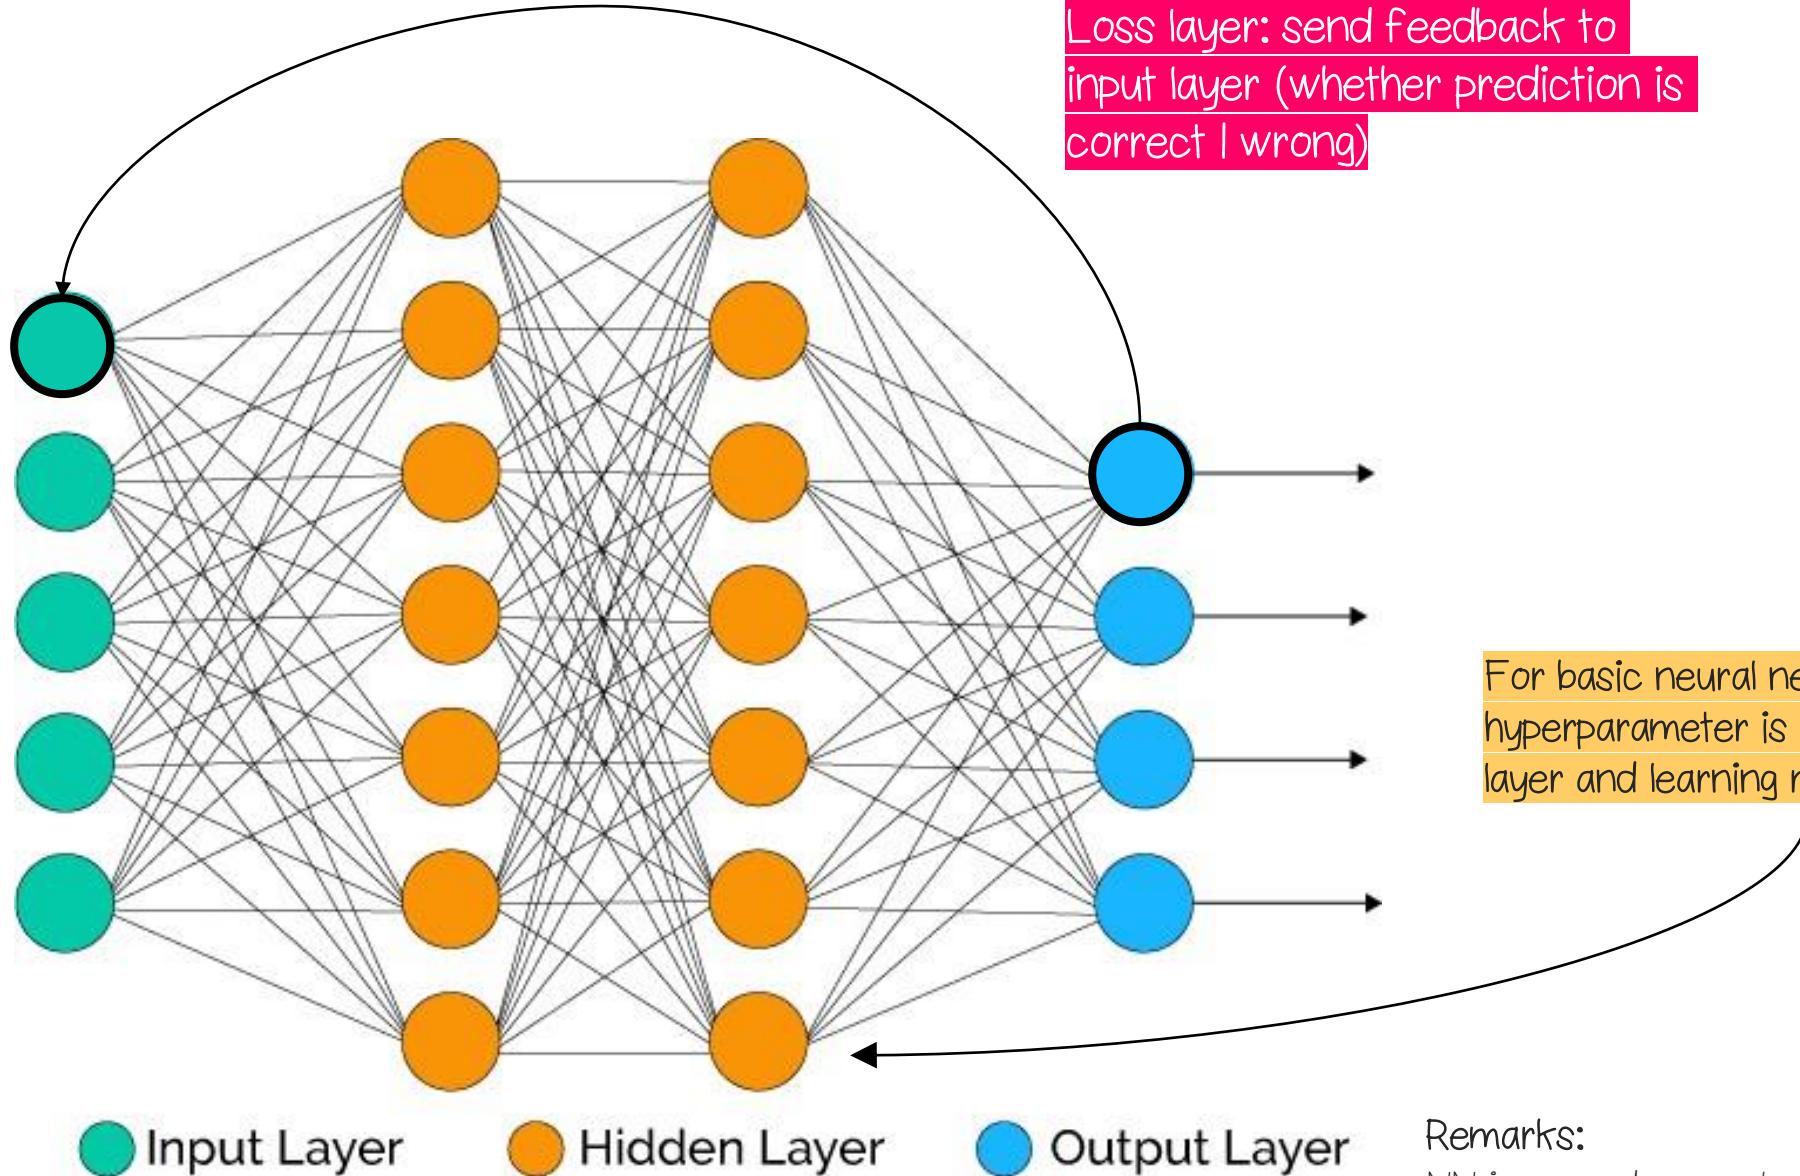


What's going on here?

$$\begin{aligned} \text{Sum all weights} &= \text{Bias} + \text{Weight}(x_1) + \text{Weight}(x_2) \\ &= 0.5 + 1 + 1 = 2.5 \text{ still less than } 4 \end{aligned}$$

: neuron not activate

: perceptron learn from this experience and adjust the weight



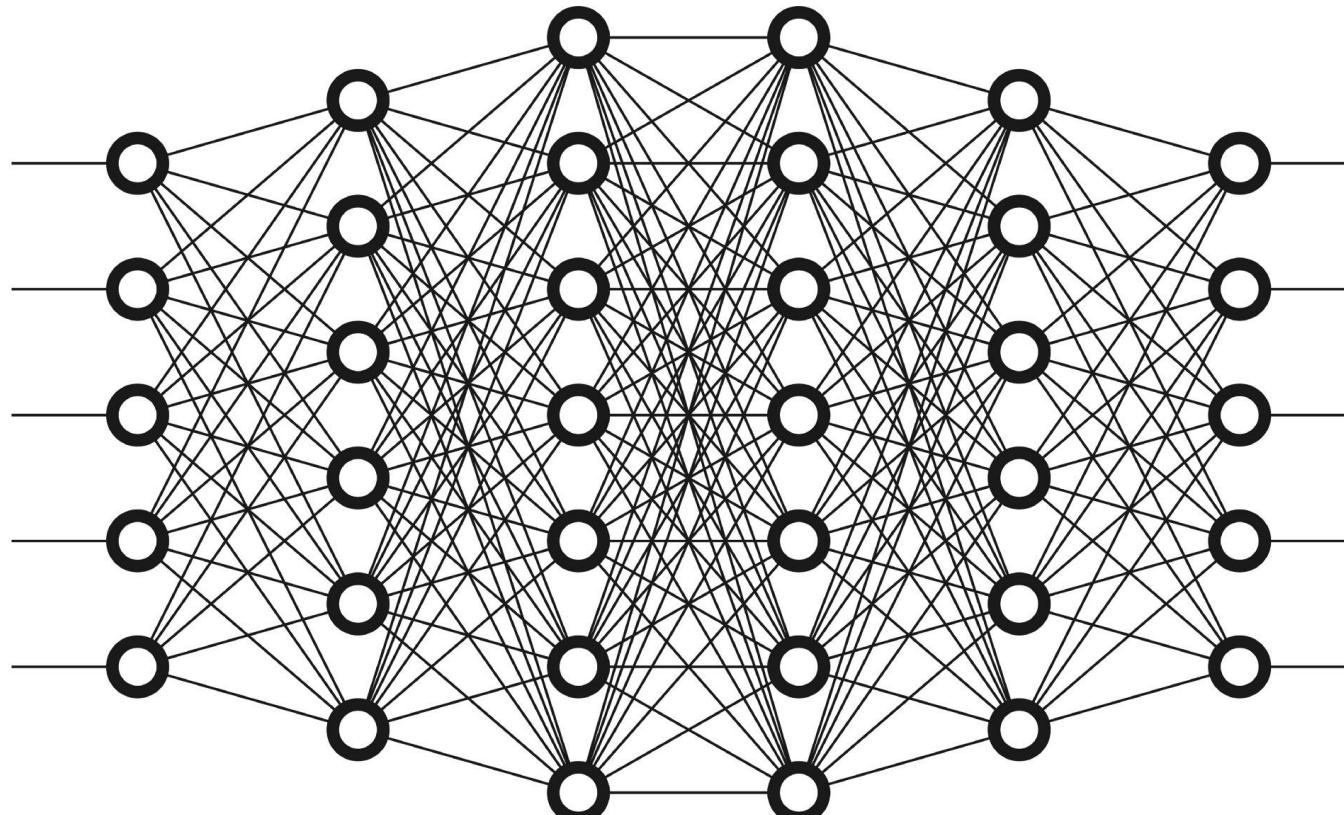
For basic neural networks,
hyperparameter is number of hidden
layer and learning rate (α)

Remarks:

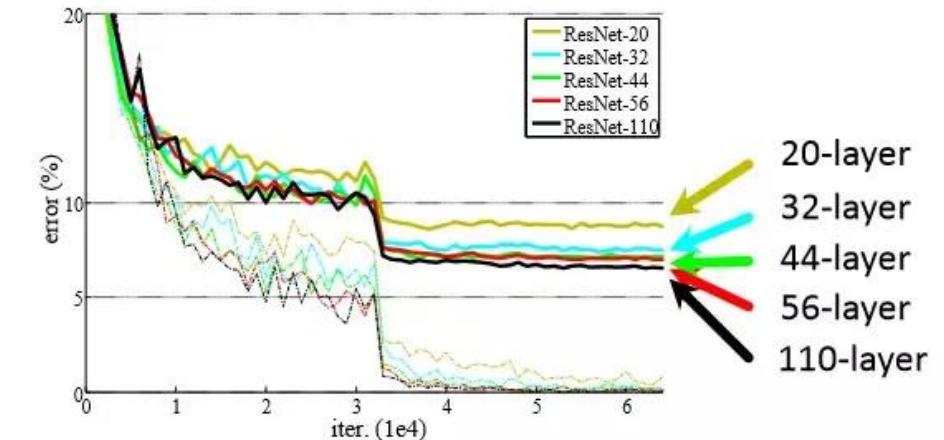
NN is a good example of supervised learning that also
a reinforcement learning (learn from experience)

Deep Learning

Expands the number of hidden layers



CIFAR-10 ResNets



Hieu Pham, Has done some machine learning
Updated Apr 29, 2017

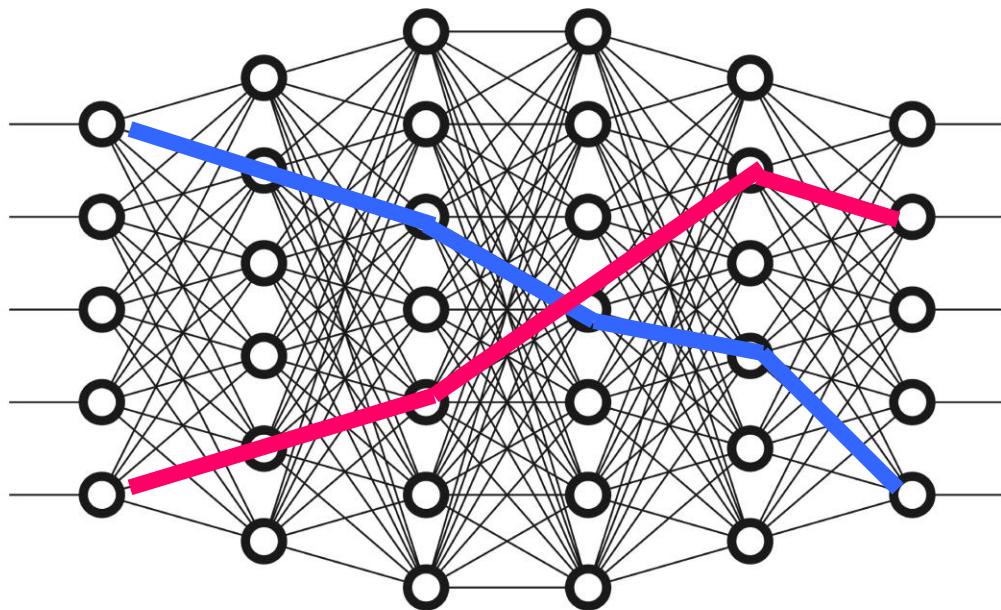


In this [paper](#) (Huang et al, 2016), the authors report that they “can increase the depth of residual networks **even beyond 1200 layers** and still yield meaningful improvements in test error (4.91% on CIFAR-10)”.

I have never heard of anything deeper than that.

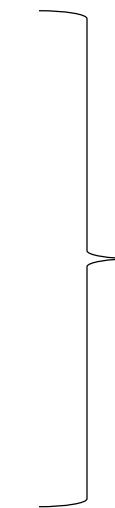
Backpropagation

The backbone of artificial neural networks



Wrong

Right

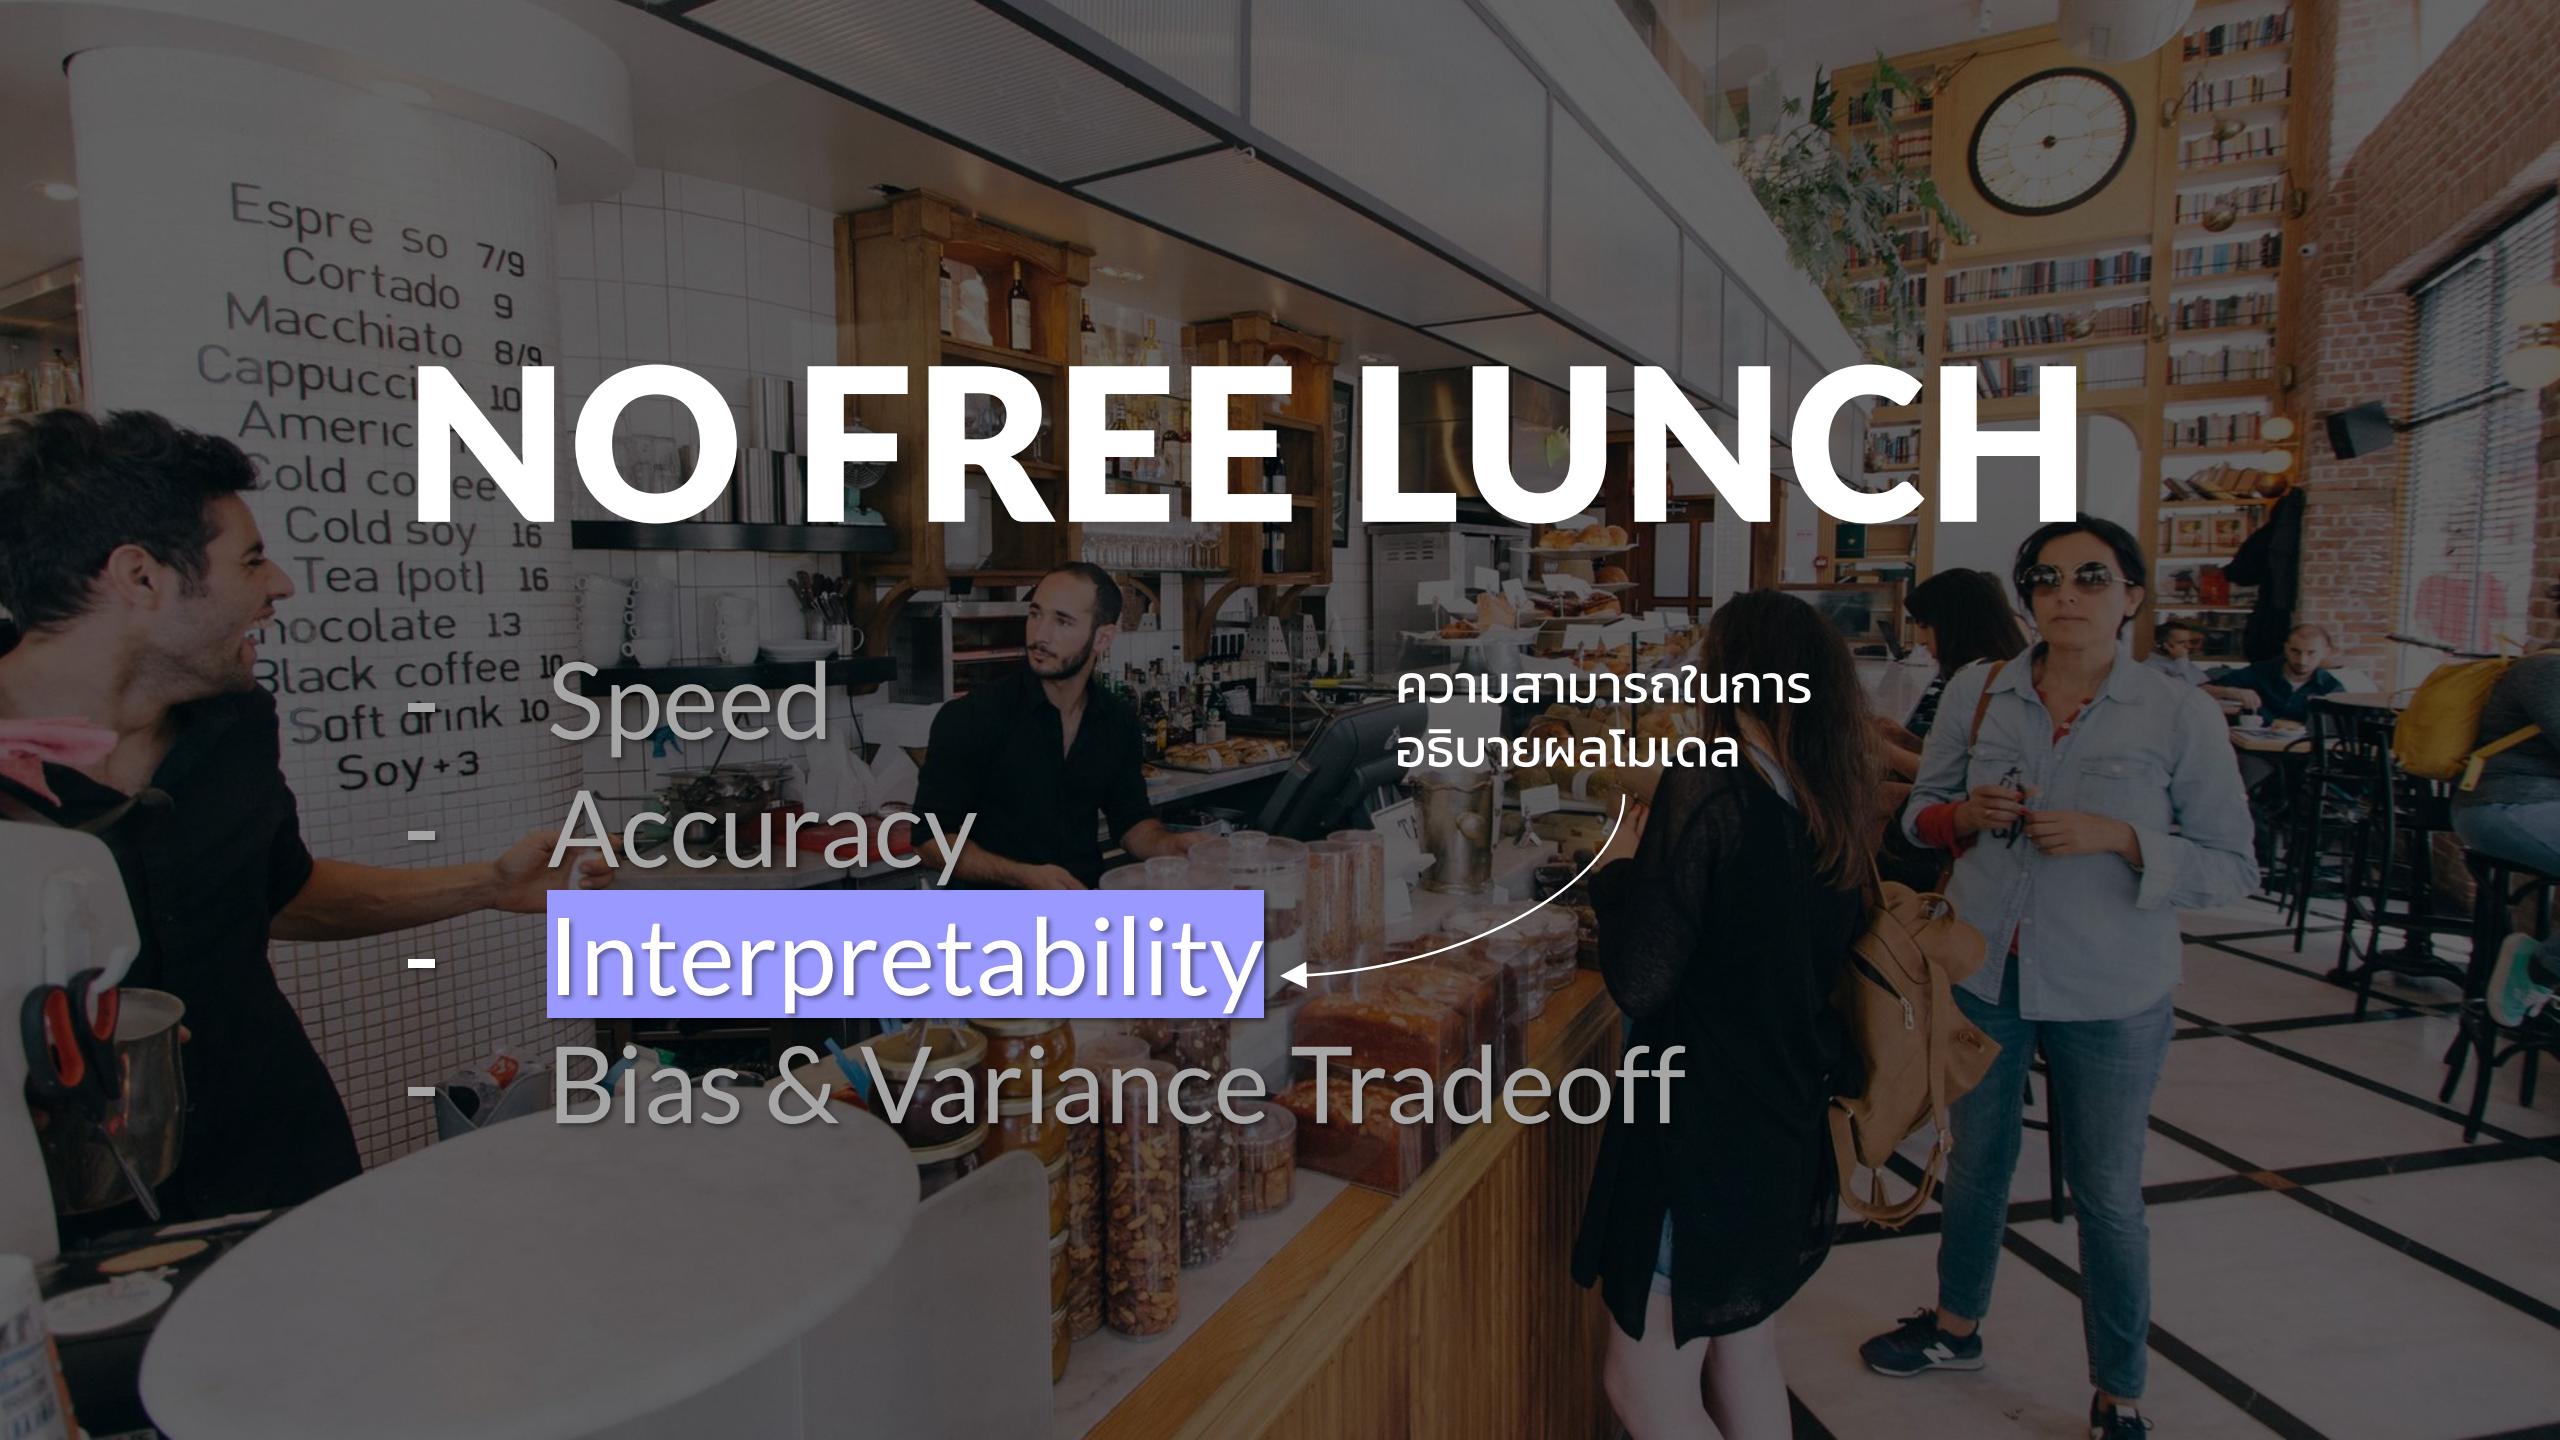


Send feedback to
input layers and
adjust the weights

NO FREE LUNCH

- Speed
- Accuracy
- Interpretability
- Bias & Variance Tradeoff

ความสามารถในการ
อธิบายผลไมเดล



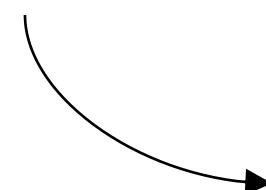
Advantages:

- Amongst the very best performers | learners in the world
- Can be developed into deep learning algorithms
- Work with both regression and classification problems

Disadvantages:

- Requires large data set
- Black Box
- Prone to overfitting problem
- Need more power if more hidden layers

Time to buy a
graphic card :D



Required Packages:

`library(caret)` – train & test models

`library(caTools)` – split data into training & testing sets

`library(nnet)` – used to train neural network in caret

Neural Network

```
method = 'nnet'
```

Type: Classification, Regression

Tuning parameters:

- `size` (#Hidden Units)
- `decay` (Weight Decay)

Required packages: `nnet`

A model-specific variable importance metric is available.

UNSUPERVISED: K-MEANS

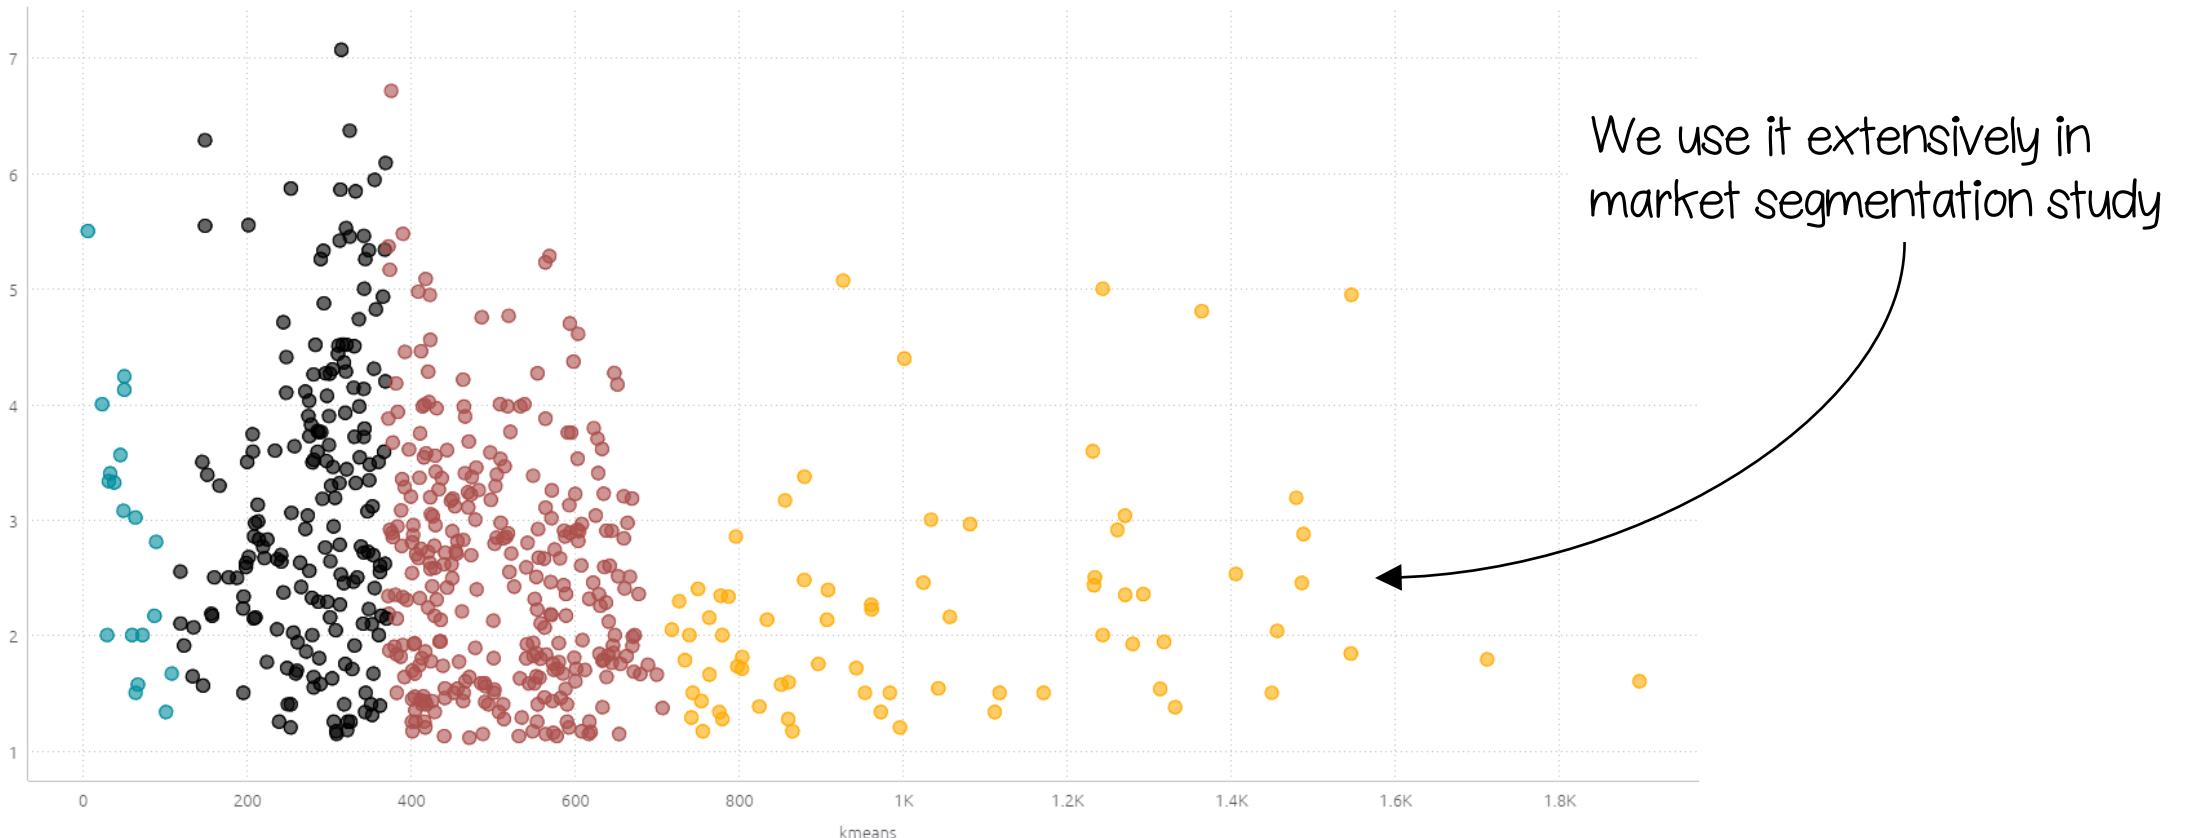
Recall the main task of unsupervised algorithm:

To find patterns in our data set

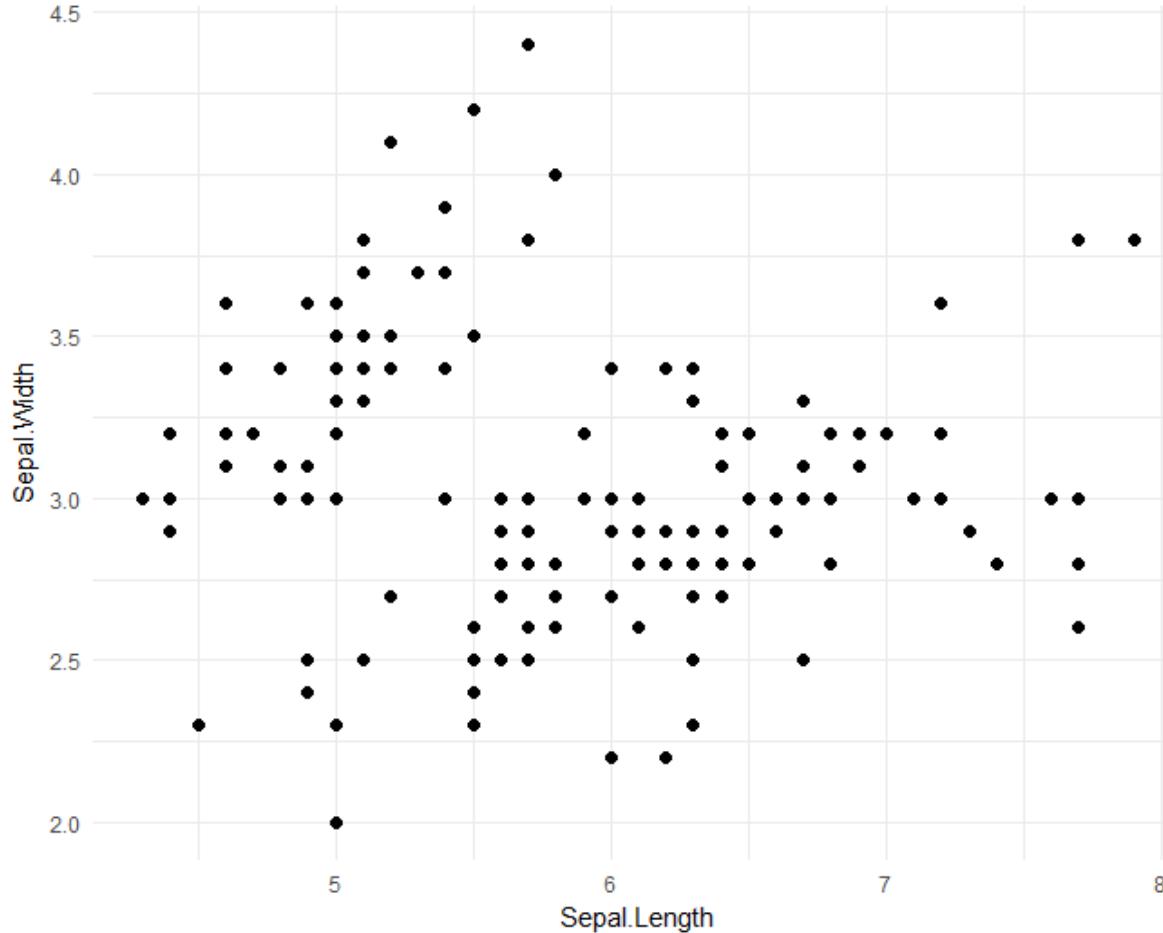
(Descriptive – no right | wrong answer) ←

It's very subjective. But there are some advanced techniques to guide our decisions (not discussed here).

K-Means is very popular unsupervised learning algorithms



K-Means is **very efficient**, easily done in one line



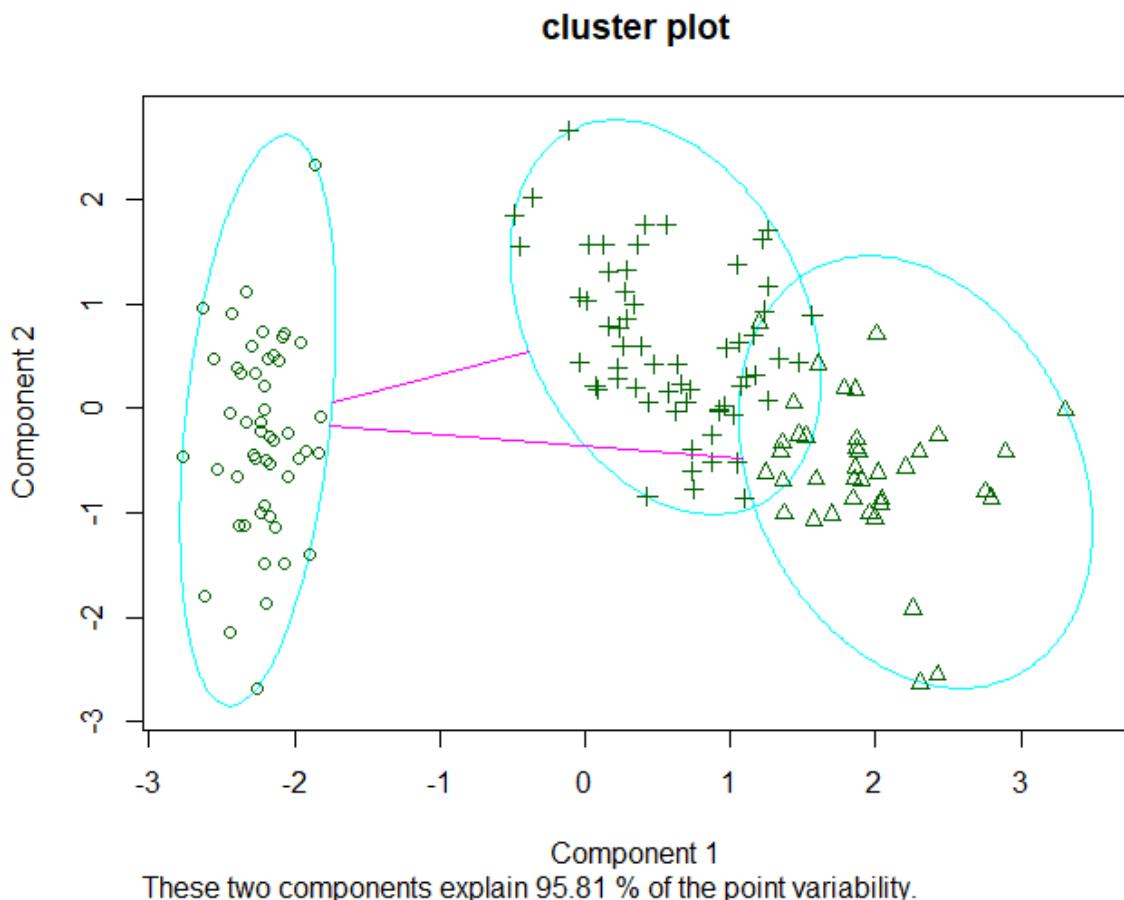
```
> Library(ggplot2)  
> ggplot(iris, aes(x=Sepal.Length, y=Sepal.Width)) +  
  geom_point(size = 2) + theme_minimal()
```

```
> set.seed(2)  
> km_fit <- kmeans(iris[1:4], centers = 3)
```

Specify (guess)
number of clusters



In reality, we don't know how accurate our model is.



```
> library(clusplot)  
> clustplot(iris[1:4], km_fit$cluster, main = "cluster plot")
```

```
> # cross-check with actual class in iris dataframe  
> table(km_fit$cluster, iris$Species)
```

K-means Cluster

Actual Class

Advantages:

- Very easy to train
- Often produce reasonable classification results

Disadvantages:

- Number of clusters sometimes difficult to determine
- Not all features are numeric (can't find mean)

ALGORITHMS ARE NOT CREATED EQUAL

- **Neural Nets** [Backpropagation]
- **SVM** [Kernel Tricks]
- **KNN** [Majority Vote]
- **Random Forest** [Randomness]
- **Decision Tree** [Recursive Partitioning]
- **Logistic Regression** [Sigmoid Function]
- **Linear Regression** [Least Squared]



*R has powerful package written by
Max Kuhn (10 years in development)*

```
# CARET  
# Classification And Regression Training  
# 238 models can be trained using caret  
install.packages("caret")  
library(caret)
```

*For more details, go to this website
<http://topepo.github.io/caret/index.html>*

Let's Recap

- **Congratulations!** Today you've learned some of the most powerful algorithms in ML world
 - *Decision Tree*
 - *Random Forest*
 - *Logistic Regression*
 - *Linear Regression*
 - *KNN*
 - *SVM*
 - *Neural Network*
 - *K-Means Clustering*

Understanding important building blocks to test and implement your own algorithms in R

Be able to discuss the concepts of ML and learning algorithms we've learned this morning

Stay tuned for new online courses on <https://datarockie.teachable.com>

Thanks for coming &
Welcome to DS journey

BY DATAROCKIE  X DATA SCIENCE CHILL CHILL 